

DATA BASE MANAGEMENT SYSTEM



Prof. Samiran Chattopadhyay

Prof. Partha Pratim Das
Computer Science & Engineering
Indian Institute of Technology, Kharagpur



INDEX

S.No	Topic	Page No.
<i>Week 1</i>		
1	Course Overview	01
2	Introduction to DBMS/1	18
3	Introduction to DBMS/2	38
4	Introduction to Relational Model/1	59
5	Introduction to Relational Model/2	80
<i>Week 2</i>		
6	Introduction to SQL/1	100
7	Introduction to SQL/2	119
8	Introduction to SQL/3	143
9	Intermediate SQL/1	166
10	Intermediate SQL/2	194
<i>Week 3</i>		
11	Advanced SQL	218
12	Formal Relational Query Languages	249
13	Entity-Relationship Model/1	277
14	Entity-Relationship Model/2	297
15	Entity-Relationship Model/3	317
<i>Week 4</i>		
16	Relational Database Design	343
17	Relational Database Design (Contd.)	366
18	Relational Database Design /3	387
19	Relational Database Design (Contd.)	412
20	Relational Database Design /5	444
<i>Week 5</i>		
21	Application Design and Development/1	469
22	Application Design and Development/2	495
23	Application Design and Development/3	518
24	Storage and File Structure/1: Storage	543
25	Storage and File Structure/2: File Structure	568
<i>Week 6</i>		
26	Indexing and Hashing/1 : Indexing/1	586

27	Indexing and Hashing/2 : Indexing/2	603
28	Indexing and Hashing/3 : Indexing/3	626
29	Indexing and Hashing/4 : Hashing	650
30	Indexing and Hashing/5 : Index Design	678

Week 7

31	Transactions/1	698
32	Transactions/2 : Serializability	716
33	Transactions/3 : Recoverability	735
34	Concurrency Control/1	759
35	Concurrency Control/2	783

Week 8

36	Recovery/1	798
37	Recovery/2	825
38	Query Processing and Optimization/1 : Processing	848
39	Query Processing and Optimization/2 : Optimization	876
40	Course Summarization	897

Database Management System
Prof. Partha Pratim Das
Department of Computer Science & Engineering
Indian Institute of Technology, Kharagpur

Lecture - 01
Course Overview

(Refer Slide Time: 00:45)

Module Objectives

- To understand the importance of database management systems in modern day applications
- To Know Your Course

SWAYAM: NPTEL-NOCO MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018

PPD

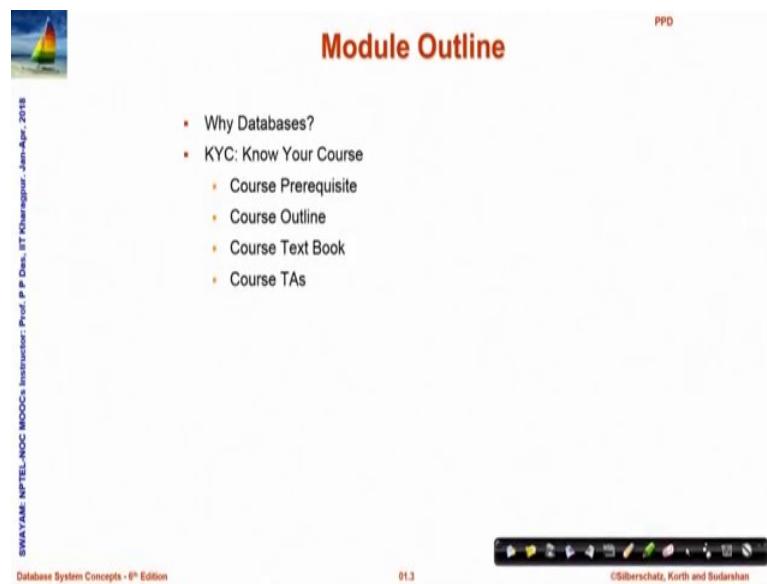
Database System Concepts - 8th Edition

01.2

©Silberschatz, Korth and Sudarshan

Welcome to database management systems. In this course, we will have 40 modules; each module would be of about half an hour. So, this is the first module, where we would talk about the overview of the course. So, we will discuss the importance of database management systems in modern day applications, and we will familiarize you with different aspects of the course.

(Refer Slide Time: 01:06)



The slide is titled "Module Outline" in red at the top right. In the top left corner, there is a small logo of a sailboat on water. The title "Module Outline" is centered above a bulleted list of course topics. The footer contains copyright information and navigation icons.

PPD

Module Outline

- Why Databases?
- KYC: Know Your Course
 - Course Prerequisite
 - Course Outline
 - Course Text Book
 - Course TAs

SWAYAM: NPTEL-NOOCs Instructor: Prof. P. P. Den, IIT Kharagpur - Jam-Apr-2018

Database System Concepts - 8th Edition

01.3

©Silberschatz, Korth and Sudarshan

So, this will be the outline. First, we will try to explain why we need databases, and then we will run through a KYC on the course prerequisites, course outline, the textbook and the TAs who will help us in this course.

(Refer Slide Time: 01:28)



The slide features a large red title "WHY DATABASES?" in the center. In the top right corner, there is a small logo of a sailboat on water. To the right of the title, there is a bulleted list of two items. The footer contains copyright information and navigation icons.

PPD

WHY DATABASES?

- Why Databases?
- KYC: Know Your Course

SWAYAM: NPTEL-NOOCs Instructor: Prof. P. P. Den, IIT Kharagpur - Jam-Apr-2018

Database System Concepts - 8th Edition

01.4

©Silberschatz, Korth and Sudarshan

So, first why do we need databases?

(Refer Slide Time: 01:32)

The slide has a title 'Database Management System (DBMS)' in red at the top right. To its left is a small sailboat icon. On the far left, vertical text reads 'SWAYAM-NPTEL-NOOCs Instructor: Prof. P. P. Doss, IIT Kharagpur' and 'Date: June-Apr., 2018'. At the bottom left is the text 'Database System Concepts - 8th Edition'. The main content is a bulleted list:

- DBMS contains information about a particular **enterprise**
 - Collection of interrelated **data**
 - Set of **programs** to access the data
 - An **environment** that is both *convenient* and *efficient* to use
- Database **Applications**:
 - Banking: transactions
 - Airlines: reservations, schedules
 - Universities: registration, grades
 - Sales: customers, products, purchases
 - Online retailers: order tracking, customized recommendations
 - Manufacturing: production, inventory, orders, supply chain
 - Human resources: employee records, salaries, tax deductions
 - ...
- Databases can be very **large**
- Databases touch **all aspects** of our lives

At the bottom right is a small toolbar with icons for search, refresh, etc.

A database management system contains information about a particular enterprise. So, it deals with collection of interrelated data. We have a set of programs which access and manipulate that data; and together a DBMS presents an environment for convenient as well as efficient use of data of the enterprise under consideration. So, there could be different database application.

Actually, if we look around in the world that we are living today, every aspect of our life has been touched by certain database application. Some of the very common and wide applications will include banking; we are doing net banking very regularly. Net banking is highly distributed database application that allow us to do different kinds of retail as well as corporate banking.

We perform different kind of booking reservation, railway reservation, airline reservation, hotel reservations all of these are now possible over internet as well these are all big database applications. When we are attending colleges, universities, the students, courses, teachers, course enrolments, performance of the students, and different courses, examination all are parts of huge university database applications.

There are different sales applications, online retailing applications like I am sure all of you have been using some of the Amazon, Flipkart, E-bay, Snapdeal, all these are online retail database applications which allow us to order to select items to make payments and to track the deliveries of different items that we have ordered. There are database

applications in terms of manufacturing, production systems, inventory management, any big factory of manufacturing need to use a huge database application to manage the supply chain, the inventory, the orders everything. There are database applications in human resources applications like LinkedIn are huge human resource application of database, the social media applications all involve different kinds of database applications.

So, in this database management system course, we are going to understand how such applications can be designed developed and managed over a period of time. Database applications are typically characterized by the fact that they are very large. If we have a relatively small set of data then possibly we will be able to manage it in terms of using an excel sheet or a couple of excel sheets, but soon when it goes beyond a certain size we need a database application. So, the fact of being large is a critical factor of any DBMS. So, with all this we as we observe the database says cover all aspects of our life and hence understanding DBMSs is a critical requirement for any computer science information technology student.

(Refer Slide Time: 05:54)

University Database Example

- Application program examples
 - Add new students, instructors, and courses
 - Register students for courses, and generate class rosters
 - Assign grades to students, compute grade point averages (GPA) and generate transcripts
- In the early days, database applications were built directly on top of file systems

SWAYAM: NPTEL-MOOCs Instructor: Prof. P. P. Chou, IIT Kharagpur - Date: April 2018

Database System Concepts - 8th Edition

01.6

Gilbert

So, if we talk about a specific database example say university database, then we will have several application programs which will allow us to do several required applications like when new students join, we would need to add the students. We will need to add new courses when courses are floated; we will need to add new instructors

when faculty join; we will need to do allotments; we will need to do registration of students for courses. We will need to conduct examinations, we will need to assign grades to students when they are graded for different courses and compute their GPA and so on. So, all the activities that a university has to deal with are application programs of varied kinds that the university database need to work with.

In the earlier days, before the days of databases, typically such applications used to be managed through file systems. We all know that all our system has a file system, where different files text as well as data files can be stored, and information can be written in these files in a certain order. And just to quickly recall file systems typically have a large number of sequential files which can be written and read in a certain order and some random access files where you can reach a particular point in the file to do certain access operation certain manipulation operations. So, in the earlier days, it was a collection of file systems which managed large enterprise data as is required.

(Refer Slide Time: 07:53)

SWAYAM NPTEL-NOC MOOCs Instructor: Prof. P. K. Khuriya Date: 17-Apr-2018

Drawbacks of using file systems to store data

- Data **redundancy** and **inconsistency**
 - Multiple file formats, duplication of information in different files
- Difficulty in **accessing data**
 - Need to write a new program to carry out each new task
- Data **isolation**
 - Multiple files and formats
- **Integrity** problems
 - Integrity constraints (e.g., account balance > 0) become "buried" in program code rather than being stated explicitly
 - Hard to add new constraints or change existing ones

Database System Concepts - 6th Edition 01.7

But over time it was observed that the file systems to store data to manage data has a lot of drawbacks. For example, if you look at in a file system, there is often a lot of data redundancy and inconsistency. These are terms which we will loosely define here and as we go across along the course, we will understand these terms better. But redundancy just to explain redundancy is a concept where the same data is written at multiple places in different forms and that may give rise to several forms of inconsistencies because if

you write the same data in multiple files, because you need to deal with many of these aspects.

So, there is a file for students, there is a file for teachers, there is a file for particular courses, there is a file for enrolment and so on several data items may be redundantly copied in multiple files. And once that is done then we can it is very easy to get inconsistent in terms of the data because you may update the data in one file, and may forget to update the data in another file, so that is one of the first problems with using the file systems.

Then this difficulty in accessing the data because as I said the data is the files often are sequential in nature, even if they are random access, then every task might need to use data from multiple files. And opening those files and reaching to the appropriate point of access is a non-trivial task. Then there is issues of data isolation, because there are multiple types of files, there are multiple formats used therein. Very importantly there are a lot of integrity problems, the database any database application need to have a lot of integrity.

For example, if you want to withdraw money from an account - bank account, then certainly the balance needs to be positive, you can withdraw only that much amount up to that much amount, which exists in the account. So, any application will need to check for this. So, if you use a file system based application to store the data then at every point wherever you are updating the balance, you will need to make such checks which make the application quite complicated, and often creates the possibility that certain integrity checks may be missed out. So, it is hard to code these new constraints over a period of time.

(Refer Slide Time: 10:40)

The slide has a header 'Drawbacks of using file systems to store data (Cont.)' with a sailboat icon. The content includes a bulleted list of drawbacks:

- **Atomicity** of updates
 - Failures may leave database in an inconsistent state with partial updates carried out
 - Example: Transfer of funds from one account to another should either complete or not happen at all
- **Concurrent access** by multiple users
 - Concurrent access needed for performance
 - Uncontrolled concurrent accesses can lead to inconsistencies
 - Example: Two people reading a balance (say 100) and updating it by withdrawing money (say 50 each) at the same time
- **Security** problems
 - Hard to provide user access to some, but not all, data

At the bottom, it says 'Database systems offer solutions to all the above problems'.

Then there are issues of atomicity of update. What atomicity means is the ability to do certain operations in a as a single unit. So, what you want is either that operation happens or the and if it happens then it happens in full in totality, whereas otherwise it may not happen at all. For example, consider that there is a funds being transformed from one account to another, so this means the account from which the funds are being transferred needs to be debited certain amount, and that same amount has to get credited to the accounts to which it is being paid.

Now, if for some reason of failure or because of the fact that there was link issues or something, if you are not able to make this whole transfer, then it is possible that you have already debited the account, but you have not been able to credit that account. Now, this will be a major cause of inconsistency in the database. So, what you want is if the transfer can happen then it must happen in totality that is what that debit and the credit must happen together or nothing should happen at all. So, there are several examples of requirement of automaticity for a update which is critical for maintaining consistency of the data.

The other aspect which has become very, very deeply required in every aspect of database is concurrency of access. If there is a database then certainly there is not a single user, there are multiple hundreds of users you think about net banking, you think about railway reservation, multiple users are trying to make bookings in multiple trains

from varied stations to vary stations in different classes and so on. So, all of this must go on at the same time that is what is called the concurrency of update.

So, it is quite possible that while you are trying to update, you check berth availability on a certain train on a certain date that you intend to travel. At the same time, someone else may be checking for the berth availability on the same train on the same date, and there could be conflict of concurrency because there may be one berth available and you are trying to book that you have seen that one berth is available.

So, you go ahead and book it try to book it, and there is another user who also saw that one berth is available and that user makes payments and tries to book that. So, concurrency needs to make sure that both of these users should not be allowed to make the booking to the same berth because then that will be a disaster. So, uncontrolled concurrency can add to several inconsistencies in the application.

Then certainly there are you all would be very familiar that today we are living under a whole lot of security threats. So, there has to be proper security that it should be possible to access the data by a user to the extent the user is allowed to do that. So, as a user you should be able to access certain parts of the data, the manager of the system, the administrator should be able to access a bigger part of the data possibly. So, security is hard to provide in terms of a file system based applications to store data. So, all these with we conclude the data based systems had those which provide solution to take care against all those above problems, and we are trying to learn how to do such things.

(Refer Slide Time: 14:37)

The slide features a sailboat icon in the top left corner. In the top right, there are two bullet points: '• Why Databases?' and '• KYC: Know Your Course'. The title 'KNOW YOUR COURSE' is centered in large red capital letters. At the bottom, there is a navigation bar with icons for back, forward, search, and other presentation controls. The footer contains the text 'SWAYAM: NPTEL-NOC MOOCs Initiator: Prof. P. P. Das, IIT Kharagpur - Jain-Apr.-2018', 'Database System Concepts - 8th Edition', '01.8', and '©Silberschatz, Korth and Sudarshan'.

So, moving on I would quickly take you through familiarizing with you with the overall plan of the course.

(Refer Slide Time: 14:47)

The slide features a sailboat icon in the top left corner. The title 'Course Prerequisites' is centered in large red capital letters. Below it, under the heading '• Essential', is a list of topics: 'Set Theory' (with sub-points: Definition of a Set, Intentional Definition, Extensional Definition, Set-builder Notation), 'Membership, Subset, Superset, Power Set, Universal Set', 'Operations on sets' (with sub-points: Union, Intersection, Complement, Difference, Cartesian Product), 'De Morgan's Law', and 'MOOCs: Discrete Mathematics' (with a link: https://npTEL.ac.in/noc/individual_course.php?id=noc16-ma01). At the bottom, there is a navigation bar with icons for back, forward, search, and other presentation controls. The footer contains the text 'SWAYAM: NPTEL-NOC MOOCs Initiator: Prof. P. P. Das, IIT Kharagpur - Jain-Apr.-2018', 'Database System Concepts - 8th Edition', '01.10', and '©Silberschatz, Korth and Sudarshan'.

So, what we first talk about are the course prerequisites. These prerequisites are kind of certain elementary level knowledge in computer science and related discrete mathematics that you should have. You should be that would make it easier for you to understand and follow the course; otherwise if you find at any stage that you are finding any of this aspect difficult to follow in the course then I would advise that you go back to

some of the background material and try to study them. So, I have tried to list down these prerequisite topics, one certainly is the set theory because that is the basic premise on which databases are designed on.

So, starting from definition of the set membership, concepts of subset, superset, power set, different operations of union, intersection, complementation, difference Cartesian product De Morgans law, all these basic set theory you should be very familiar and conversant with. If you are not I have mentioned one MOOC's course this is a past course, but you can access the videos and the contents which has a very nice discussion on these aspects of set theory which you may refer to.

(Refer Slide Time: 16:13)

Course Prerequisites

PPD

- * Essential
 - Relations and Functions
 - › Definition of Relations
 - › Ordered Pairs and Binary Relations
 - Domain and Range
 - Image, Preimage, Inverse
 - Properties – Reflexive, Symmetric, Antisymmetric, Transitive, Total
 - › Definition of Functions
 - › Properties of Functions – Injective, Surjective, Bijective
 - › Composition of Functions
 - › Inverse of a Function
 - MOOCs: Discrete Mathematics:
https://nptel.ac.in/noc/individual_course.php?id=noc16-ma01

SWAYAM NPTEL-MOOCs Instructor: Prof. P. P. Doshi, IIT Kharagpur - 2016

Database System Concepts - 8th Edition

01.11

©Silberschatz, Korth and Sudarshan

Moving on the next, which goes on top of the sets are the concept of relations and functions. We all know that a relation is a subset of a set. So, if I have a set A, then I can define a binary relation over that set A, which is basically the pair of elements from set A or in other words the relation binary relation is a subset of the cross product of A with itself where the domain and the range are related together. So, there are concepts of the image of a domain the pre-image of the range the inverse relation, and several basic properties of relations like the relation being reflexive, symmetric, anti-symmetric, transitive, total relations and so on.

You should be familiar with these; otherwise, you will find difficult to follow major concepts in the database systems because this tribal system primarily the one that we are

going to take you through in this course is relational in model. So, it is heavily based on relations and functions. And you should understand one specifically a relation becomes a function, and when what is meant by functions being injective, subjective, bijective, what is meant by composition, and inverse of functions and so on. Again if you need, you can refer to this MOOC's course on discrete mathematics to brush up your knowledge about relations and functions.

(Refer Slide Time: 17:52)

The screenshot shows a presentation slide titled "Course Prerequisites". At the top right is a small "PPD" icon. On the left is a decorative image of a sailboat on water. The slide content is organized into sections:

- Essential**
 - Propositional Logic**
 - Truth Values & Truth Tables
 - Operators: conjunction (and), disjunction (or), negation (not), implication, equivalence
 - Closure under Operations
 - MOOCs: Discrete Mathematics:**
https://nptel.ac.in/noc/individual_course.php?id=noc16-ma01
 - Predicate Logic**
 - Predicates
 - Quantification – Existential, Universal
 - MOOCs: Discrete Mathematics:**
https://nptel.ac.in/noc/individual_course.php?id=noc16-ma01

At the bottom of the slide, there is footer text: "SWAYAM-NPTEL-MOOCs Instructor: Prof. A.P.Datta, IIT Kharagpur - Jan-Apr-2015", "Database System Concepts - 6th Edition", "01.12", and "©Silberschatz, Korth and Sudarshan".

We also need you to have a basic understanding of the propositional logic which is truth values true and false; and the different operations of conjunction, disjunction, negation that is and or not, what is meant by implication, what is meant by equivalence. You know that given two variables, which have or two propositions which can take a value true or false the conjunction of them can be represented in terms of a truth table where we say that only when both these propositions are true, then the resultant conjunctive proposition becomes true; otherwise, a conjunctive proposition is false. So, you should be familiar with these concepts; if you are not, please brush up your ideas about propositional logic.

We need a little bit of predicate logic as well a predicate logic in contrast to propositional logic deals with quantification that the knowledge of existential and universal quantifier. When you say that whether certain proposition predicates hold for all values in the domain or for some value in the domain whether there exists some value for which it

holds or whether for all values it holds. And based on that predicate logic is build up we do not need very advanced concept here just basic level familiarization will help and the same MOOC's course on discrete mathematics would be of your help in case you need to brush it up further.

(Refer Slide Time: 19:33)

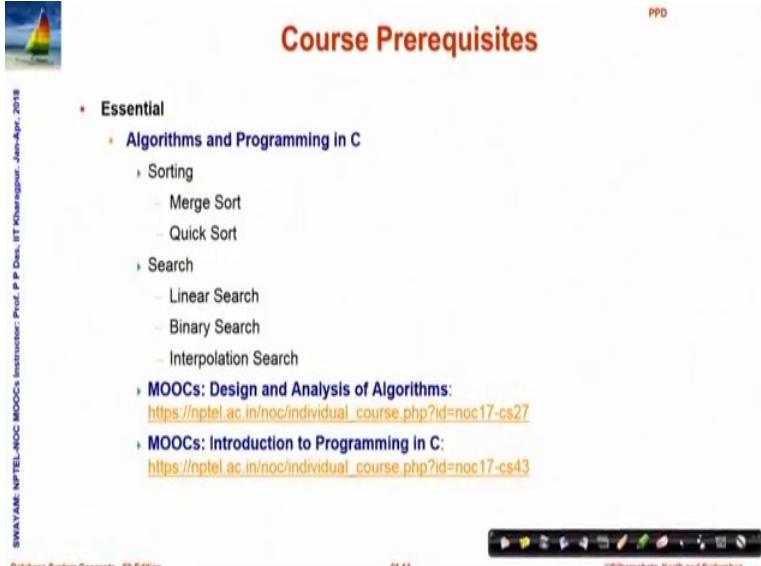
The screenshot shows a presentation slide titled "Course Prerequisites" in red at the top right. On the left, there is a small logo of a sailboat on water. The main content is a bulleted list under the heading "Essential".

- * Essential
 - Data Structures
 - Array
 - List
 - Binary Search Tree
 - Balanced Tree
 - B-Tree
 - Hash Table / Map
 - MOOCs: Design and Analysis of Algorithms:
https://nptel.ac.in/noc/individual_course.php?id=noc17-cs27
 - MOOCs: Fundamental Algorithms – Design and Analysis:
https://nptel.ac.in/noc/individual_course.php?id=noc16-cs24

On aspects of computer science, certainly you need a good familiarity with the data structures array list particularly binary search tree, what is called a binary search tree, what is meant by a height of a binary search tree, when we say that a binary search tree is balanced. What are the ways and conditions of balancing is particularly the B-trees for organizing good search trees hash tables, what is hashing we need you to be familiar with this concept, because the databases will be heavily designed based on the concepts of B trees and hash tables and so on.

There are courses on design and analysis of algorithms, fundamental of algorithms have mentioned two excellent courses in MOOC's from which you can brush up if you need to.

(Refer Slide Time: 20:31)



The slide is titled "Course Prerequisites" in red at the top right. It features a small sailboat icon in the top left corner. The background is white with a thin vertical border on the left side. The text is organized into sections with bullet points.

Essential

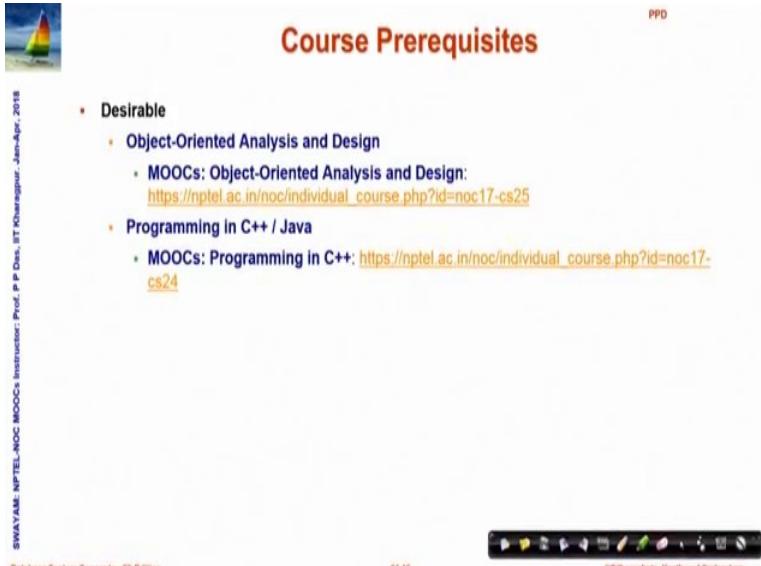
- Algorithms and Programming in C
 - Sorting
 - Merge Sort
 - Quick Sort
 - Search
 - Linear Search
 - Binary Search
 - Interpolation Search
- MOOCs: Design and Analysis of Algorithms:
https://nptel.ac.in/noc/individual_course.php?id=noc17-cs27
- MOOCs: Introduction to Programming in C:
https://nptel.ac.in/noc/individual_course.php?id=noc17-cs43

SWAYAM-NPTEL-MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - 2018
PPD

Database System Concepts - 8th Edition 01.14 ©Silberschatz, Korth and Sudarshan

Certainly you need certain familiarity with common algorithms particularly I had mentioned sorting and searching algorithms, because these are critical for database applications. And again the same MOOC's courses would be of your help. And it will be good to have familiarity with programming in C, because several of the applications need some application high level application programming to be performed. And we would assume that those aspects we described in C because that is a fundamental and most commonly known language.

(Refer Slide Time: 21:11)



The slide is titled "Course Prerequisites" in red at the top right. It features a small sailboat icon in the top left corner. The background is white with a thin vertical border on the left side. The text is organized into sections with bullet points.

Desirable

- Object-Oriented Analysis and Design
 - MOOCs: Object-Oriented Analysis and Design:
https://nptel.ac.in/noc/individual_course.php?id=noc17-cs25
- Programming in C++ / Java
 - MOOCs: Programming in C++:
https://nptel.ac.in/noc/individual_course.php?id=noc17-cs24

SWAYAM-NPTEL-MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - 2018
PPD

Database System Concepts - 8th Edition 01.15 ©Silberschatz, Korth and Sudarshan

Besides these prerequisites which I have marked as essential, because they will certainly be required for the major parts of the course, it will be good if you have some familiarity with object oriented analysis; and design and some language which is more heavily object oriented object aligned in nature like C++ or java. So, again some MOOCs courses the related MOOCs courses are mentioned here in case you need to brush up.

(Refer Slide Time: 21:46)

Course Outline

Week	Topics
Week 1	Course Overview Introduction to RDBMS
Week 2	Structured Query Language (SQL)
Week 3	Relational Algebra Entity-Relationship Model
Week 4	Relational Database Design
Week 5	Application Development Case Studies Storage and File Structure
Week 6	Indexing and Hashing Query Processing
Week 7	Query Optimization Transactions (Serializability and Recoverability)
Week 8	Concurrency Control Recovery Systems Course Summarization

SWAYAM: NPTEL-MOOCs Instruction: Prof. P. P. Desai, IIT Kanpur

PPD

Application Programmer

DBA / DBMS Developer

Database System Concepts - 8th Edition

01.16

©Silberschatz, Korth and Sudarshan

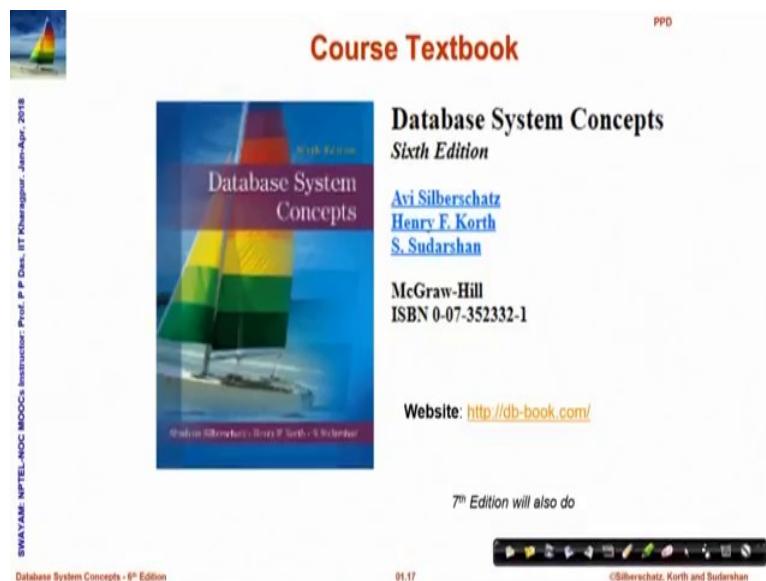
Moving on this is your course outline. So, as I said the course comprise 40 modules. So, it is divided into 8 weeks. So, this plan is given based on what we do in different weeks from week 1 to week 8. So, as the course unfolds, you will be able to we will take you through these modules on these topics. And on the right, you can find that I have marked that the initial part of the course which we cover from week one to first half of week 5 is primarily meant for application programming which it means that the database system has already been designed and basic premises the schemas and constraints have been set up.

But now you want to write different data query different data manipulation applications and that is where the large volume of database engineers work. So, they are called application programmers, so that is I should say the first level in terms of a database understanding. And you must target to become a master of application programming to get started with.

The other half of the course which start from the middle of week five with the storage and file structure and goes on till the next is meant for the analysts who are responsible either for designing a particular database which the application programmer can use tune that for performance index it properly design queries to be efficient. So, these kind of analysts will be involved the more with the understanding of the second part of the course.

And the second part of the course would also be useful for the database a DBMS designers. If you want to really become an advanced programmer we want to work on database engineering in terms of creating database management systems not merely creating databases or database applications, then you need to have a good initial command over the second half of the course. So, while you prepare for the course where you go through the modules please keep this in mind that your familiarity with the application programming must be at the highest level. And in the later parts will be relatively little advanced, but they are required for good design and good development of consistent efficient system in future.

(Refer Slide Time: 24:39)



We will follow a textbook. This is as you can see this is a sixth edition that I am following in this course. It is called Database System Concept by Silberschatz, Korth and Sudarshan. This is kind of a classical book in database systems; current version actually is the seventh edition. So, if you get access to the seventh edition, you can use that as

well, but whatever we are following in this course sixth edition is good enough. So, I advise the, that you try to get a copy of this book to yourself.

(Refer Slide Time: 25:24)

The screenshot shows a presentation slide with the following details:

- Title:** Course TAs
- Content:**
 - Srijoni Majumdar, majumdarsrijoni@gmail.com, 9674474267
 - Himadri B G S Bhuyan, himadribhuyan@gmail.com, 9438911655
 - Gurunath Reddy M, mgurunathreddy@gmail.com, 9434137638
- Navigation:** A standard presentation navigation bar with icons for back, forward, search, and other controls.
- Page Number:** 01/18
- Source:** ©Silberschatz, Korth and Sudarshan
- Other:** A vertical watermark on the left side reads "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Chakrabarti - June-Apr. 2018".

So, moving on we have different three TAs we will help you in this course Srijoni Majumdar, Himadri Bhushan Bhuyan and Gurunath Reddy. So, these are their emails; I have also put their mobile numbers. However, I would advise that unless you are really stuck avoid calling them on the mobile, because they are research students as well, busy with their own work as well, but they can certainly put all your questions on the forum, which will be promptly responded by some of these TAs or by myself. And in case you would have very specific follow ups to do, you can write email to one or all of these TAs. So, this is about the course overview.

(Refer Slide Time: 26:17)

The slide is titled "Module Summary" in red at the top right. In the top left corner, there is a small logo of a sailboat on water. The top right corner has the text "PPD". On the left side, there is vertical text: "SWAYAM NPTEL-NOOCs Instructor: Prof. P. P. Doss, IIT Kharagpur", "Date: 2018", and "Version: 1.0". The main content area contains a bulleted list:

- Elucidates the importance of database management systems in modern day applications
- Introduced various aspects of the Course

At the bottom left, it says "Database System Concepts - 8th Edition". At the bottom center, it shows "01.19". At the bottom right, it says "©Silberschatz, Korth and Sudarshan". A standard Windows-style toolbar is visible at the very bottom.

So, in this module we have discussed about the importance about database management systems in the modern day applications. And we have tried to familiarize you with different aspects of the course. And I reiterate that please give due consideration to the prerequisites as mentioned, we have floated in an assignment called assignment zero where there are questions on different aspects of these prerequisites, please try to solve that assignment and see your performance. This assignment I would like to mention that this assignment will not go in the final evaluation of the performance of the course; this is just to give you an idea for self assessment of your preparedness for this course.

So, if you find that on questions on certain topics say on relation function or on data structure, if you have not been able to answer the questions well then it will be good to check back and go through the prerequisites once more. But please keep in mind that database management system is a course which depends on these background knowledge quite heavily. So, if you have gaps in understanding those topics, then all through the course you will get into several difficulties in understanding and problem solving. So, that is about our first module. So, from the next module, we will start introducing the database management system. Enjoy the course, and try to learn, try to become a good database engineer.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science & Engineering
Indian Institute of Technology, Kharagpur

Lecture - 02
Introduction to DBMS/1

(Refer Slide Time: 00:42)

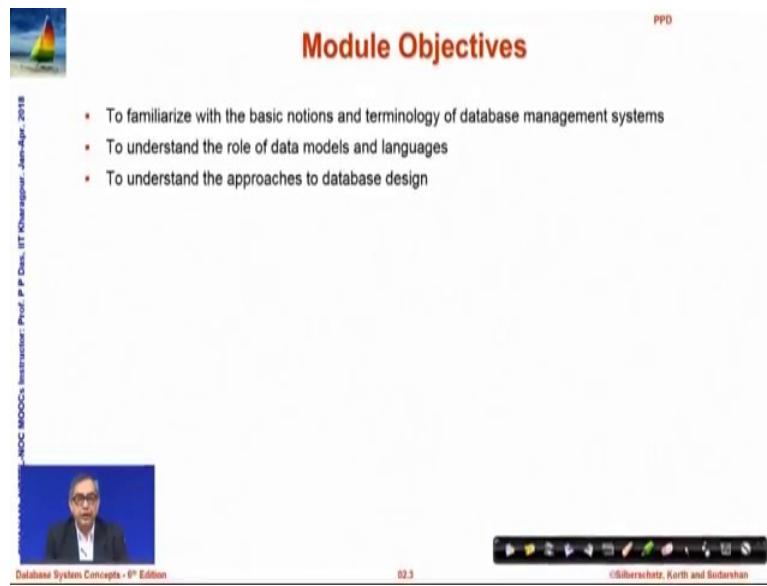
The slide is titled "Module Recap" in red at the top right. It features a small sailboat icon in the top left corner. On the far right, there is a watermark-like text: "PPD". The main content is a bulleted list under the heading "Module Recap". The list includes:

- Why Databases?
- KYC: Know Your Course
 - Course Prerequisite
 - Course Outline
 - Course Text Book
 - Course TAs

At the bottom of the slide, there is footer text: "SWAYAM: NPTEL-MOOCs Initiatives Prof. P. P. Das, IIT Kharagpur - June-Aug. 2018", "Database System Concepts - 8th Edition", "02.2", and "©Silberschatz, Korth and Sudarshan". There is also a set of navigation icons at the bottom.

Welcome to module two of database management systems. In this module, we will talk primarily about the introduction to the DBMS. This discussion will span two modules that is the current one module 2 and the next one that is module 3. Just to quickly recap in the last module we have discussed about why we need databases, and we have introduced you to different aspects of the course.

(Refer Slide Time: 00:54)

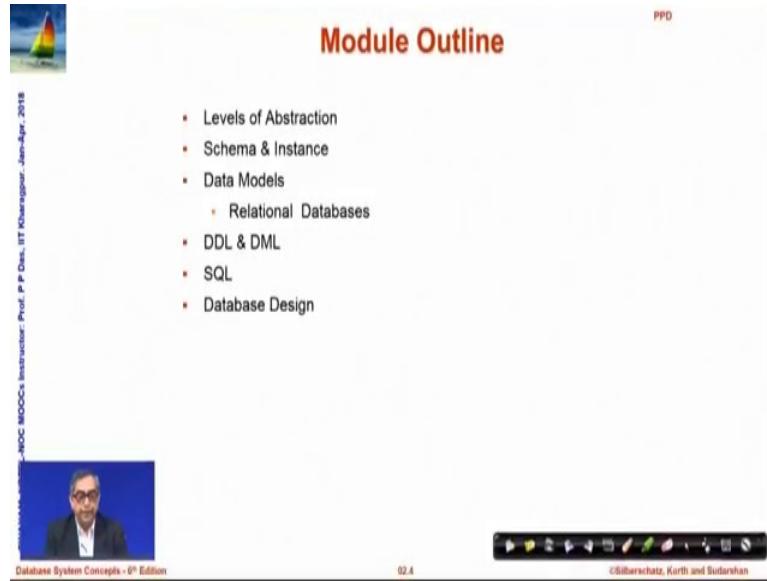


This slide is titled "Module Objectives" in red at the top right. It features a small sailboat icon in the top left corner. On the far left, there is vertical text: "APU MOOCs", "APU MOOCs Instructor", "Prof. P. P. Das, IIT Kharagpur", and "2018". At the bottom left is a video frame showing a man in a suit. The bottom right contains the text "Database System Concepts - 8th Edition", "02.3", and "©Silberschatz, Korth and Sudarshan". A navigation bar with various icons is at the bottom right.

- To familiarize with the basic notions and terminology of database management systems
- To understand the role of data models and languages
- To understand the approaches to database design

In view of this, in this module, we would familiarize you with basic notions and terminology of a database management system just at an introductory level. We will try to understand the role of data models and specific languages for database systems. And we would also outline the approach to database design.

(Refer Slide Time: 01:21)



This slide is titled "Module Outline" in red at the top right. It features a small sailboat icon in the top left corner. On the far left, there is vertical text: "APU MOOCs", "APU MOOCs Instructor", "Prof. P. P. Das, IIT Kharagpur", and "2018". At the bottom left is a video frame showing a man in a suit. The bottom right contains the text "Database System Concepts - 8th Edition", "02.4", and "©Silberschatz, Korth and Sudarshan". A navigation bar with various icons is at the bottom right.

- Levels of Abstraction
- Schema & Instance
- Data Models
 - Relational Databases
- DDL & DML
- SQL
- Database Design

So, the module outline would be like this. And as we go along you will be able to follow which particular aspect of this outline we are discussing about.

(Refer Slide Time: 01:35)

The slide features a title 'LEVELS OF ABSTRACTION' in large red capital letters. To the left of the title is a small video thumbnail showing a person's face. On the right side, there is a vertical list of topics under the heading 'PPD': 'Levels of Abstraction', 'Schema & Instance', 'Data Models', 'DDL & DML', 'SQL', and 'Database Design'. The footer of the slide includes the text 'Database System Concepts - 8th Edition', '02.5', and '©Silberschatz, Korth and Sudarshan'.

So, to start with we talk about levels of abstraction.

(Refer Slide Time: 01:40)

The slide has a title 'Levels of Abstraction' in red. Below the title is a code snippet for a record type 'instructor' with fields ID, name, dept_name, and salary. The slide also lists three levels of abstraction: Physical level, Logical level, and View level. The footer contains the text 'Database System Concepts - 8th Edition', '03.8', and '©Silberschatz, Korth and Sudarshan'.

```
type instructor = record
    ID : string;
    name : string;
    dept_name : string;
    salary : integer;
end;
```

- **Physical level:** describes how a record (e.g., instructor) is stored
- **Logical level:** describes data stored in database, and the relationships among the data
- **View level:** application programs hide details of data types
 - Views can also hide information (such as an employee's salary) for security purposes

A database system like any other system is conceptualized in terms of three levels of abstraction. At the lower most level is the, what we say is a physical data level or the physical level which describes how a record is actually stored, so that is about the physical storage in the system. At the next level, we talk about it we say it is a logical level which describes the data stored in databases and its relationship amongst the data. So, you can any data that is stored you can think about it as a record. So, if we here we

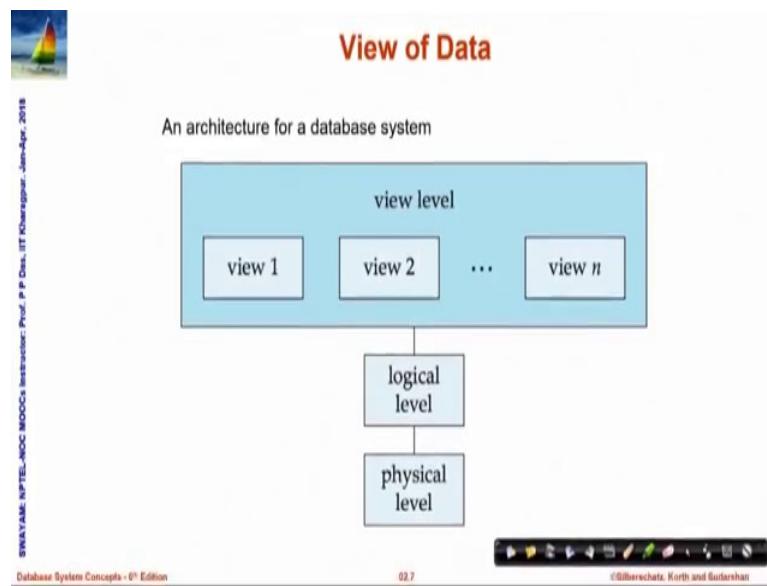
are illustrating the record of an instructor who teaches a course. So, as you know the record is a collection of multiple fields of different types. So, here we have field to describe the ID identifying number or shrink of an instructor, we have the name of the instructor, the name of the department, the salary and so on.

So, this logically says this is the entity this is a record or this is the structure of a record that I want at a logical way. So, this in contrast to the physical level, logical level is not concerned particularly with how that these data, how this string the number and all that will be actually stored, and how these multiples of hundreds and thousands of records will actually be stored so that they can be efficiently used. But we are just concerned with the logical view that I should be able to deal with records as they are.

At the third level which is it can say the topmost level is called the view level where the application program tries to view the data. And when the application program tries to view the data, it deals with the details that it needs to; and rest of the details are usually hidden from this view. For example, if here we talked about the university database in the last module, so if you are talking about the university database, and then you are a student. And you are when you access the database you should be able to see what all courses you are enrolled in or where is that course being held who is the instructor of that course and so on. But you should not be able to access or see the view of what is the instructors salary or for that matter, what are the grades that are obtained to by different students in different courses and so on.

Whereas, an instructor may be able to view the performance of the students in multiple different courses particularly the ones that he or she is involved in evaluation, but she again may not be allowed to check the salary of other instructors. So, view level is a high level where of abstraction where you try to provide the information about the data in terms of what the application needs, what the users of that application need, but you do not actually deal with the details of all the records that the logical level has.

(Refer Slide Time: 05:29)



So, these are three levels form the basic structure of a database system architecture of a database system. As you can see the physical level using that a logical level of records are formed. Physical level typically is in terms of database files is binary in nature, the organization of those files. The logical level deals with the records and the different fields of the records the schema of the database and so on.

And the view level is something which is constructed from the logical level in terms of different views that the different applications need. I am sure at this stage you may not understand the whole of these levels and their implications, but this is just to give you an idea of the existence of three levels, and the need to deal with the three levels. And as we go along, we will see that we will refer to these levels more and more and discuss about the specific aspects of those.

(Refer Slide Time: 06:27)

The slide has a header with a sailboat icon and the text 'PPD'. A vertical sidebar on the left contains the text 'SWAYAM: NPTEL-NOC MOOCs Initiator: Prof. P. P. Das, IIT Kharagpur; Admin-Help: Jahn-Bert - 2018'. The main title 'SCHEMA AND INSTANCE' is in red. Below it is a video player showing a man speaking. The footer includes 'Database System Concepts - 8th Edition', '02.8', and '©Gillberschattz, Korth and Sudarshan'.

Next, let us talk about schema and instance.

(Refer Slide Time: 06:34)

The slide has a header with a sailboat icon and the text 'PPD'. A vertical sidebar on the left contains the text 'SWAYAM: NPTEL-NOC MOOCs Initiator: Prof. P. P. Das, IIT Kharagpur; Admin-Help: Jahn-Bert - 2018'. The main title 'Schemas and Instances' is in red. The content lists components of a schema: 'Similar to types and variables in programming languages', 'Schema' (with 'Logical Schema' defined as the overall logical structure of the database, analogous to type information in a program, with examples of customers and accounts in a bank), 'Customer Schema', 'Account Schema', and 'Physical Schema' (the overall physical structure of the database). The footer includes 'Database System Concepts - 8th Edition', '02.8', and '©Gillberschattz, Korth and Sudarshan'.

We will very regularly keep on referring about schemas and instance. The schema is in a very simple terms say if we talk about first a logical schema, it is a way a certain data needs to be organized, it is a plan for organizing data. So, if you can draw a parallel then say when a building is constructed, a plan is prepared. And according to that plan several buildings in a say residential complex may be constructed. So, there is a difference between the plan which gives you the layout of where different rooms should be where

there is a staircase where is the courtier and so on and the actual building when or the group of buildings which are constructed. So, the schema primarily talks about what is the plan to organize the data.

So, if we talk about a customer schema, it has multiple different fields, it should talk about the name of the customer, ID of the customer, it is account possibly the other ID the mobile number and so on. So, the fact that these the fields need to be present for describing a customer, forms a customer schema. Whereas, when we talk about a specific schema of account that the customer holds with a bank, then we need the account number, account type, interest rate, minimum balance, the current balance and so on. These are the fields of information that we need; and we need to know what is the type of every such field, and all those and those kind of information from the schema information.

And again in line with the abstractions of physical logical and view as we did, schema also can be at a logical schema which is corresponding to the logical level of abstraction. And we may also have a physical schema which tells actually how the data is physically organized in the database, what are the different database files, how they are indexed and so on.

So, all these information which we can say is kind of a metadata information. This is not actually that it is not the customer schema is not saying who the customers are, the account schema is not saying, what are the accounts, what is their balance. But it is saying that if a customer needs to be defined; then what is the information that you need to know, what is the information that you need to manage. If an account need to be described then what is the different fields that are important. So, this schematic or this metadata is called the schema of a database.

(Refer Slide Time: 09:25)

The slide has a header 'Schemas and Instances' and a footer with course details: SYNA1101: Database System Concepts - IIT Kharagpur, Prof. P.P. Chakrabarti, Prof. A.M. Korattiri, and Prof. S. Sudarshan. The footer also includes 'Database System Concepts - 8th Edition', '02.10', and '©Silberschatz, Korth and Sudarshan'.

Instance

- The actual content of the database at a particular point in time
- Analogous to the value of a variable
- Customer Instance

Name	Customer ID	Account #	Aadhaar ID	Mobile #
Pavan Laha	6728	917322	182719289372	9830100291
Lata Kala	8912	827183	918291204829	7189203928
Nand Prabhu	6617	372912	127837291021	8892021892

Account #	Account Type	Interest Rate	Min. Bal.	Balance
917322	Savings	4.0%	5000	7812
372912	Current	0.0%	0	291820
827183	Term Deposit	6.75%	10000	100000

Now, based on this schema specific instances of databases happen, instances when you actually have populated different records according to the schema. Now, naturally once the schema is fixed, your records will need to have values in all of those fields that the schema has; and every value must be of the type that the particular field is specified with. So, I have just shown here certain sample instance of customer schema, where you can see three customers with their name, customer ID, account number, other ID, and mobile number, these are all fictitious data of course, but this is just to give you an idea of how this customer instance would look like.

Similarly, we have shown what is a accounts instance, so you can see that there is some kind of a relationship that you can see between these instances. For example, the first customer ID on the customer instance can be seen as a first entry first record in the account instance and so on. So, when we actually populate the schema with different records and this is what keeps on changing.

So, certainly when we do operations on the database, then certainly very regularly new records will get added, some records might get deleted from this instance, and fields of certain records may keep on changing. For example, in an accounts instance very regularly whenever a transaction is done, the account balance will get changed; maybe at a certain time the bank might decide to change the interest rate for a certain type of

account then the interest rate field will get changed, new customers may come into the customer instance and so on. So, instances keep on regularly being updated manipulated added deleted updated, but the schema remains unchanged.

So, change of the schema is very rare in a database and needs to be done only when the database is designed or when it is being upgraded. Because once you change the schema, it changes the way you look at the whole world, you look at the whole database scenario. So, if you are changing the schema at a logical level, then naturally the your view will also get affected, because you are using these schemas to decide how you would like to present a transaction application to the user or for a balance check application to the user and so on.

(Refer Slide Time: 12:22)

The slide has a title 'Schemas and Instances' in red. To the left is a small logo of a sailboat on water. On the right is a list of bullet points under the heading 'Physical Data Independence'. At the bottom, there is footer text: 'SWAYAM-NPTEL-ANOC-MOOCs Initiator: Prof. P. P. Desai, IIT Kharagpur', 'Database System Concepts - 8th Edition', '02.11', and 'Übersicht, Korth und Sudarshan'.

- **Physical Data Independence** – the ability to modify the physical schema without changing the logical schema
 - Analogous to independence of 'Interface' and 'Implementation' in Object-Oriented Systems
 - Applications depend on the logical schema
 - In general, the interfaces between the various levels and components should be well defined so that changes in some parts do not seriously influence others.

But, of course, what do we would want is between the physical schema and the logical schema we normally would want certain independence. What it means is the logical schema is what you need to deal with, because it is linked with the view that you have at a higher level of abstraction. On the other end, the logical schema is based on the physical schema; physical schema is how you are actually organizing the information in terms of the binary files the database files.

Now, certainly you want that logical schema not to change because if it changes then at a view level all your applications will have to change. But it is quite possible that your physical schema the way you have organizing files and so on might need a change,

because maybe it is just that you had designed the database for 10,000 records and you already have 9000 records and you would like to expand it to maybe 1,00,000 records.

And the physical this system needs to be different you may need to reorganize the files and so on, you may need to index it in a different way and all this, but you would like to do that change in the physical level without requiring any change at a logical schema. So, this property of a database schema is very required which we say is a physical data independence or the ability to change the physical schema without actually affecting the logical schema or the view level. So, that will be a critical factor that will have to keep in mind.

(Refer Slide Time: 14:16)

PPD

- Levels of Abstraction
- Schema & Instance
- Data Models
- DDL & DML
- SQL
- Database Design

DATA MODELS

SWATANTRUPTELL-NOCs Institute: Prof. P. P. Dr. S. K. Khatri - 2018

Database System Concepts - 8th Edition

02.12

©Silberschatz, Korth and Sudarshan

So, next is data models.

(Refer Slide Time: 14:20)

The slide has a header 'Data Models' in red. On the left, there is a small logo of a sailboat on water. The main content is a bulleted list of items related to data models:

- A collection of tools for describing
 - Data
 - Data relationships
 - Data semantics
 - Data constraints
- **Relational model** (we focus in this course)
- Entity-Relationship data model (mainly for database design)
- Object-based data models (Object-oriented and Object-relational)
- Semi-structured data model (XML)
- Other older models:
 - Network model
 - Hierarchical model

At the bottom, there is footer text: 'SYNTHOME: INFTEL-NOC: IIMODC's Infrastructure', 'Prof. P.P. Doshi, IIT Kharagpur', 'Database System Concepts - 8th Edition', '02.13', and '©Giberschattz, Korth and Sudarshan'.

Data models are a collection of tools that describe the data because we are talking about a database system. So, certainly the main focus here is to be able to model that it to be able to represent the data, so that talks about relationships between data, it talks about the meaning of the data the data semantics, it talks about data constraints. For example, it is an account balance, we just refer to account balance in the account schema in the instance, now, the question is will can the account balance be negative, the answer is yes or no.

If the bank mandates that I can only withdraw as much money up to which I have deposited, then account balance cannot be negative, but if the bank is giving me the facility to overdraft then I may be able to draw more money than I have actually deposited. So, my account balance might note negative. In some banks it could be that the bank says that well there is a minimum balance. So, minimum balance is 5,000. So, which means that my account balance cannot go below five thousand rupees the bank will not allow those transactions. So, these are examples of different constraints that might exist in the real world, and therefore, will be required in terms of the data model representation.

So, there are several data models that exist today. The most widely used the most popular and most powerful in terms of a certain section of database applications which we are commonly interested in is a relational model and that is what we focus in this course. We

will have a major discourse in terms of the relational model and lot of things will be developed based on that. But it is not easy to directly design a database in terms of the relational model, because you first need to understand the real world in which the design is happening.

So, we normally start with a less formal model known as the entity-relationship data model or an ER model, ER diagram, these are commonly called. So, if you recall your knowledge about object oriented systems, and if you happen to know UML, you already know about ER models and corresponding class diagrams that result. So, we will talk briefly about your model and show you how to do modelling on the real world in terms of the ER diagrams. But, then they are not actually models which the database systems directly used they are subsequently converted to some relational or other model and which the database systems will use.

Next is a object-based data models. You all would be knowing familiar with the fact that objects give a better power to represent the system which objects are not like simple strings or numbers or characters like that, they are encapsulated concept of an entity which can be manipulated in a certain way. So, like in the real world, you have several objects, it would have been nice to have similar objects in the databases. So, quite a few database system have been designed used which are object-based database systems. So, there are models for those. However, we will do little of that in this course, because it is little bit advanced in notion.

The other model, which has become very popular is called the semi-structured data model. It is primarily in terms of XML. I am sure all of you would be familiar with the basics of XML, which is extensible mark up language in which you can create use tags and use different mark ups to describe the data. You can say this is the field and this is the value kind of. And this is become a very nice way to represent the data. And XML or the semi-structured models are particularly useful today to exchange data between different systems.

I may be using a my SQL kind of database system, my friend may be using an oracle system, and we need to exchange data tables between these two, these two systems will represent the data in the physical schema which are not may not be compatible. So, we can represent both of them in terms of XML models convert the data. So, I convert the

data into XML, give it to my friend and my friend can import from that XML into the database. So, it is a XML is a data model which is frequently used in terms of data exchanges.

Then there are several other models like the network model, hierarchical model which used to be very popular in the early days of database systems before relational model came into force. They still exist in terms of the some of the databases. And some of the newer emerging big data databases actually we have started using this old concept in a new way again. So, this is a overall set of data model.

(Refer Slide Time: 19:58)

Relational Model

- All the data is stored in various tables
- Example of tabular data in the relational model

ID	name	dept_name	salary
2222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58853	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

So, here I am just showing an example of a relational model data which is simply looks like a table. So, you can see that on top row in the blue are the names of the different fields which describe the schema. It says that it has an ID, it has a name, it has a departments name, it has salary. They are trying to describe a particular instructor, and then a whole lot of records rows in that table, which are every row individually is a particular data entry or a record. So, columns are attributes, and rows are records that the relational model described.

(Refer Slide Time: 20:40)

The slide features a title 'A Sample Relational Database' at the top right. On the left, there is a small sailboat icon and a vertical watermark reading 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Doshi, IIT Kharagpur - June-Aug., 2015'. In the center, there are two tables. Table (a) 'The instructor table' has columns ID, name, dept_name, and salary, with 14 rows of data. Table (b) 'The department table' has columns dept_name, building, and budget, with 7 rows of data. At the bottom, there is a navigation bar with icons and the text 'Database System Concepts - 8th Edition'.

(a) The instructor table

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Sald	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58883	Califeri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(b) The department table

dept_name	building	budget
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

Some more of that the instructor table along with the department table. So, the table below describes details of a department, so that has its own schema and the individual records. We have of course seen similar instances already in terms of the customer and accounts instance that we have just discussed a couple of while ago.

(Refer Slide Time: 21:03)

The slide features a title 'DDL AND DML' at the top right. On the left, there is a small sailboat icon and a vertical watermark reading 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Doshi, IIT Kharagpur - June-Aug., 2015'. In the center, there is a bulleted list of database concepts: Levels of Abstraction, Schema & Instance, Data Models, DDL & DML, SQL, and Database Design. At the bottom, there is a navigation bar with icons and the text 'Database System Concepts - 8th Edition'.

- Levels of Abstraction
- Schema & Instance
- Data Models
- DDL & DML
- SQL
- Database Design

Let me introduce these two terms DDL and DML.

(Refer Slide Time: 21:10)

Data Definition Language (DDL)

- Specification notation for defining the database schema
 - Example: `create table instructor (`
 `ID char(5),`
 `name varchar(20),`
 `dept_name varchar(20),`
 `salary numeric(8,2))`
- DDL compiler generates a set of table templates stored in a **data dictionary**
- Data dictionary contains metadata (i.e., data about data)
 - Database schema
 - Integrity constraints
 - ↳ Primary key (ID uniquely identifies instructors)
 - Authorization
 - ↳ Who can access what

Syntax: INTEL-DOC: Books & Instruments Prof. P. Dr. S. K. Khatri - 2018
Database System Concepts - 8th Edition

02.17 ©Silberschatz, Korth and Sudarshan

DDL talks about data definition language. So, what the concept wise what we are trying to say is certainly we have a schema and we have instance. So, we need certain language constructs to be able to define a schema and certain other language constructs to be able to manipulate the data in that schema or they are basically manipulate the instances. So, DDL is the language or part of the language which is used to define and manipulate the schema of a database that is why schema is a way to define a database. So, it is called a data definition language.

So, you can define that I will am going to have a table called `instructor` which will have four different fields, each having certain types of data. So, it says that the `ID` will be a five character data; the `name` would we would have a variable length, because you cannot say that the `name` will be of a fixed length, but it will be a variable length that is what `varchar` is, but the length will not exceed 20. And similarly, `salary` will be a numeric data with up going up to 8 figures, and having a decimal part having two parts.

So, this way of defining the schema in terms of the different attributes and their types or columns in the table or trying to define the structure of that table is a main issue of the data definition language. So, the data definition language compiler who generates a set of tables in the data dictionary, where the data dictionary basically contains metadata as I said the schema is nothing but a metadata about the database tables. So, which will have the database schema, it will have different integrity constraints, it could say that well the

account balance cannot be negative or account balance cannot be less than the minimum balance. So, these are different integrity constraints. It could say that this is the primary key, we will talk more in more depth in terms of what is key. And it could also specify the authorization as to who is allowed to access which part of the data and so on, so that is these all are part of the schema definition and forms the DDL of the language.

(Refer Slide Time: 23:36)

The slide has a title 'Data Manipulation Language (DML)' in red at the top right. To the left of the title is a small logo of a sailboat on water. The main content is a bulleted list under the heading '• Language for accessing and manipulating the data organized by the appropriate data model'. This list includes: '• DML also known as query language', '• Two classes of languages', '• Pure – used for proving properties about computational power and for optimization' (with sub-points: '» Relational Algebra (we focus in this course)', '» Tuple relational calculus', '» Domain relational calculus'), and '• Commercial – used in commercial systems' (with sub-point: '» SQL is the most widely used commercial language'). At the bottom left is a video thumbnail showing a person speaking. The bottom right shows a navigation bar with icons for back, forward, search, and other presentation controls. The footer contains the text 'SWAYAM MOOCs Instructional Prod. P. P. Desai, ST Khemchand Patel', 'Database System Concepts - 3rd Edition', '02.18', and '©Baburao, Karth and Sudarshan'.

In contrast, the data manipulation language is a language for accessing and manipulating the data organized. So, it is for access, update, addition of new records, deletion of existing records and so on. And very commonly we will refer to the data manipulation language as a query language, because this is what you want to know what exists in the database. So, the query language will be designed, they are designed primarily in one of the two ways. One group of languages is known as a pure language, they are more mathematical in nature. They have a formal basis that can you can prove that whatever do you do in these languages are correct, and will give you the correct result.

So, they are different languages based on the relational model, they are called relational algebra, tuple relational calculus, domain relational calculus and so on. Of these three, we will in this course deal only with relational algebra. There are mathematical proof which show that whatever you can do in relational algebra you can do it in tuple relational calculus and vice versa. Similarly, whatever you can do in relational algebra, you can do in domain relational calculations and so on. In one sense that these languages

are equally powerful; the same thing can be done in any one of them, but we will just take the simplest of them and study here in terms of the relational algebra, but these are more mathematical representations are not easy to write as a program. So, normally we will use certain commercial query language which is called SQL for most of our applications and we will do the coding in that.

(Refer Slide Time: 25:19)

The screenshot shows a presentation slide with a title 'SQL' in red at the top right. To the left of the title is a small icon of a sailboat on water. Below the title is a bulleted list of facts about SQL:

- The most widely used commercial language
- SQL is NOT a Turing machine equivalent language
- To be able to compute complex functions SQL is usually embedded in some higher-level language
- Application programs generally access databases through one of
 - Language extensions to allow embedded SQL
 - Application program interface (e.g., ODBC/JDBC) which allow SQL queries to be sent to a database

On the left side of the slide, there is vertical text that reads: 'SWAYAM-NIPTEL-MOOCs Initiator: Prof. P. P. Doshi, IIT Kharagpur - June-Aug. 2018'. At the bottom left is a video player showing a person speaking, with the text 'Database System Concepts - 8th Edition' above it. At the bottom right is a navigation bar with icons for back, forward, search, and other presentation controls.

So, SQL which is a most widely used commercial language and mind you this is not a Turing equivalent language which means that everything that can be that need to be computed cannot be computed in SQL, there are certain computations which SQL cannot do. It is a limitation; it is a restricted language. So, often SQL is used in conjunction with some common high-level programming language like C or C++ and so on. So, whatever is there can be done in SQL in terms of data manipulation will be done in terms of the relational model, but there could be additional logic that needs to be built in, in terms of the high level language. So, application programs are typically written through them. So, we can do this through a process of embedding that is put in SQL as a part of a C program or use certain libraries which can actually take a query from C, and fire it on the SQL database.

(Refer Slide Time: 26:26)

The slide features a title 'DATABASE DESIGN' in large red capital letters. To the right of the title is a bulleted list of topics: '• Levels of Abstraction', '• Schema & Instance', '• Data Models', '• DDL & DML', '• SQL', and '• Database Design'. In the top right corner, the text 'PPD' is visible. On the left side, there is a vertical column of text: 'SYNTHAINE, INFLTEL-NOC, MOOCs Institute', 'Prof. P.P. Drs. IIT Kharagpur - June-Apr. - 2018', and 'Database System Concepts - 8th Edition'. At the bottom left is a small video thumbnail showing a man speaking. The bottom right contains a navigation bar with icons and the text '©Güberachitz, Korth and Sudarshan'.

So, we will see how to do this in the course of time.aspect.

(Refer Slide Time: 26:32)

The slide has a title 'Database Design' in red. Below the title is a definition: 'The process of designing the general structure of the database:'. A bulleted list follows, divided into two main categories: '• Logical Design – Deciding on the database schema.' and '• Physical Design – Deciding on the physical layout of the database.' The first category includes sub-points: 'Business decision' (with 'What attributes should we record in the database?'), 'Computer Science decision' (with 'What relation schemas should we have and how should the attributes be distributed among the various relation schemas?'). The second category includes a single point: 'Physical Design – Deciding on the physical layout of the database'. The slide's layout is identical to the previous one, with the same sidebar text and footer elements.

Coming to the database design this is a process through which the databases need to be designed. And certainly the first part of the design is the logical design where you want you need to identify what are the schemas and you know what are the constraints that apply, what is authorization required. And first set of decisions those are related to the business as we say. Business means it is basically comes from the domain. So, it is if I am doing a university database, the business decisions will come in terms of you know I

have courses, students, instructors, and the instructor teach courses, can an instructor teach multiple courses, can a course be taught by multiple instructors these kind of business decisions are critical for the database design.

And then there is a whole set of computer science decision or the data based decisions to decide as to if this is the kind of business information that you want to keep in the database, then what is the kind of relation, what is the kind of schemas that we should use what are should be the attributes, which attribute should be of what type what should be strain, what should be numbers and so on. So, these are formed the basis of the physical logical design. And of course, we then need a physical design which decides on the physical layout of the data, what are the different database files, how they should be indexed and so on.

(Refer Slide Time: 27:53)

Database Design (Cont.)

- Is there any problem with this relation?

ID	name	salary	dept_name	building	budget
22222	Einstein	95000	Physics	Watson	70000
12121	Wu	90000	Finance	Painter	120000
32343	El Said	60000	History	Painter	50000
45565	Katz	75000	Comp. Sci.	Taylor	100000
98345	Kim	80000	Elec. Eng.	Taylor	85000
76766	Crick	72000	Biology	Watson	90000
10101	Srinivasan	65000	Comp. Sci.	Taylor	100000
58583	Califieri	62000	History	Painter	50000
83821	Brandt	92000	Comp. Sci	Taylor	100000
15151	Mozart	40000	Music	Packard	80000
33456	Gold	87000	Physics	Watson	70000
76543	Singh	80000	Finance	Painter	120000

Database System Concepts - 8th Edition

So, here we are showing an example table. So, it has multiple fields. It shows the instructors expanded form of the instructor table you saw earlier. It is expanded with the departments name and the building in which it is housed. So, if you look carefully that this certainly comes from the business decision that you need to know the department to which an instructor belongs and certainly you need to know the building in which that department exists. So, knowing the department of the instructor and the building of where that department is are critical, but the question is this a good design, is this so we

will discuss as to when why this is a good, this may not be a good design to represent the data.

(Refer Slide Time: 28:41)

The slide is titled "Module Summary" in red at the top right. On the left, there is a small sailboat icon and some vertical text on the left edge. The main content area contains a bulleted list of learning objectives:

- Familiarized with the basic notions and terminology of database management systems
- Introduced the role of data models and languages
- Introduced the approaches to database design

At the bottom left, there is a small video thumbnail showing a man speaking. The bottom right corner shows a navigation bar with icons and the number "02.24".

So, in this module, we have taken you through the basic notions and terminology of database management systems, highlighting primarily the levels of abstraction, the schema an instance, the basic data models the languages that you need DDL, DML and the commercial SQL language. And we have also tried to give you a glimpse of the approach that is required in terms of the database design. We will elaborate on this more in the second part of our introduction to DBMS which will be taken up in module 3.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science & Engineering
Indian Institute of Technology, Kharagpur

Lecture - 03
Introduction to DBMS/2

Welcome to module 3 of Database Management Systems course. We started discussions introducing the database management systems in module 2. This is the second and concluding part of that discussion.

(Refer Slide Time: 00:37)

Module Recap

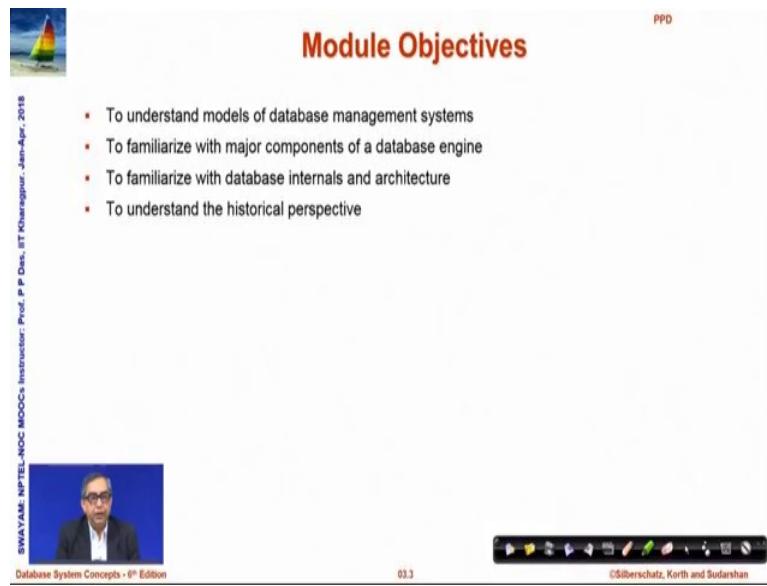
- Levels of Abstraction
- Schema & Instance
- Data Models
 - Relational Databases
- DDL & DML
- SQL
- Database Design

SWAYAM: NPTEL-NOC
Prof. P P Das, IIT Kharagpur - Instructor
Database System Concepts - 8th Edition

03.2 ©Silberschatz, Korth and Sudarshan

So, this is what, these are the aspects that we are discussed earlier starting from level of abstraction to the outline of database design.

(Refer Slide Time: 00:47)



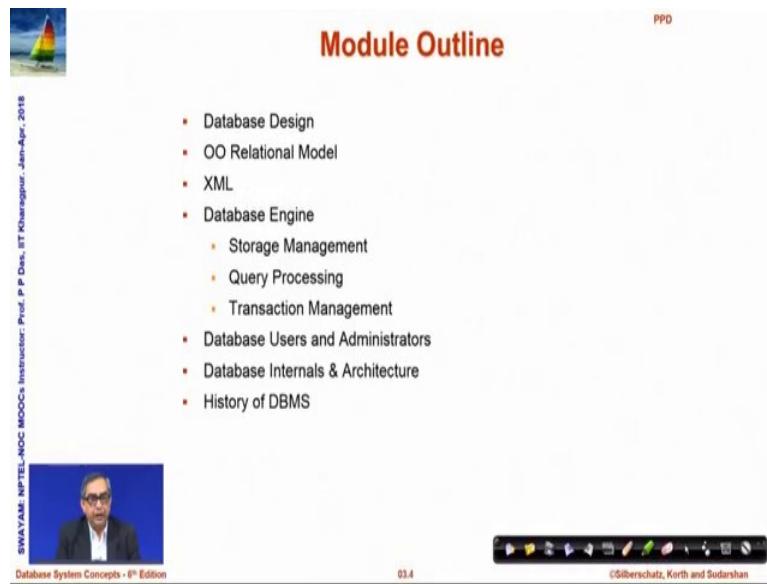
This slide is titled "Module Objectives" in red at the top right. It features a small sailboat icon in the top left corner. On the far right, there is some very small, illegible text. The main content is a bulleted list of objectives:

- To understand models of database management systems
- To familiarize with major components of a database engine
- To familiarize with database internals and architecture
- To understand the historical perspective

At the bottom, there is a video frame showing a person speaking, the text "Database System Concepts - 8th Edition", the number "03.3", and the copyright notice "©Silberschatz, Korth and Sudarshan".

In the current module, we would like to understand the modules of database management systems little bit more and we will try to familiarize with the concept of major components of database engine, will elaborate on those and will familiarize. So, with the basic architecture of a database management system, some of the internal components and we will present a brief historical perspective of the DBMS.

(Refer Slide Time: 01:21)



This slide is titled "Module Outline" in red at the top right. It features a small sailboat icon in the top left corner. On the far right, there is some very small, illegible text. The main content is a bulleted list of outline items:

- Database Design
- OO Relational Model
- XML
- Database Engine
 - Storage Management
 - Query Processing
 - Transaction Management
- Database Users and Administrators
- Database Internals & Architecture
- History of DBMS

At the bottom, there is a video frame showing a person speaking, the text "Database System Concepts - 8th Edition", the number "03.4", and the copyright notice "©Silberschatz, Korth and Sudarshan".

So, this is the outline that we will follow.

(Refer Slide Time: 01:25)

The slide features a small sailboat icon in the top left corner. In the top right, the text 'PPD' is visible. A vertical column of text on the left side reads: 'SWAYAM: NPTEL-NOOCs MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Apri- 2018'. On the right, a bulleted list of topics is shown:

- Database Design
- OO Relational Model
- XML
- Database Engine
- Database Users and Administrators
- Database Internals & Architecture
- History of DBMS

DATABASE DESIGN

Database System Concepts - 8th Edition

03.5

©Silberschatz, Korth and Sudarshan

(Refer Slide Time: 01:34)

The slide includes a sailboat icon in the top left and a video player interface in the bottom right showing a man speaking. The main title 'Database Design' is centered at the top. Below it, a definition is provided: 'The process of designing the general structure of the database:'. The content is organized into two main sections:

- **Logical Design**
 - Deciding on the database schema. Database design requires that we find a "good" collection of relation schemas.
 - Business decision
 - What attributes should we record in the database?
 - Computer Science decision
 - What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
- **Physical Design**
 - Deciding on the physical layout of the database

Database System Concepts - 8th Edition

03.6

©Silberschatz, Korth and Sudarshan

(Refer Slide Time: 01:46)

So, we have already discussed about the database design, I would like to raise a few further issues about that. So, we have seen that, there is a logical design which is driven by the business decisions and refined by the computer science decisions, there is a physical design as well; and based on that we had presented this particular table asking, whether, this database is a good design or not.

So, let us have a little look into this for example, we have introduced the department name and the building in which the department is housed. So, if we look at there are multiple instructors say, let us say Professor Einstein, who teaches in the Physics department, that is housed in the Watson building and if we look through there is a Professor Gold, who also teaches in the Physics department and naturally that is housed in the Watson building.

Now, the question is; so, Physics department, if it is housed in the Watson building then, all the instructors in this table, all the instructors who are part of the Physics department, would have their department housed in the Watson building. So, there is a certain issue of redundancy in these two, that is, the same information is given more than once, which is not a very desirable thing.

The consequence of this could be suppose, tomorrow the university decides to move this Physics department from Watson to the Taylor building. This will mean that once this is moved, then this Watson will have to be changed to Taylor, also this Watson will also have to be changed to Taylor. All instances of Watson that corresponded to the Physics

department in this table, will have to be changed to Taylor, and that is not a good scenario. So, it is not only that we have redundancy, we have potential for anomaly; that is program the application programmer might forget to update the building say, for this entry then, we will be in an inconsistent database. So, to put it in simple terms that this is not a good design and there are several issues to consider, in terms of whether some design is good or some design needs refinement.

(Refer Slide Time: 04:22)

The slide has a title 'Design Approaches' in red at the top right. To the left of the title is a small icon of a sailboat on water. On the far left, there is vertical text: 'SWAYAM: NPTEL-MOOCs Instructor: Prof. P. P. Doshi, IIT Kharagpur - Jan-Apr - 2015'. At the bottom, there is footer text: 'Database System Concepts - 8th Edition', '03.8', and '©Silberschatz, Korth and Sudarshan'.

Design Approaches

- Need to come up with a methodology to ensure that each of the relations in the database is "good"
- Two ways of doing so:
 - Entity Relationship Model (Chapter 7)
 - Models an enterprise as a collection of *entities* and *relationships*
 - Represented diagrammatically by an *entity-relationship diagram*:
 - Normalization Theory (Chapter 8)
 - Formalize what designs are bad, and test for them

So, we need to come up with a methodology to ensure that, each of the relations in the database is good. So, we primarily follow two approaches in doing this. One is, using the entity relationship model, which models the enterprise as a collection of entities or concepts or if you are familiar with the object orientation classes and the relationships that hold between these entities.

So, in a university database the entities are students, courses, teachers and the relationships are a teacher teaches a set of courses, a student attends a set of courses and so on, the teachers supervise a set of students for projects and so on and then represent them diagrammatically in terms of a ER diagram entity relationship diagram and once that has been done then, we try to follow a certain normalization theory.

This normalization theory tries to capture that what are the properties that must hold in this database design, that must be satisfied on this database design, in terms of what is known as database dependencies, there are, varied forms of dependencies functional

dependencies, multi value dependencies, joint dependencies and so on and try to formalize and evaluate whether a design is good or it is bad, test them for quality and then normalize to make them better; make them the best possible that can happen.

So, this is something that is starting from the entity relationship model, which captures the real world to the actual database schema, there is a process of representation and then capturing of ground truths, that should hold in the database system and then normalize the database is a basic requirement of the design approach.

(Refer Slide Time: 06:34)

The slide features a small sailboat icon in the top left corner. In the top right corner, the text 'PPD' is visible above a list of topics: Database Design, OO Relational Model, XML, Database Engine, Database Users and Administrators, Database Internals & Architecture, and History of DBMS. The main title 'OBJECT-RELATIONAL DATA MODELS' is centered in large red capital letters. At the bottom, there is footer text: 'SWAYAM: NPTEL-NOOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Apr- 2015', 'Database System Concepts - 8th Edition', '03.9', and '©Silberschatz, Korth and Sudarshan'.

(Refer Slide Time: 06:39)

The slide features a small sailboat icon in the top left corner. The main title 'Object-Relational Data Models' is centered in large red capital letters. Below the title is a bulleted list of points: Relational model: flat, "atomic" values; Object Relational Data Models; Extend the relational data model by including object orientation and constructs to deal with added data types; Allow attributes of tuples to have complex types, including non-atomic values such as nested relations; Preserve relational foundations, in particular the declarative access to data, while extending modeling power; Provide upward compatibility with existing relational languages. At the bottom, there is footer text: 'SWAYAM: NPTEL-NOOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Apr- 2015', 'Database System Concepts - 8th Edition', '03.10', and '©Silberschatz, Korth and Sudarshan'.

We have talked about object relational data models for a few more points about that, that in a relational model everything is flat, every value is atomic, in the sense that everything if you, look back and think in terms of C then, every field is a value which can be a simple, you know, built in type like integer, like fixed length string, variable length string, a floating point number like that, but I cannot have a composite you know, object kind of fields.

But in a relational data model we extend in the object relational data model, we extend the relational model by including the object orientation and the constituent constructs to deal with added data types, higher data types where attributes are allowed to have complex types, non atomic values that may allow things like nested relation; that is a value could itself be a relation, could itself be a table but, we try to preserve the relational foundation and we will see what those foundations mean and provide upward compatibility to existing relational databases. So, this is what the basic concept of object relational data models are and as I said that, we will just glimpse through it, but this is not the primary objective that we will try to cover.

(Refer Slide Time: 08:04)

The slide has a white background with a small sailboat icon in the top-left corner. On the right side, there is a vertical list of topics:

- Database Design
- OO Relational Model
- XML
- Database Engine
- Database Users and Administrators
- Database Internals & Architecture
- History of DBMS

At the bottom, there is a red banner with the text "XML: EXTENSIBLE MARKUP LANGUAGE". The footer contains the text "SWAYAM: NPTEL-NOC's Instructional Prof. P. P. Chakraborty - Jan-Apr., 2018", "Database System Concepts - 8th Edition", "03.11", and "©Silberschatz, Korth and Sudarshan".

(Refer Slide Time: 08:07)

The slide has a header 'XML: Extensible Markup Language' with a sailboat icon. The main content is a bulleted list of facts about XML. At the bottom left is a video player showing a speaker, and at the bottom right is a navigation bar.

• Defined by the WWW Consortium (W3C)
• Originally intended as a document markup language not a database language
• The ability to specify new tags, and to create nested tag structures made XML a great way to exchange **data**, not just documents
• XML has become the basis for all new generation data interchange formats
• A wide variety of tools is available for parsing, browsing and querying XML documents/data

In contrast, the XML extensible markup language was defined by W3C, and it was originally intended for marking up document languages. It was not designed as a database language, it was designed for marking up. So, it is kind of saying that this particular element should be put in capital, this should be in blue colour, this means a verb, this means a paragraph, there should be a page break here, those kind of markups, But subsequently, it turned out that the way XML deals with different components in terms of tags, and the ability to create nested tags, makes a great language for exchange of data, As I explained in the last module also.

So, has become the basis for all kinds of new generation data interchange format. So, as I explained, that any database should be able to convert the data instances of the tables in terms of corresponding XML format and then you take it to some other database, in which you are intending to interchange the data and that target database should be able to import from that XML structure, and it has become widely available that you have different tools for parsing, browsing and querying XML content document data and so on. So, if you are familiar with C programming, I hope so you are! You can look up certain XML parsing and try out, there are great tools to learn.

(Refer Slide Time: 09:53)

The slide features a title 'DATABASE ENGINE' in large red capital letters at the top center. To the left of the title is a small video window showing a man speaking. On the right side, there is a vertical list of topics under the heading 'PPD': Database Design, OO Relational Model, XML, Database Engine, Database Users and Administrators, Database Internals & Architecture, and History of DBMS. The bottom of the slide includes a navigation bar with icons and the text 'Database System Concepts - 8th Edition', '03.13', and '©Silberschatz, Korth and Sudarshan'.

Moving on, let us briefly look at what is the core of a database management system, the database engine.

(Refer Slide Time: 10:03)

The slide has a title 'Database Engine' in red at the top center. Below the title is a bulleted list of three components: Storage manager, Query processing, and Transaction manager. To the left of the list is a video window showing a man speaking. The bottom of the slide includes a navigation bar with icons and the text 'Database System Concepts - 8th Edition', '03.14', and '©Silberschatz, Korth and Sudarshan'.

The database engine, primarily contains 3 major components; the storage manager, the query processing engine, sub engine and the transaction manager.

(Refer Slide Time: 10:14)

The slide has a header 'Storage Management' with a sailboat icon. On the left, there is vertical text: 'SWAYAM, NPTEL-NOOC, IIT-Kharagpur', 'Prof. P. P. Desai', 'Jain-Agarwal', and '2018'. The main content is a bulleted list:

- **Storage manager** is a program module that provides the interface between the low-level data stored in the database and the application programs and queries submitted to the system.
- The storage manager is responsible to the following tasks:
 - Interaction with the OS file manager
 - Efficient storing, retrieving and updating of data
- Issues:
 - Storage access
 - File organization
 - Indexing and hashing

Below the list is a video player showing a person speaking. The video player has a progress bar at '03.15' and a copyright notice '©Siberschatz, Korth and Sudarshan'.

The storage manager, is a module or a collection of modules in a database management system, that provide the interface between the low level data and the application program. So, we have looked at the storage manager is the one, which is a bridge between the physical level of abstraction and the logical level of abstraction, then finally, to the view level of abstraction. So, the storage manager has to deal with interactions with the operating system on which the DBMS is kept, the file manager of the operating system, it is responsible for efficient storage, retrieval update of the data, it is responsible to make sure that if there are certain problems in the file system, then the data is not corrupted and so on.

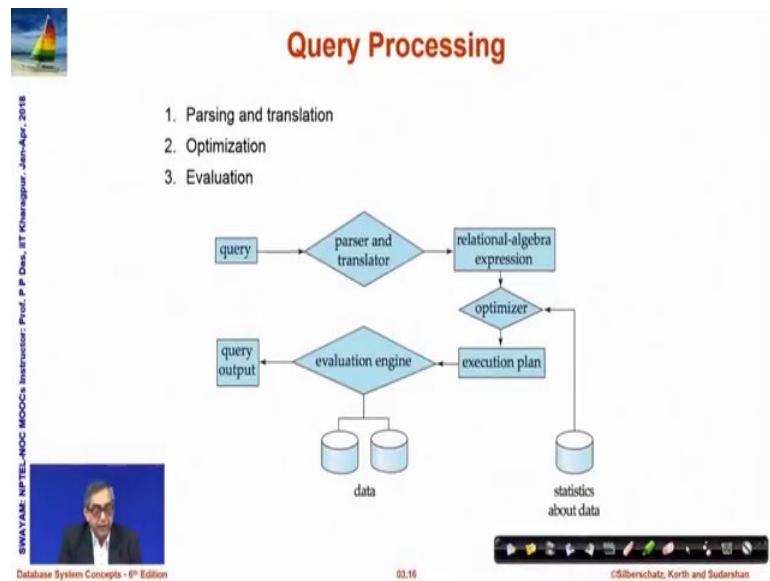
So, the issues certainly that involve are :- the access to the storage, the organization of the files and very importantly indexing and hashing and we will talk about the concept of indexing later in the course. It primarily says that, if I want to, for example, you can simply understand that if you have a large chunk of data that you want to organize for efficient search then, you can use the binary search tree in simple algorithm terms. The binary search tree needs to be organized in terms of one data component.

We say that, well there is one value based on which you can say that, comparison is done. So, that at every node if that value is smaller, you go to the left sub tree, if that value is larger, you go to the right sub tree and so on. So, if we want to organize the

records of a database system in terms of such a binary search tree, then the question certainly is, which field do I use for the search tree comparison.

Now, whatever field I used for search tree comparison, on that field the searching would be very efficient, but if I want to search on a value for a different field, the searching would not remain that efficient. So, indexing is a mechanism by which, you can actually create auxiliary search trees on multiple fields. So that, the search on multiple fields can be made efficient and we will talk about this later when, that particular module comes up, but the storage manager has to deal with such issues.

(Refer Slide Time: 12:43)



Moving on, the query processing is, if we have already talked about the language; the DDL, the DML, the query language. So, it is some kind of, like the C program, it is some kind of a text based programming code. So, naturally that code needs to be parsed and translated, as we typically do in a C compiler. So, there needs to be a query compiler. So, it parses and analyses the code, but translated unlike the C program, which translates the C program into an intermediate code and then into the binary instructions of the machine, the assembly binary instructions of the machine.

The query translator, translates the query into relational algebra expressions. I said that, there are two kinds of languages :- the commercial query language and the pure language. So, it translated in terms of a program in the pure language, which could be a

relational algebra language and then it tries to optimize. So, that is a critical term to be noted that there is an optimizer.

So, this optimizer is a critical component, which tries to make sure that the query, when it is run on your data will run with the most you know, least amount of time in an effective manner. So, then an execution plan needs to be decided, we will be able to understand this when we go to the actual relational algebra execution plan, basically says that if, there are multiple operations in that query to be performed, then how those operations, in which order they should be performed and where should temporary tables be used, where they should be skipped and so on and then, once that has been done then it passes on to an evaluation engine which actually runs that query on the data that you have, the instances of the data that you have, and that brings out the resultant query output which is another table of results that we get. So, this query processing is a core part of a database engine, which actually allows us to write text based queries and relative data efficiently, change-update data efficiently, insert data efficiently, and so on.

(Refer Slide Time: 15:12)

So, when we do this, we need to look at alternative ways of evaluating a query. There could be different ways to write the same thing, these are called equivalence expression, equivalent expressions and what are the good algorithms for doing each and every operation, there is a cost between good and bad way of evaluating. So, this has to be understood that, the same thing you can compute in a, you have seen this similar

concepts in normal programming language also, I mean we have seen for example, for sorting the several ways to sort and some are better some are not as efficient.

So, if the similar things in terms of a query need to be evaluated and the cost between good and bad ways need to be figured out. So, then we need to estimate the cost of every operation. It depends on the information of what has happened in the past, the statistical information and need to estimate those statistics for intermediate results. These are a couple of things that the query processing sub engine in a database will do.

(Refer Slide Time: 16:36)

The slide has a title 'Transaction Management' in red. To the left of the title is a small sailboat icon. On the right side of the slide, there is a vertical column of text that reads:
SWATAMI, INFLTEL-NOC, MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jain-Agarwal
Database System Concepts - 8th Edition
03.18 ©Silberschatz, Korth and Sudarshan

- What if the system fails?
- What if more than one user is concurrently updating the same data?
- A **transaction** is a collection of operations that performs a single logical function in a database application
- **Transaction-management component** ensures that the database remains in a consistent (correct) state despite system failures (e.g., power failures and operating system crashes) and transaction failures.
- **Concurrency-control manager** controls the interaction among the concurrent transactions, to ensure the consistency of the database.

At the bottom of the slide is a video player interface showing a thumbnail of a person speaking, the text 'Database System Concepts - 8th Edition', the time '03.18', and the copyright notice '©Silberschatz, Korth and Sudarshan'.

So, beyond the storage manager and the query processor we have a transaction management system, which is very critical and core of the database system. It primarily has to deal with two fundamental issues of a database. One, what if a system fails; see, database systems are unlike the programs that you have written so far. A program starts execution ends, the program always deals with transient data, the data did not exist before your program started, it ceases to exist after your program ends. So, a program; however, complicated; however, important has a limited lifetime.

A database in contrast, has a much longer lifetime which deals with persistent data, that is very important to understand ,that is the each application whether I am doing a bank fund transfer, whether I am making a credit card payment, to whether I am checking the balance or I am booking a railway ticket, whether, I am purchasing a book from Amazon, each one of the applications are like the normal program, it has a fixed lifetime.

If I start it, I do certain operations, I am done with it, but the data that is behind it the data of my accounts, my account balance, my transactions, my different bank charges, all that need to stay on and on and on and beyond every particular operation that I have done on the database. So, which means that if this database system fails, at some stage for some reason, then we have an enormous impact of that and that is not something that we can absorb that something that we can accept.

So, a database system has to come with the concept of recovery. It must be possible, if the system fails, it must be possible to recover it, to a certain earlier point where it is consistent. So, transaction management system is responsible to guarantee this kind of recover ability of databases. Then, the other question that we have discussed about earlier also, is a multiple users are you accessing the same database, the same set of data, at the same time. So, what how to make sure that more than one user can concurrently use an update without the data getting inconsistent, that is as I had mentioned, there is only one seat available, one berth available on a particular train, on a particular date and two users at the same time has initiated a booking.

It should not happen, that both of them get the booking. So, one should get, one should not get and that needs to be the complexity is high, for this kind of you know, decisions because in the databases, as applications are significantly distributed, Indian railways have no idea of who is going to do what booking, of which berth, from where at, which point of time. So, transaction management system is, as the name suggests defines something called, a transaction which always keeps the database consistent and operable. So, a transaction is a collection of operation, that performs a single logical function in a database application. This is very critical. It is a collection of operations and performs a single logical function.

So, it does not do anything and everything, it just does a single particular logical operation and that is what forms the transaction. So, a transaction management system is a component, that ensures that with all these transactions happening, hundreds and thousands of transactions happening every second in a database system, in a typical database system; the data should still remain consistent, in a consistent state, in spite of failures, in spite of concurrences, it must make sure that at no point of time it should happen that, an amount has been debited from an account and has not been credited to account or an amount has been credited to an account and has not been debited to the

corresponding account or the same seat same berth is booked by two persons at the same time and so on, so forth. So, this also includes concurrency control manager, which controls the interaction among different concurrent transactions, which ensures the consistency in the database and provides the total safety.

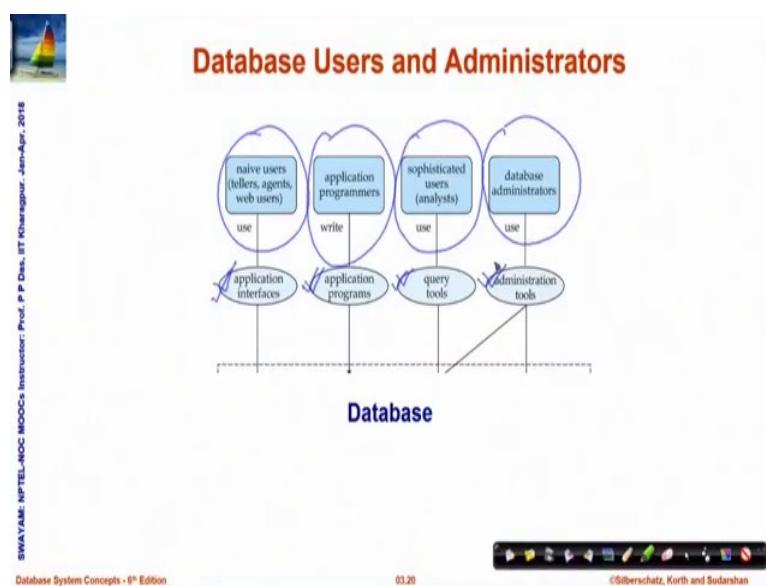
(Refer Slide Time: 21:32)



The slide features a small sailboat icon in the top left corner. On the right side, there is a vertical list of topics under the heading 'PPD': Database Design, OO Relational Model, XML, Database Engine, Database Users and Administrators, Database Internals & Architecture, and History of DBMS. The main title 'DATABASE USERS AND ADMINISTRATOR' is centered in large red capital letters. Below the title, there is a navigation bar with icons for back, forward, search, and other presentation controls. The footer contains the text 'Database System Concepts - 8th Edition', the date '03.19', and the copyright notice '©Silberschatz, Korth and Sudarshan'.

So, in total, we have seen the different components of the database engine comprising the storage manager, comprising the query processor, and the transaction manager. Now, we will just have a quick look in terms of who are the typical users of a database system.

(Refer Slide Time: 21:57)



So, if we see grossly, the users of a database system can be grouped into, I mean you can group it in multiple different ways, but this is a typical way to group that, you have the naive users, those like the secretarial staff, who sits at the teller of the bank. Now, that person does not know database management system, but that person just needs to know the particular application. He knows a few set of screens; graphical screens, what needs to be filled up, where which button needs to be clicked and so on and can use this database through an application interface.

So, this is a lowest level of user. Then, you have the set of application programmers, about whom I talked about in my course overview presentation, that application programming is a big chunk of you know, IT services that databases need, who actually write the application programs, while the naive user is similarly using it application programmers are responsible for writing coding this application program. So, they need to understand the database designs they need to understand how to write the query language, how to fit with the application data, input, output all the systems.

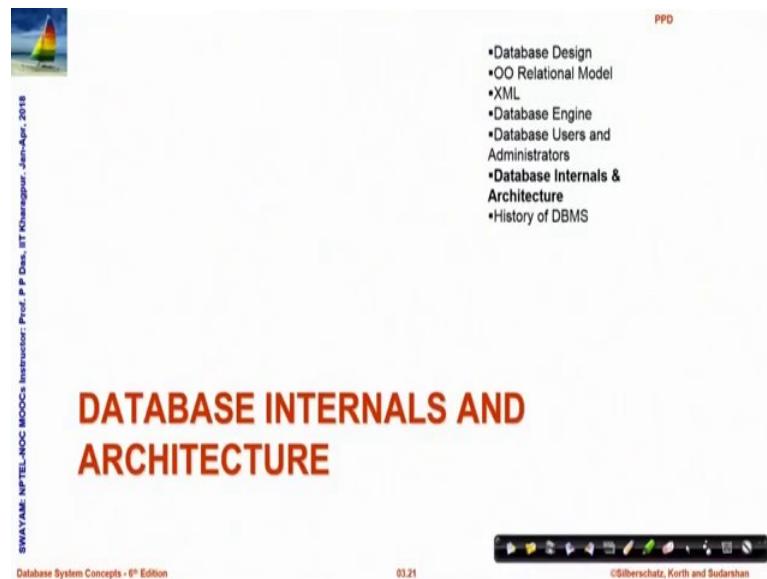
The next levels are the analysts, who are called the sophisticated users. So, they design different kinds of query tools, they are responsible for the design of the database, that is a schema the different constraints, the authorizations and so on and manages that over a period of time, when the application requirements change they might need to redesign the schema, migrate the data from an old schema to a new schema. So, analysts are higher level of programmers, they have far more solid understanding of the database management system to be able to design different kind of query tools that, the application programmer will use and at the end there are database administrators.

So, database administrators are people with specialized rights, who can do a lot of privileged operation on the database. For example, taking backups of databases, for example, creating different users. For example, if there has been a failure then how to do the failure recovery scripts, recovering the database and so on. So, they do all kinds of administration tasks, but not the usual day to day data maintenance and you know, query processing and so on.

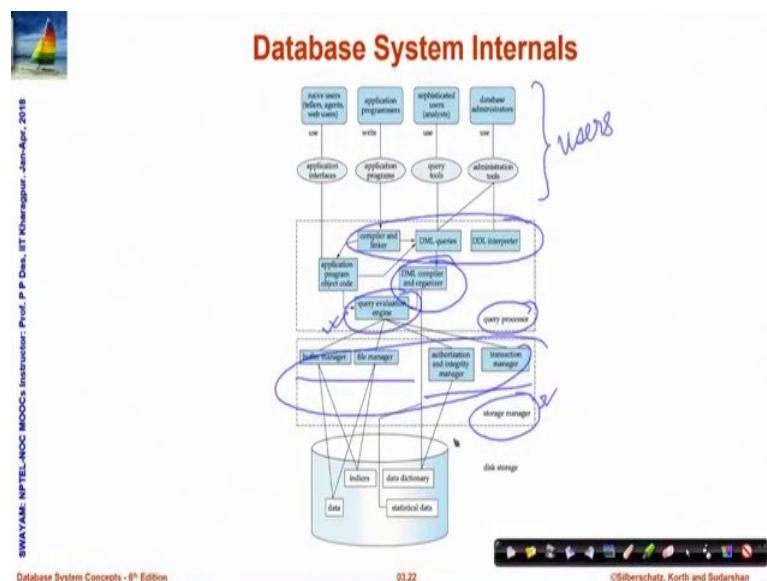
So, if you would like to know your positions then I would say that by this course, you are going to position yourself amongst the application programmers and the analysts and as I mentioned that the first half of the course is focused on application programming aspect

and the second half would be more focused on the analysis and some of the administrative will do little bit of administration, but not nearly serious administration tasks.

(Refer Slide Time: 25:33)



(Refer Slide Time: 25:35)

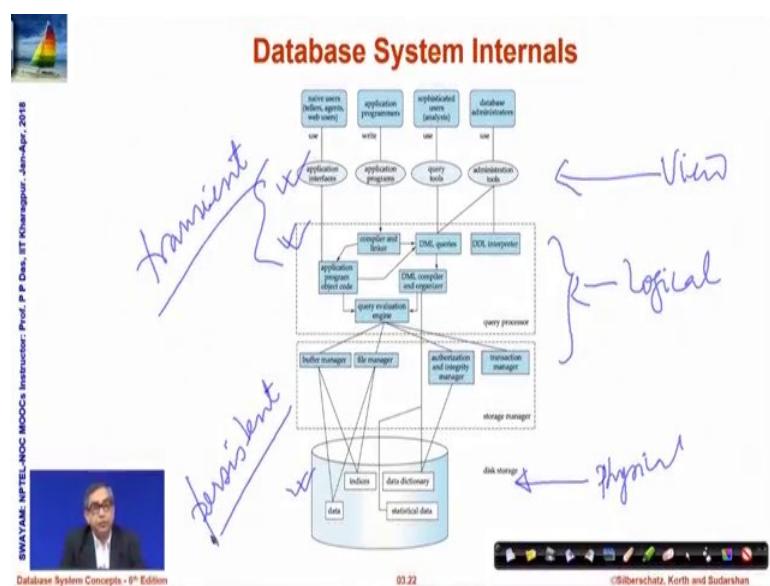


Now, we will take a quick look into the database internals and architecture. So, I will take you to this diagram. I am sorry this diagram is little small, in terms of the script size. So, please refer to the actual presentation. So, if you look at the top here is the users. So, it is trying to show what different users use. This is a query processor, that we have

known. So, the query processor gets a query and so that naturally, this query comes from the application program. So, the compiler link these are the basic processing, then the compiler organization, the evaluation engine which actually takes care of the processing of the whole query and then it goes to the storage manager which is now taking this whatever the evaluation engine needs to do, has to go through different modules in the storage manager, which we will talk about these modules, when we do have a discussion module on the storage management.

Later in this course, and we will then talk about what is a file manager and what is authorization and so on, but these are the sub components of the storage management needs to do and then the storage manager is the only one who deals with the actual data, the actual disk storage the different files and so on.

(Refer Slide Time: 27:05)



So, as you can see that, the whole system is kind of layered in terms of so, this is your basic physical layer that you have, and this is your final view layer that you have and in between this is a logical layer, that you deal with. So, you can see that the abstractions, as we had talked about are also mapped in terms of them, way the actual database system architecture is managed and finally, the data stays in the disk storage which ensures that whatever data I have is actually persistent in nature. It does not go away. Any data that stays within here or within the application interfaces transient.

So, these data are transient, they will disappear as soon as the application is over, but these data are persistent, they exist beyond this and architectures supports that whole gamut from of all applications or transiency of the applications based on the persistence of the data at the storage. So, that is a basic architectural view of a typical database systems. The actual architecture we form more complex, but we just want to take a schematic view. So, that we can understand it better.

(Refer Slide Time: 28:24)

The slide features a title 'History of Database Systems' in red, accompanied by a small sailboat icon. Below the title is a bulleted list of historical milestones. On the left edge of the slide, there is vertical text indicating the source: 'SWAYAM: NPTEL-MOOCs Instructor: Prof. P. P. Desai, Jai Narayan Kharatmalkar - Jan-Apr - 2018'. At the bottom, there is footer text: 'Database System Concepts - 8th Edition', '03.25', and '©Silberschatz, Korth and Sudarshan'.

- 1950s and early 1960s:
 - Data processing using magnetic tapes for storage
 - Tapes provided only sequential access
 - Punched cards for input
- Late 1960s and 1970s:
 - Hard disks allowed direct access to data
 - Network and hierarchical data models in widespread use
 - Ted Codd defines the relational data model
 - Would win the ACM Turing Award for this work
 - IBM Research begins System R prototype
 - UC Berkeley begins Ingres prototype
 - High-performance (for the era) transaction processing

So, the actual architecture may again vary, based on the computer system that you are using. It could be centralized, we will talk about some of these at later modules. Centralized in the sense that there could be one database server or you know, a group of database servers at the same physical location connected together to which, all applications, all users will connect to, it could be in terms of a client server model which is a very typical client server model, that the programming systems are. So, that typically for example, any of the net based internet based database applications we are looking at are necessarily client server in nature.

What you are doing in the browser is a client and there is a server back far back there the multiple tiers in between them. Databases could be single processes or for performance they could be in terms of multiprocessor parallel databases, the data sets themselves could be so large that, it may not be possible to search on them through a single

processor, in a reasonable time, they could be distributed the data itself could be distributed, tables could get so large that I may not be able to keep them at a single point.

So, these are different aspects of you know, complex real life database system that exist and we will deal with some of those aspects over the course of time, but you have to keep in mind that a final architecture of a database and its associated application will have some of these factors that you will need to know and maybe decide on in terms of history well and I will not go through each and every point here. This is more for completeness and to give you an essence of how things have been going. So, database system started in the 50's and early 60's, then major developments of these relational models and all that started happening in the 70's.

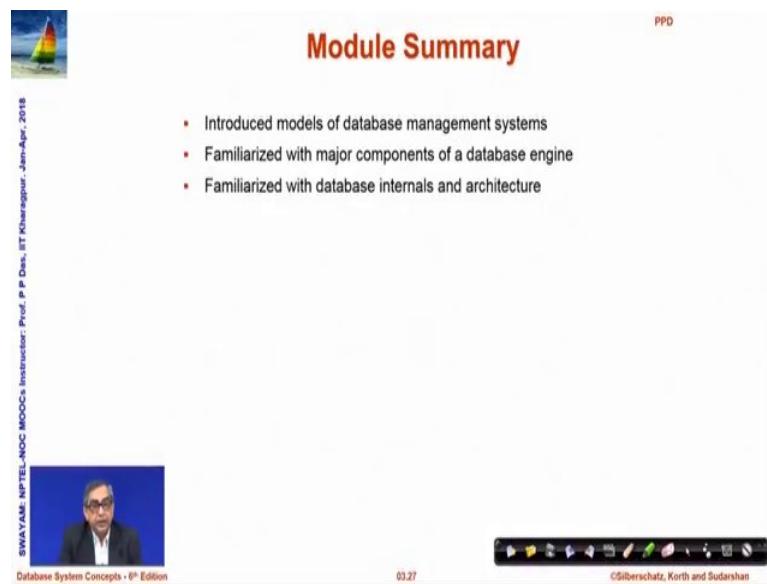
(Refer Slide Time: 30:36)

The slide is titled "History (cont.)" in red. It features a vertical timeline on the left with a sailboat icon at the top. The timeline is divided into four main periods: "1980s:", "1990s:", "Early 2000s:", and "Later 2000s:". Each period has a list of developments. At the bottom, there is a navigation bar with icons for back, forward, search, and other presentation controls.

- 1980s:
 - Research relational prototypes evolve into commercial systems
 - SQL becomes industrial standard
 - Parallel and distributed database systems
 - Object-oriented database systems
- 1990s:
 - Large decision support and data-mining applications
 - Large multi-terabyte data warehouses
 - Emergence of Web commerce
- Early 2000s:
 - XML and XQuery standards
 - Automated database administration
- Later 2000s:
 - Giant data storage systems
 - Google BigTable, Yahoo PNuts, Amazon, ..

And in 80's, really it proliferated large in terms of prototypes and commercial systems, parallel distributed systems, object based systems started happening in 80's 90's. It really exploded in terms of large decision support, data mining applications, you know applications widespread use of internet based data applications and systems emergence of Google and all that. From 20's, early in early 2,000 it has been XML and automated database administration and now what we are facing in at the big data, which are certainly aspects of other courses, but at the back of back of back at the last layer then often is a strong relational system that exists.

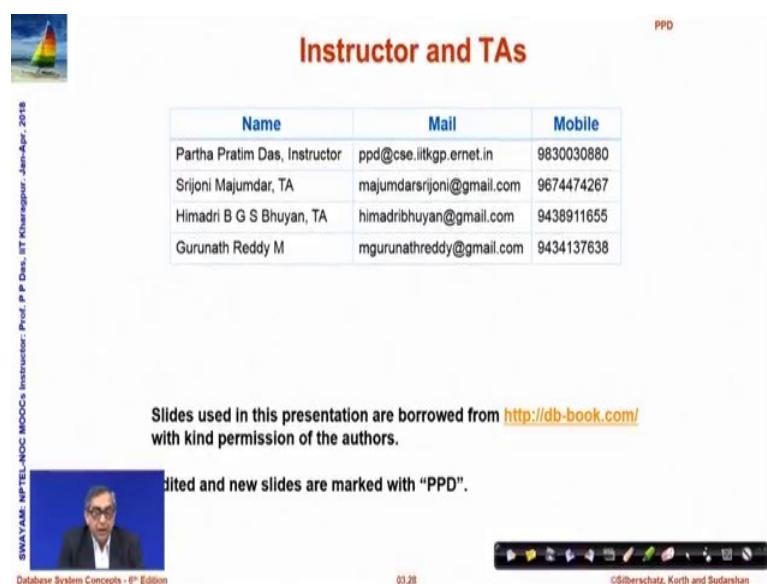
(Refer Slide Time: 31:28)



This slide is titled "Module Summary" in red at the top right. It features a small sailboat icon in the top left corner. On the left side, there is vertical text: "SWAYAM: NPTEL-NOOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apri- 2018". At the bottom left is a video thumbnail showing a man in a suit. The bottom center contains the text "Database System Concepts - 8th Edition". The bottom right shows a navigation bar with icons and the text "03:27" and "©Silberschatz, Korth and Sudarshan". A red "PPD" mark is in the top right corner.

- Introduced models of database management systems
- Familiarized with major components of a database engine
- Familiarized with database internals and architecture

(Refer Slide Time: 31:41)



This slide is titled "Instructor and TAs" in red at the top right. It features a small sailboat icon in the top left corner. On the left side, there is vertical text: "SWAYAM: NPTEL-NOOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apri- 2018". At the bottom left is a video thumbnail showing a man in a suit. The bottom center contains the text "Database System Concepts - 8th Edition". The bottom right shows a navigation bar with icons and the text "03:28" and "©Silberschatz, Korth and Sudarshan". A red "PPD" mark is in the top right corner.

Name	Mail	Mobile
Partha Pratim Das, Instructor	ppd@cse.iitkgp.ernet.in	9830030880
Srijoni Majumdar, TA	majumdarsrijoni@gmail.com	9674474267
Himadri B G S Bhuyan, TA	himadribhuyan@gmail.com	9438911655
Gurunath Reddy M	mgurunathreddy@gmail.com	9434137638

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with "PPD".

So, to summarize, in this module, we have introduced the models of database management system, the major components of a database engine and discussed about the internals and architecture, and this will conclude our discussion on the internals of database management systems, and next we will move on to exposing more on the relational model.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science & Engineering
Indian Institute of Technology, Kharagpur

Lecture - 04
Introduction to Relational Model/1

Welcome to module 4 of database management systems, in the last two modules, we have introduced the basic notions of DBMS in this module and the next; we make an introduction to the relational model, which we said is a major data model that we are going to use.

(Refer Slide Time: 00:46)

Module Recap

- Database Design
- OO Relational Model
- XML
- Database Engine
 - Storage Management
 - Query Processing
 - Transaction Management
- Database Users and Administrators
- Database Internals & Architecture
- History of DBMS

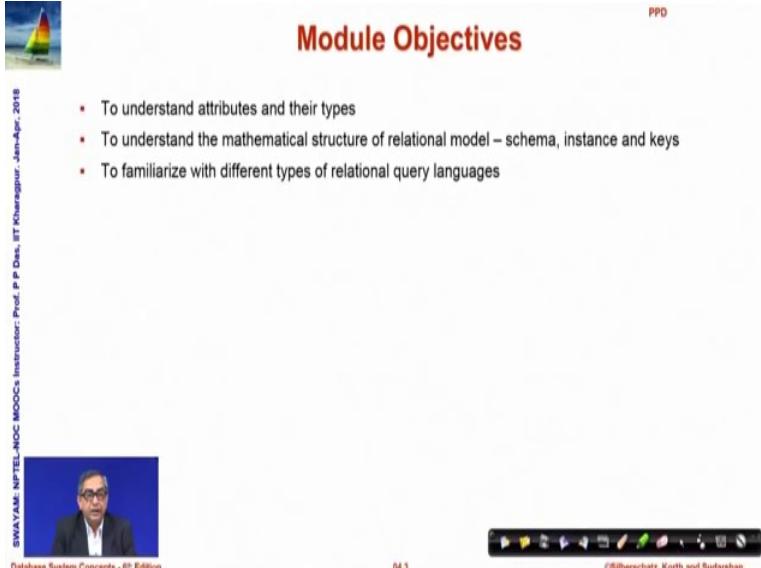
SWAYAM: IITTELMOOCs Instructor: Prof. P P Das, IIT Kharagpur - Jan-Apr. 2018

Database System Concepts - 8th Edition

04.2 ©Silberschatz, Korth and Sudarshan

So, this is what we did in the last module.

(Refer Slide Time: 00:52)

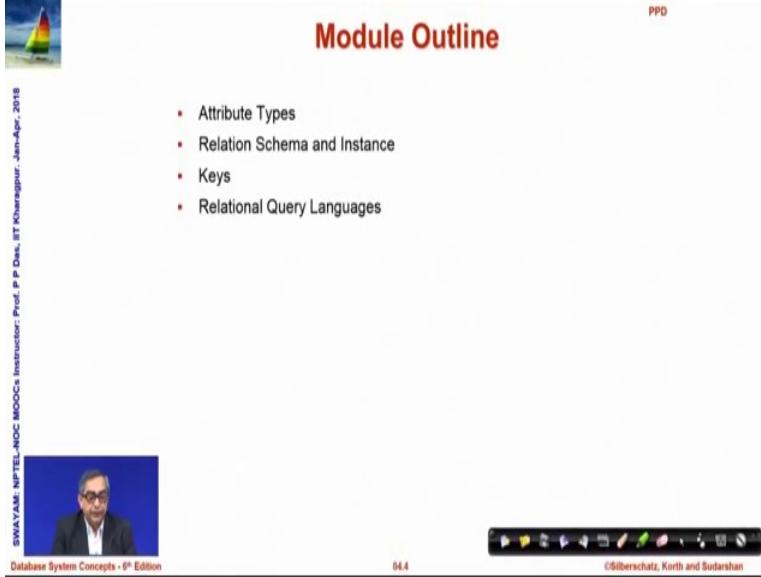


This slide is titled "Module Objectives" in red at the top right. It features a small sailboat icon in the top left corner. On the left side, there is vertical text: "SWAYAM: NPTEL-MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Apri- 2018". At the bottom left is a video frame showing a man in a suit. The bottom right contains the text "Database System Concepts - 8th Edition" and "©Silberschatz, Korth and Sudarshan". A navigation bar with icons is at the bottom right, and the page number "04.3" is at the bottom center.

- To understand attributes and their types
- To understand the mathematical structure of relational model – schema, instance and keys
- To familiarize with different types of relational query languages

And in the current one, our objective will be to understand key concepts of relational model that is attributes and their types, the basic mathematical structure of instance schema and what is known as keys and to familiarize with different types of relational query languages. This is a module outline that we will follow.

(Refer Slide Time: 01:17)



This slide is titled "Module Outline" in red at the top right. It features a small sailboat icon in the top left corner. On the left side, there is vertical text: "SWAYAM: NPTEL-MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Apri- 2018". At the bottom left is a video frame showing a man in a suit. The bottom right contains the text "Database System Concepts - 8th Edition" and "©Silberschatz, Korth and Sudarshan". A navigation bar with icons is at the bottom right, and the page number "04.4" is at the bottom center.

- Attribute Types
- Relation Schema and Instance
- Keys
- Relational Query Languages

(Refer Slide Time: 01:22)



Example of a Relation

SWAYAM-NPTEL-MOOCs Instructor: Prof. P. P. Doshi, IIT Kharagpur - Jan-Apr - 2018

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

attributes
(or columns)

tuples
(or rows)

Database System Concepts - 8th Edition

64.5

©Silberschatz, Korth and Sudarshan

So, we again repeat this example from past, this is an example of instructors, in a university a table of instructors given by attributes or columns ID name, department name and salary. These are the four columns, the four IDs and multiple rows, which are specific rows or we often refer to them as tuples. So, you can say that since there are 4 attributes, that is every row has 4 columns.

So, this is a 4 tuple that we have and such a table is called a relation is as simple as that. So, this is it, whenever we talk about a relation, we have a number of fields number of attributes number of columns, whatever way, we said of a table and that table according to those columns, it has multiple 0 1 or any number of rows of values, filled in and that is what is a relation.

(Refer Slide Time: 02:34)

The slide features a small sailboat icon in the top left corner. In the top right, it says 'PPD'. Below that is a bulleted list under the heading 'Attribute Types':

- Relation Schema and Instance
- Keys
- Relational Query Languages

At the bottom, there's a navigation bar with icons and the text 'Database System Concepts - 8th Edition' and '04.6'.

Now; so, let us look at attributes more specifically.

(Refer Slide Time: 02:37)

The slide has a sailboat icon and 'PPD' in the top right. The title 'Attribute Types' is centered below it. A bulleted list follows:

- The set of allowed values for each attribute is called the **domain** of the attribute
 - Roll #: Alphanumeric string
 - First Name, Last Name: Alpha String
 - DoB: Date
 - Passport #: String (Letter followed by 7 digits) – nullable
 - Aadhaar #: 12-digit number
 - Department: Alpha String
- Attribute values are (normally) required to be **atomic**; that is, indivisible
- The special value **null** is a member of every domain. Indicated that the value is "unknown"

Below the list is a table with columns: Roll #, First Name, Last Name, DoB, Passport #, Aadhaar #, and Department. Data rows are shown for two students. The 'Passport #' column for the second student contains the value 'null', which is circled in blue. Red arrows point from the list items to the corresponding table columns.

So, attributes each column is an attribute as you said, this every attribute has a domain. The domain is a set of possible values that attribute can take. So, if you just look into the example here, so, I am trying to define a table having different students. So, there is a roll number for a student, there is a first name last name, the date of birth; DOB, the passport number, the Aadhaar card number, the department to which the student belongs and so on. So, let us say these 1 2 3 4 5 6 7 are the different attributes.

Now, if we look into every attribute, then every attribute has a set of possible values of which some value is entered in a particular row. For example, the roll number is an alphanumeric string, as you can see, it has numeric as well as it has letters whereas, the first name or the last name are simple alpha strings. In fact, we can also say that the roll number actually is not only alphanumeric, it has a fixed length, here is it has length of 9. So, you can say alphanumeric strings of length 9 are eligible for being values of this domain.

There could be more restrictions, but that the domain will be certain collection of values which are possible as values of that attribute, when you talk about D o B that certainly has to be a date. So, it is written in the form of d d m m m y y y y that is two digit date, three letter month codes and four digit year, the passport number is a string, a letter followed by seven digits. The other number is a twelve digit number, the department is alpha string and so on. So, the domain is a set corresponding to an attribute, which define that all possible values that attribute can take ok.

Now, these attribute values if you look at they are atomic in nature that is you cannot divide them into smaller parts. So, what I mean is say when we are talking about date of birth the whole date of birth type the date type is one atomic value. For example, if you were to code this in C what you could do you could possibly create a structure with three fields; one is date, one is a month, one is a year and we will say that this composite record composite structure is actually my date.

They can do a ls type def, you could do if you are working in C++, you will define a class called date, which has these components and as well as operations with them, but that kind of types are not allowed in a relational database, it has to be an atomic type. So, in a relational database will give you atomic type called date, but all of these are pre specified and has to be taken as one unit, other atomic types are integer like we do not have an integer field here. There are strings; there are numerical values, which are kind of floating point values and so on.

Now, some attribute may have a special value called the null value, which is the member. It is domain; actually every attribute of any domain can have this special value. The null value is not actually a value; it is actually an absence of a value. So, it says that this value is not known. So, if you look into the example above, then you will see that for

passport we have said that the passport is a string letter followed by seven digits and it is nullable, which means that in the passport field, I may have a value, may have this null value which means that it is not that, the passport is null, what it is saying is this passport number for this particular student, the row number 2 is not known, is unknown.

Now, all fields may or may not be nullable. For example, will not allow D O B to be nullable, date of birth has to be there, will not allow roll number to be nullable, will not allow first name to be nullable, but we may allow last name to be nullable. It is been a style, let not to use your last name, many people just use one name. So, you could allow that, it is not known, it is not there whereas, department may not be nullable, it must be there. So, null is a very critical concept and what it actually does? It actually creates a lot of issues and complications in terms of defining many operations. So, understanding null as a value in terms of an attribute is a critical requirement for the design.

(Refer Slide Time: 08:03)

The slide is titled "SCHEMA AND INSTANCE" in large red capital letters. In the top right corner, there is a small "PDF" icon. To the left of the title, there is a vertical sidebar with the following text:
SWAYAM NPTEL MOOC Instructor: Prof. P. P. Datta, IIT Kharagpur - Jam-Apr-2018
Database System Concepts - 8th Edition
04.8
©Bücherschutz, Korth und Söderström
The main content area contains a bulleted list of topics:

- Attribute Types
- Relation Schema and Instance
- Keys
- Relational Query Languages

Now, coming to the schema and instance, we have discussed about the basic understanding of schema and instance. So, understanding them formally now, we say that if we have a schema. So, it is like a table having multiple columns say, there are n columns, having names A 1 to A n, then this A 1 to A n are the attributes.

(Refer Slide Time: 08:26)

The slide has a header 'Relation Schema and Instance' and a bullet point: '• A_1, A_2, \dots, A_n are attributes' with a handwritten checkmark. To the right is a hand-drawn diagram of a table with four columns and three rows, with the last cell in the third row crossed out with a large 'X'. On the left, there's a vertical footer with text: 'SWAYAM: NPTEL-NOOC MOOCs Instructor: Prof. P. D. Patnaik, IIT Kharagpur - Jam-Apr-2018'. At the bottom, it says 'Database System Concepts - 8th Edition' and shows a progress bar at 04.9.

So, these are the different attributes. So, if I have this, then it basically means that I have a table, where these are the columns A 1 A 2 A n like this.

(Refer Slide Time: 08:50)

The slide has a header 'Relation Schema and Instance' and two bullet points: '• A_1, A_2, \dots, A_n are attributes' and '• $R = (A_1, A_2, \dots, A_n)$ is a relation schema'. Below this is a section 'Example:' with the definition 'instructor = ($ID, name, dept_name, salary$)'. Another bullet point states: '• Formally, given sets D_1, D_2, \dots, D_n a relation r is a subset of $D_1 \times D_2 \times \dots \times D_n$ '. Below this is the text 'Thus, a relation is a set of n -tuples (a_1, a_2, \dots, a_n) where each $a_i \in D_i$ '. Handwritten notes include 'A \rightarrow D' with arrows pointing from A to D. On the left, there's a vertical footer with text: 'SWAYAM: NPTEL-NOOC MOOCs Instructor: Prof. P. D. Patnaik, IIT Kharagpur - Jam-Apr-2018'. At the bottom, it says 'Database System Concepts - 8th Edition' and shows a progress bar at 04.8.

So, then a relational schema is a collection of these attributes. So, it is a collection of all these attributes. So, we said R is a relational schema, which has attributes A 1 to A n. Now, every attribute A i has a domain D i. So, for every attribute, I have a set of values that are possible. So, if you, if you recall then here we had different, these are the different attributes and these are their different domain. So, D o B is an attribute and the

domain is date. So, any possible date, other is an attribute and this is the domain, which is A. So, all attributes, each attribute will need to have certain domain and those are marked by the D Sets. So, we will say that a particular relation a particular relation R.

(Refer Slide Time: 10:01)

Relation Schema and Instance

- A_1, A_2, \dots, A_n are attributes
- $R = (A_1, A_2, \dots, A_n)$ is a relation schema
- Example: $\text{instructor} = (ID, name, dept_name, salary)$
- Formally, given sets D_1, D_2, \dots, D_n a relation R is a subset of $D_1 \times D_2 \times \dots \times D_n$. Thus, a relation is a set of n -tuples (a_1, a_2, \dots, a_n) where each $a_i \in D_i$.

$$R = \left\{ (a_1, a_2, \dots, a_n) \in D_1 \times D_2 \times \dots \times D_n \right\}$$

SWAYAM, NPTEL-NOC, Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr., 2018
Database System Concepts - 10th Edition
Gliberschatz, Korth and Sudarshan

So, R is a schema. So, a particular relation R is a subset of $D_1 \times D_2 \times \dots \times D_n$. So, recall the mathematical notion of relation which says that a relation is basically a subset of a set. So, these are the possible values. So, the first attribute can take values from D 1, second attribute can take values from D 2 and so on and the nth attribute can take values from D n. So, any specific row, any specific record is a set of values for A 1 A 2 A n and therefore, is a member of this Cartesian product and the relation is a subset of that. So, this is a D value is an n tuple, which is a subset of this (a_1, a_2, \dots, a_n) .

This particular record is an element of this Cartesian product set and R necessarily is a set of such tuples that is a mathematical view of the schema and the instance. So, this is the schema and this is the instance corresponding to that schema based on the different domains of the different attributes and this is the notion that we will continue using. So, please try to follow this carefully.

(Refer Slide Time: 11:46)

The slide features a small sailboat icon in the top left corner. The title 'Relation Schema and Instance' is centered at the top in a red font. Below the title is a bulleted list of definitions and examples:

- A_1, A_2, \dots, A_n are **attributes**
- $R = (A_1, A_2, \dots, A_n)$ is a **relation schema**

Example:

- $\text{instructor} = (ID, name, dept_name, salary)$
- Formally, given sets D_1, D_2, \dots, D_n a **relation r** is a subset of $D_1 \times D_2 \times \dots \times D_n$. Thus, a relation is a set of n -tuples (a_1, a_2, \dots, a_n) where each $a_i \in D_i$.
- The current values (**relation instance**) of a relation are specified by a table
- An element t of r is a **tuple**, represented by a **row** in a table

At the bottom of the slide, there is a video player interface showing a video of a professor and some navigation icons.

Now, whenever we have an instance, we mark that as a table and every such table.

(Refer Slide Time: 11:56)

The slide features a small sailboat icon in the top left corner. The title 'Relations are Unordered' is centered at the top in a red font. Below the title is a bulleted list of points:

- Order of tuples is irrelevant (tuples may be stored in an arbitrary order)
- Example: *instructor* relation with unordered tuples

Below the list is a table with handwritten annotations. The table has columns labeled ID , $name$, $dept_name$, and $salary$. The rows contain data for various instructors. Handwritten labels a_1, a_2, a_3, a_4 are placed next to the first four rows, with a circled 'C' above them, indicating they are different tuples in an unordered set.

ID	$name$	$dept_name$	$salary$
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

At the bottom of the slide, there is a video player interface showing a video of a professor and some navigation icons.

So, here you have now understood it very well. So, these are my attributes. So, this is A 1, this is A 2, this is A 3, this is A 4 and any one i name is at the different values a 2 a 3 a 1 a 2 a 3 a 4 98345 is a 1 Kim is a 2 and so on. Now, naturally this, it is not visible from the instance, because we are taking an instance view, we are not being able to see, what that domain is that will be visible? If we look at the corresponding D D 1, the definition language description of the schema, which must have specified I D as a numeric value,

the name as a string value the department name as another string value whereas, salary as a numeric value and so on. Now, what is important to note here is a relation necessarily is a set as we said is a set, which is the as the relation R is a set, this is a set, which is a subset of this set.

(Refer Slide Time: 13:04)

The screenshot shows a presentation slide with the following elements:

- PPD** logo in the top right corner.
- A small sailboat icon in the top left corner.
- A vertical sidebar on the left containing the text: "SWAYAM_NIPTEL-NOC MOOCs Instructor: Prof. P. P. Dand, IIT Kharagpur - Jan-Apr. 2016".
- KEYS** is the main title in large red capital letters.
- A video player window showing a man speaking, with the text "Database System Concepts - 8th Edition" below it.
- A navigation bar at the bottom with icons for back, forward, search, and other controls.
- Page number 04.11 and copyright information "©Bilberschatz, Korth and Sudarshan" at the bottom right.

So, we know the elements in a set are do not have any ordering, they are unordered. So, a relation is necessarily unordered. So, it does not really matter that in terms of this collection of rows, which row is at what position, if I reorder them, the relation does not change it is just that they are a collection of this set of rows. So, that lack of ordering is critical information that we will have to remember in mind next concept is key.

(Refer Slide Time: 13:49)

The slide has a header 'PPD' in red at the top right. The title 'Keys' is in red at the top center. On the left, there is a small logo of a sailboat on water. The main content is a bulleted list:

- Let $K \subseteq R = (A_1, A_2, \dots, A_n)$
- K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each possible relation $r(R)$
- Example: {ID} and {ID, name} are both superkeys of *instructor*

On the left side of the slide, there is vertical text: 'SWAYAM', 'NPTEL-NOOC', 'MOOCs', 'Instructor', 'Prof. P. P. Desai, IIT Kharagpur', and '2018'. At the bottom left, it says 'Database System Concepts - 8th Edition'. In the bottom right corner, there is a video player interface showing a timestamp '04.12' and the copyright notice '©Silberschatz, Korth and Sudarshan'.

So, R as we have seen is a relational schema, which is a collection of attributes **A 1, A 2A n**. Now, K let K is a subset of R . So, it is one or more attributes, it has to be a non-empty subset. Now, we will say that K is a super key of R , if we consider the values of different tuples in the attributes of K and we find that there cannot be two tuples, which are different, but match on these attributes, which mean that the values of the attributes of K , uniquely identify each row of the relation, then we will say that K is a super key of R .

So, the instructor table that we have seen, ID is a super key similarly. So, K can be taken as a singleton set of attribute ID or K can be thought of as the set comprising ID and name both of them are super keys of instructor.

(Refer Slide Time: 15:32)

PPD

Keys

- Let $K \subseteq R$
- K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each possible relation $r(R)$
 - Example: $\{ID\}$ and $\{ID, name\}$ are both superkeys of *instructor*
- Superkey K is a **candidate key** if K is minimal
 - Example: $\{ID\}$ is a candidate key for *Instructor*
- One of the candidate keys is selected to be the **primary key**
 - Which one?
- A **surrogate key** (or synthetic key) in a database is a unique identifier for either an *entity* in the modeled world or an *object* in the database
 - The surrogate key is *not* derived from application data, unlike a *natural* (or *business*) key which is derived from application data

Now, we say a super key K is a candidate key, if K is minimal. So, the idea is like this that this is a key, super key, this is also a super key, but certainly this is a subset of this. This is smaller than this. So, we will say this is a candidate key, but this is not a candidate key, because it does not satisfy the minimality condition.

There could be multiple candidate key in a relation, if there are multiple candidates key then we select one to be the primary key. Now; obviously, there is a question of which one we select, but anyone can be selected as a primary key, which is the key of the relation and we will see that in some cases, there is concept of surrogate keys.

So, if I have a relation where there is no attribute, whose value can uniquely identify each and every row of the table then I might synthetically generate a value for example, like a serial number, I can generate a serial number and say that this is my value. So, that serial number or that computer generated field value has no business implication, the real world did not have this value, it is not like a Aadhaar card number or like a passport number, but it is a value which is purely generated to identify every row uniquely. So, such keys are known as surrogate keys or synthetic keys.

(Refer Slide Time: 17:40)

The slide has a header 'PPD' at the top right. On the left, there is a small logo of a sailboat and some vertical text: 'SWAYAM-NPTEL-NCOC', 'Instructor: Prof. P. P. Desai', and 'Date: 20-Apr-2018'. The main title 'Keys' is in red at the top center. Below it is a bulleted list of key types:

- **Super Key:** Roll #, (Roll #, DoB)
- **Candidate Keys:** Roll #, (First Name, Last Name), Passport #, Aadhaar #
 - Passport # cannot be a key. Why?
- **Primary Key:** Roll #
- **Secondary / Alternate Key:** (First Name, Last Name), Aadhaar #
- **Simple key:** Consists of a *single attribute*
- **Composite Key:** (First Name, Last Name)
 - Consists of more than one attribute to uniquely identify an entity occurrence
 - One or more of the attributes, which make up the key, are not simple keys in their own right

Below the list is a table with student data:

Roll #	First Name	Last Name	DoB	Passport #	Aadhaar #	Department
15CS10026	Lalit	Dubey	27-Mar-1997	L4032464	1728-6174-9239	Computer
16EE30029	Jatin	Chopra	17-Nov-1996	null	3917-1836-3816	Electrical
C10016	Smriti	Mongra	23-Dec-1996	G5432849	2045-9271-0914	Electronics
E10038	Dipti	Dutta	02-Feb-1997	null	5719-1948-2918	Civil
S30021	Ramdin	Minz	10-Jan-1997	X8811623	4928-4027-6024	Computer

At the bottom of the slide, there is a footer with the text 'Database System Concepts - 8th Edition', '04.13', and '©Silberschatz, Korth and Sudarshan'.

Now, let us look at some examples, this is again the same student database, I just shown a while ago the same set of columns, but I have added few more rows. Now, if we look at what could be a super key there are several candidates, but I have just written a few roll number is certainly a key, because I am assuming that the university assigns roll numbers to uniquely identify every student.

So, there cannot be two rows in this table, which match in the value of the roll number and does not match in the values of the other fields. So, roll number can uniquely identify, if it can then any super set of attributes, which continual number will also be a super key. So, roll number and date of birth together is a super key that can also unique to identify every row trivial. What are the candidate keys? Now, there are of course, that could be several other super keys that has to be kept in mind, the candidate keys are roll number is a candidate key, the first name last name together, we can say is a candidate key.

So, we are saying that not only the first name, but if we take this pair, you remember the key, the set of attributes forming a super key is a set. It is not an individual field. So, I say the first name last name together, from say, key well. This does make some assumption, because if I say the first name last name together from say key; that means, that there cannot be two records in this student table, where the first name and last name match, but the records are different. So, which mean that no, two students having the

same first name and last name can be enrolled in the university. This is a restrictive assumption right, but I am just making that assumption to illustrate the different possibilities; then what is the other possibility passport number? Everybody has a unique passport number. So, passport number could also be a key, could be a candidate key. Aadhaar number; everybody has a unique Aadhaar number. So, that can be a key and so on.

So, these are called the candidate keys. Now of course, we can observe that given the data it is clear and it was also mentioned when the schema was designed this passport number cannot be a key. Why can it not be a key? Can two students have same passport number? Of course, not every student has a unique passport number, but it is possible that some student does not have a passport. So, if some student does not have a passport then the passport number field of that student is a null, the passport number is a nullable field, if the passport number is null then it is possible that multiple students may not have passports.

So, as we can see here that this student Jatin Chopra does not have a passport. So, similarly, Dipti Dutta does not have a passport either. So, certainly if this were to be the key then for all records, for which passport number is nil, this value would not be able to distinguish them in terms of the rows of the table. So, we have to say that passport number cannot be a key or in other words, we can say that no key can be a nullable field.

No key attribute or a participant to a key attribute could be a nullable field right. So, this is one observation here. And, so that clearly also implies that, if we say that Aadhaar number is a valid candidate key that will mean that for admission to that university having Aadhaar number, would be mandatory, if somebody does not have a Aadhaar number that will have to be null, which is not allowed.

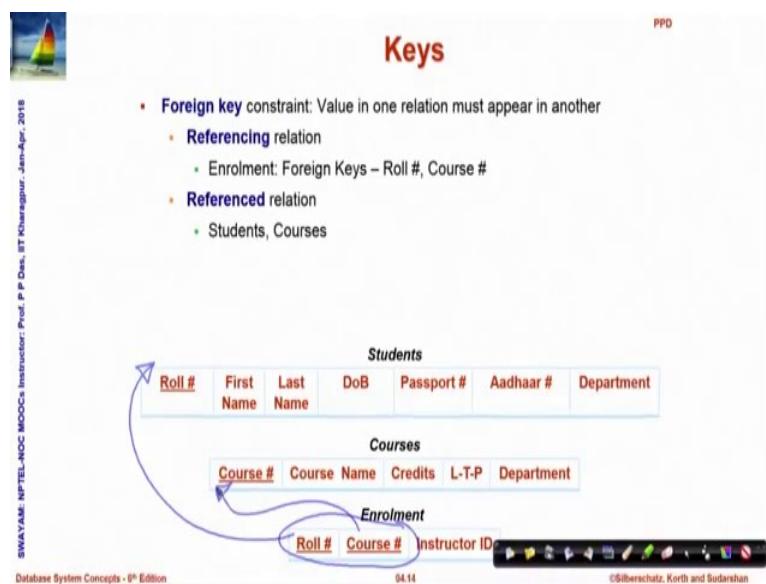
So, let us move on. So, one of these candidate keys have to be made the primary keys. Let us say; we make roll number, the primary key and since, we make roll number the primary key in the schema. We underlined the roll number attribute; this would be a common way to show that roll number is a primary key. So, the others that are not taken as a primary key are called the secondary or alternate key. So, first name last name pair

could be an alternate key Aadhaar number could be an alternate key and so on. A key is said to be simple, if it consists of a single attribute.

So, roll number is a simple key, Aadhaar numbered is a simple key, if it were taken to be primary, but first name last name pair, if we take that to be a primary that will not be considered as simple key, because it has more than one attribute naturally the other, if you have a simple key.

They have other side is a composite key is one, which has more than one field such that none of those fields individually can act as a key, but together they can act as a key. So, first name itself cannot be a key last name itself cannot be a key, but together. They can be a key of course, under the assumption that no two students with the same first name last name are given admission. So, these are the different types of keys that can happen.

(Refer Slide Time: 23:32)



Let us have some more views with the keys, we extend the schema and besides a student I introduce two more schema; one is called the courses, which is given by course number, course name credits L T P. L T P is number of hours of lectures tutorials and practical's and the department. So, these are the different fields and from the convention already stated you can figure out that course number is the key primary key of this relation. I use another schema, which is enrolment, which describes which student is attending which course. So, it has a roll number and the course number.

So, roll number of the student attending the particular course number and it also has an instructor ID as to who is teaching that course given this. You can see that in the enrolment relationship, I have this pair roll number and course number, which will certainly be the key for enrolment, because if I have two rows in enrolment. How they will be distinguished, they cannot be distinguished by roll number, because a particular student may take multiple courses.

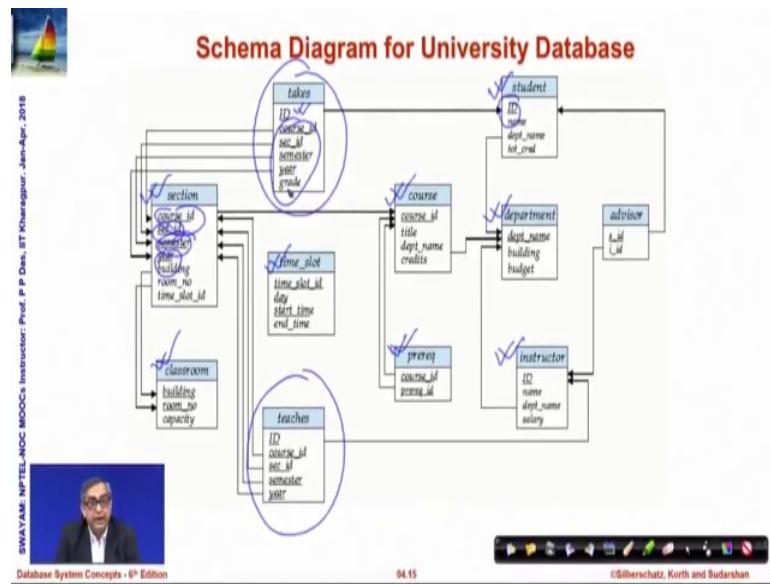
So, there will be multiple records having the same roll number, but different course number, the course number by itself cannot be the key, because every course will have multiple students. So, there will be multiple rows having the same course number, but all different roll numbers, but if we take this together, roll number and course number together then that forms a key.

Now, such a key such a key having roll number the roll number itself is a key of another relation the course number itself is a key of another relation. So, when we take the keys of other relations to form the key of a relation then we say that these are foreign keys. So, roll number and course number are foreign keys in student and course and since from enrolment the student and courses are being referenced are being referred.

So, we say enrolment is a referencing relation and students and courses other reference relation and we will often like to also mention as to what is a foreign key of a relational schema, because that will help us understand how the different schemas are interrelated and we will see that, this will come out directly from the notion of entities and relationships of a year model of a year diagram, a key is called, to be said to be compound, if it consists of more than one attribute to uniquely identify an entity occurrence.

So, each attribute which makes up the key is a simple key, in its own right, mind you there is a subtle, it sounds very similar. We talked about composite key; earlier, we talked up; we are talking about compound key here, the subtlety of the difference is in a composite key, every component attribute is not a simple key by itself, but the components come from the same table in a compound key. The components are simple key in their own right, in some other table and are put together as a compound key in the given table. So, the roll number, course number in the enrolment table is a compound key.

(Refer Slide Time: 27:31)



So, with this I would request you to spend some time with this relatively elaborated schema compared to what we have done already of the university database. So, every rectangular box shows a relational schema on top of each in blue, is written the name of that relation relational schema. So, it has a relational schema like courses, the students the instructors, the departments, the prerequisites, the time slots, the classrooms and so on.

The sections and the relationships between them for example, the relationship is takes is a relationship, which relates students with different sections, with courses, teachers is another relationship, which relates to instructors with sections. So, it is showing you directly as to how the keys of this, what are the attributes? What are the key attributes primary key attributes and also what are the foreign keys that we have in this for example, intakes this is a foreign key, which is featured here, course I D section, I D semester here are the foreign key part of the takes that exists here. So, please study this schema. We will keep on regularly referring to this schema in future as well. So, this is what we have here.

(Refer Slide Time: 29:29)

The slide features a small sailboat icon at the top left. On the right side, there is a vertical list of topics: Attribute Types, Relation Schema and Instance, Keys, and Relational Query Languages. The main title 'RELATIONAL QUERY LANGUAGES' is centered in large red capital letters. Below the title, the text 'Database System Concepts - 8th Edition' is visible. At the bottom right, there is a navigation bar with icons and the text '©Silberschatz, Korth and Sudarshan'. The footer contains the text 'SWAYAM: NPTEL-MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Apr.- 2016'.

Now, we move on to the relational query language, we briefly talk about the relational query language.

(Refer Slide Time: 29:33)

The slide has a small sailboat icon at the top left. The main title 'Relational Query Languages' is centered in red capital letters. Below the title, the text 'Procedural vs. Non-procedural or Declarative Paradigms' is displayed. A bulleted list follows, comparing procedural and declarative paradigms. The list includes: Procedural programming (requires telling the computer what to do), Declarative programming (requires a descriptive style), and an example for square root of n. The example for square root of n is divided into Procedural and Declarative sections. The footer contains the text 'SWAYAM: NPTEL-MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Apr.- 2016', a video thumbnail of a professor, 'Database System Concepts - 8th Edition', the time '04:17', and the copyright notice '©Silberschatz, Korth and Sudarshan'.

Now, we will have to in this the key thing that we need to understand is the relational query language is somewhat different from the programming languages that you have studied so far which are procedural in nature, in contrast the relational query language is non-procedural or declarative in nature, a procedural programming language requires that the programmer tell the computer how to get the output given, the input a pro

program is about finding output, for a given input and you write a procedure, the sequence of steps that need to be done. So, that given the input, you can compute the output.

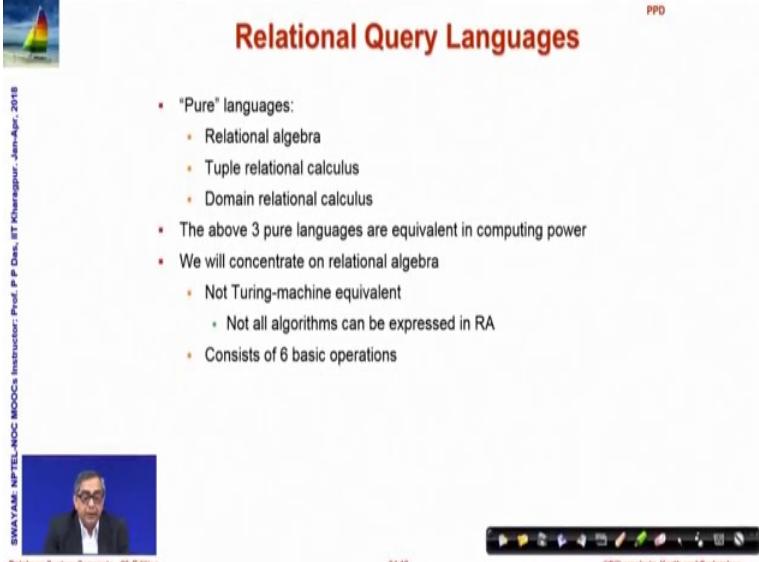
So, you say how that computation has to happen and the programmer must know that algorithm in contrast, in declarative programming, you say what you want? You will do not say how that needs to be computed? How that will be computed? You may not even know that, you may not even know a single algorithm to compute the output, but you specify what output you need. So, this distinction between how and what of programming differentiates procedural and declarative programming.

So, all that you have studied so far in terms of C C++, java python and all that are procedural programming, where you necessarily have to specify, how you will have necessarily; have to specify what the algorithm is, but in declarative you just say what you need. So, just a simple you know pathological example to understand this difference. Suppose, we were interested in computing the square root of a number n assuming n is a positive number, the procedural step would be something like; this is an algorithm that you guess a X or, which is a square root, which is close to the root of n .

I mean some guess you make and then you repeatedly refine this estimate by taking the arithmetic mean of the estimate and the quotient of the division of n by this estimate. So, you take an arithmetic mean and find the new estimate and repeat the steps, I mean as long as the difference between the two conservative estimates is more than a certain value δ , this is a procedural algorithm, you are giving an algorithm. So, given and following this algorithm, we will find the square root declaratively, you can just say that what is the result? I want to result m such that $m^2 = n$.

So, you are again asking for the same feel, you are expecting the same output, but the way you are saying is not an algorithm, you are rather specifying a predicate, which must be true in your output. You are saying that the predicate is m square must be n . So, whatever m is that square of it must equal n . So, this style is known as declarative whereas, the earlier style is known as procedure.

(Refer Slide Time: 32:38)



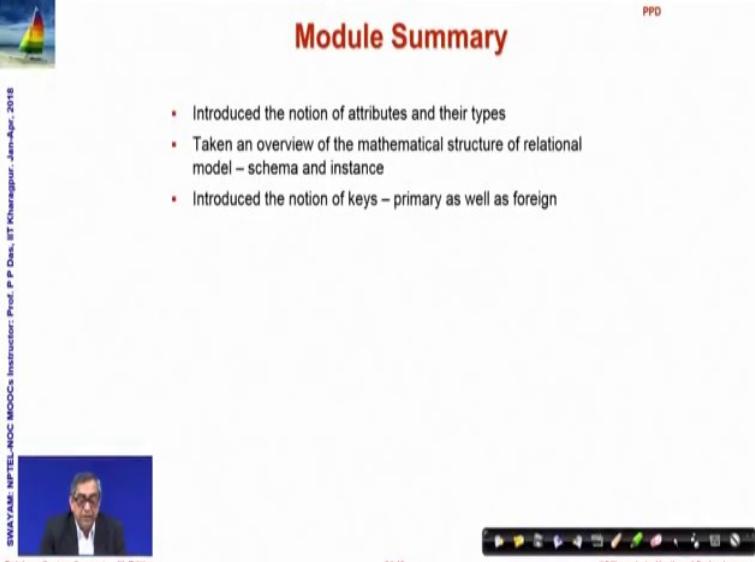
The slide is titled "Relational Query Languages" in red at the top right. It features a small sailboat icon in the top left corner. On the left edge, there is vertical text: "SWAYAM_NPTEL-MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apri- 2018". At the bottom left, it says "Database System Concepts - 8th Edition". The main content is a bulleted list:

- "Pure" languages:
 - Relational algebra
 - Tuple relational calculus
 - Domain relational calculus
- The above 3 pure languages are equivalent in computing power
- We will concentrate on relational algebra
 - Not Turing-machine equivalent
 - Not all algorithms can be expressed in RA
 - Consists of 6 basic operations

At the bottom right, it says "©Silberschatz, Korth and Sudarshan". A navigation bar is at the very bottom.

All query languages, relational query languages are declarative in nature. We have talked about the pure languages, are they are all equivalent? We mentioned that earlier also and also again to remember that none of them are actually Turing equivalent; that means, that not all algorithms can be expressed in them or specifically relational algebra, which we will look at in more depth and the relational algebra will consist of six basic operations, which we will discuss in the next module.

(Refer Slide Time: 33:11)



The slide is titled "Module Summary" in red at the top right. It features a small sailboat icon in the top left corner. On the left edge, there is vertical text: "SWAYAM_NPTEL-MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apri- 2018". At the bottom left, it says "Database System Concepts - 8th Edition". The main content is a bulleted list:

- Introduced the notion of attributes and their types
- Taken an overview of the mathematical structure of relational model – schema and instance
- Introduced the notion of keys – primary as well as foreign

At the bottom right, it says "©Silberschatz, Korth and Sudarshan". A navigation bar is at the very bottom.

So, to sum up we have introduced the notion of attributes and their types, we have taken an overview of the mathematical structure of the relational model schema and instance, we would say mathematically they are relations, mathematically they made a mapping and we have introduced the very important concept of keys and in that very specifically, what is a primary key as well as what is a foreign key? In the next module, we will discuss about the different operations of relational model relational logic.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science & Engineering
Indian Institute of Technology, Kharagpur

Lecture -04
Introduction to Relational Model/2

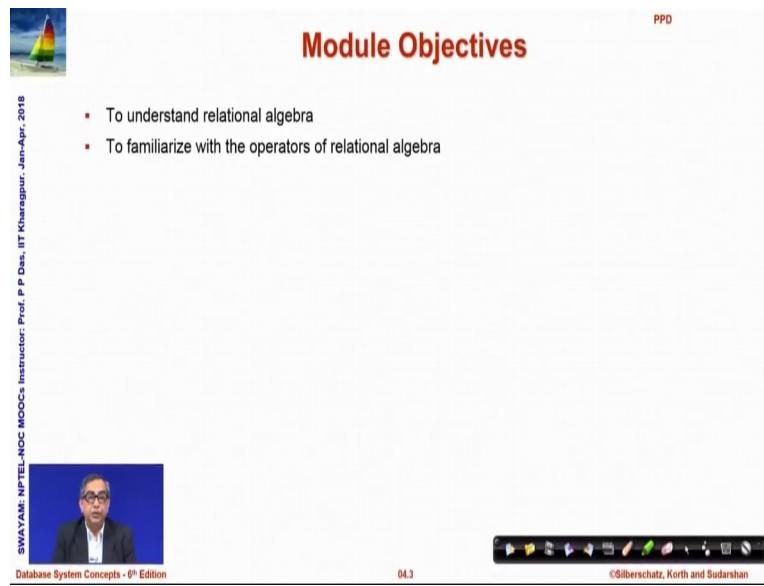
Welcome to module 5 of database management systems. In the previous module, we started discussions on introducing relational model we will conclude that in this module.

(Refer Slide Time: 00:36)

The slide has a title 'Module Recap' in red at the top right. Below it is a list of four bullet points: 'Attribute Types', 'Relation Schema and Instance', 'Keys', and 'Relational Query Languages'. At the bottom left, there is a vertical footer with the text 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P P Das, IIT Kharagpur. Jan-Apr. 2018'. At the bottom center, it says 'Database System Concepts - 8th Edition' and '04.2'. At the bottom right, it says '©Silberschatz, Korth and Sudarshan' and shows a set of small icons.

So, in the last module we have talked about attributes relational schemas and instances in mathematical form and very importantly we have tried to introduce discuss about the concept of keys.

(Refer Slide Time: 00:53)



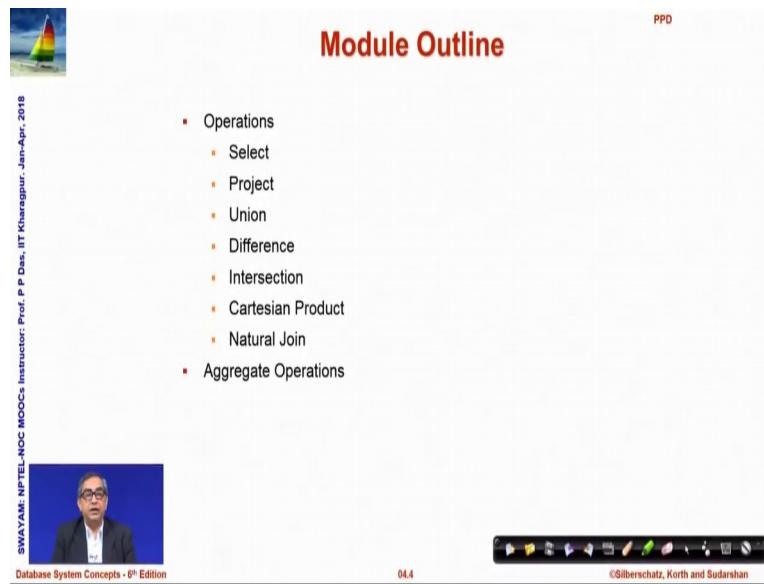
This slide is titled "Module Objectives" in red at the top right. It features a small sailboat icon in the top left corner. On the far right, there is some very small, illegible text. The main content is a bulleted list of objectives:

- To understand relational algebra
- To familiarize with the operators of relational algebra

The footer contains the text "SWAYAM: NIPTEL-NOC MOOCs Instructor: Prof. P Das, IIT Kharagpur, Jam-Apri. 2018", "Database System Concepts - 8th Edition", "04.3", and "©Silberschatz, Korth and Sudarshan". There is also a decorative toolbar at the bottom.

In this module we, will try to understand more on the relational algebra and familiarize with the operations of relational algebra.

(Refer Slide Time: 01:00)



This slide is titled "Module Outline" in red at the top right. It features a small sailboat icon in the top left corner. On the far right, there is some very small, illegible text. The main content is a bulleted list of operations:

- Operations
 - Select
 - Project
 - Union
 - Difference
 - Intersection
 - Cartesian Product
 - Natural Join
- Aggregate Operations

The footer contains the text "SWAYAM: NIPTEL-NOC MOOCs Instructor: Prof. P Das, IIT Kharagpur, Jam-Apri. 2018", "Database System Concepts - 8th Edition", "04.4", and "©Silberschatz, Korth and Sudarshan". There is also a decorative toolbar at the bottom.

So, these are the different operations that we will look at select project union and so on. Some of them are simple set theoretic operations some are newly defined operations and we look at the aggregators.

(Refer Slide Time: 01:18)

The slide features a logo of a sailboat on the left. The title 'Select Operation – selection of rows (tuples)' is at the top right. Below it, a note says 'Relation r'. A table 'r' is shown with columns A, B, C, D and rows containing tuples (α, α, 1, 7), (α, β, 5, 7), (β, β, 12, 3), and (β, β, 23, 10). To the right, a handwritten note shows the selection operation $\sigma_{A=B \wedge D > 5}(r)$. Below this, another table shows the result of the selection: rows (α, α, 1, 7) and (β, β, 23, 10). Handwritten notes next to the result table say 'σ: propositional condition'.

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur. Jam-Apr. 2018

Select Operation – selection of rows (tuples)

Relation r

A	B	C	D
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

$\sigma_{A=B \wedge D > 5}(r)$

A	B	C	D
α	α	1	7
β	β	23	10

σ: propositional condition

So, relational operators, so what is a relation as we have seen already? A relation is nothing but a table. It has a set of columns and it has a set of rows or records that fill up data according to those columns. Select is an operation, which chooses a subset of rows from a relation based on a certain condition. So, it is written in terms of in relational algebra we write it with a notion of notation of σ and following a parenthesis we put the name of the relation.

So, we say we are selecting from the relation r and then we put a condition here, which is a propositional condition. So, if for this all rows of r will be checked if a row will satisfy this condition. Then it will be included in the result, if it does not satisfy the condition, then it will not be included in the result.

(Refer Slide Time: 02:53).

The slide illustrates the selection operation on a relation r . The relation r is defined as:

$$\begin{array}{|c|c|c|c|} \hline A & B & C & D \\ \hline \alpha & \alpha & 1 & 7 \\ \alpha & \beta & 5 & 7 \\ \beta & \beta & 12 & 3 \\ \beta & \beta & 23 & 10 \\ \hline \end{array}$$

A handwritten condition $\Theta: A = B \wedge D > 5$ is shown next to the relation. The result of the selection query $\sigma_{A=B \wedge D > 5}(r)$ is:

$$\begin{array}{|c|c|c|c|} \hline A & B & C & D \\ \hline \alpha & \alpha & 1 & 7 \\ \beta & \beta & 23 & 10 \\ \hline \end{array}$$

The slide also includes a logo for SWAYAM-NIPEL-NOC MOOCs, the name of the instructor Prof. P. P. Das, and the date Jan-Apr. 2018. It also mentions the book 'Database System Concepts - 8th Edition' by Silberschatz, Korth and Sudarshan.

So, let us look at this example. So, our condition is sorry let this put this back. So, our condition Θ is $A=B \wedge D>5$. So, you are saying that that any row to be selected, the value of it is A attribute should equal the value of it is B attribute and when that happens the value of it is $D>5$. So, you can easily if we look through this we can easily say by the first condition $A=B$ you can say that this row does not satisfy this condition.

Because, A is α and B is β whereas, these 3 rows satisfy because $A=\beta$. Then you again look at D we find that this $D<5$. We say this also does not satisfy, because it fails the second condition both have to hold. So, we finally, come to that this record and this record r the selected record in the result $\alpha \alpha 1 7, \beta \beta 23 10$. In both of these records $A=B$ in both of these records $D>5$.

So, selection is a process which selects a subset of rows from a table, from a relation and creates a new relation. Based on a selection condition that must be true for all the rows for all the records that have been selected, but the set of columns do not change they remain the same.

(Refer Slide Time: 04:57)

The slide title is "Select Operation – selection of rows (tuples)". It shows a relation r with fields A, B, C, D and four tuples: $\alpha\ \alpha\ 1\ 7$, $\alpha\ \beta\ 5\ 7$, $\beta\ \beta\ 12\ 3$, and $\beta\ \beta\ 23\ 10$. A handwritten note "σ_{A=B} (r)" is shown above the relation. Below it, a query $\sigma_{A=B \wedge D > 5} (r)$ is shown, with the result being the second tuple: $\alpha\ \beta\ 5\ 7$. Handwritten notes "σ_{A=B} (r)" and "σ_{D>5}" are also present.

So, for example, if I if in the contrary if we say, that $\sigma_{A \neq B} (r)$ then naturally, I will have a relation with fields A B C D and that relation will be only this row. Where that is the only row where $A \neq B$ we can.

(Refer Slide Time: 05:26)

The slide title is "Select Operation – selection of rows (tuples)". It shows a relation r with fields A, B, C, D and four tuples: $\alpha\ \alpha\ 1\ 7$, $\alpha\ \beta\ 5\ 7$, $\beta\ \beta\ 12\ 3$, and $\beta\ \beta\ 23\ 10$. A handwritten note "σ_{A=B} (r)" is shown above the relation. Below it, a query $\sigma_{A=B \wedge D > 5} (r)$ is shown, with the result being the second tuple: $\alpha\ \beta\ 5\ 7$. Handwritten notes "σ_{A=B} (r)" and "σ_{D>5}" are also present.

I can have a selection saying r where $C > 0$. Naturally, this will satisfy, this will satisfy, this will satisfy, this will satisfy. So, this whole relation would be the result of the selection. So, it is possible that it is not necessary that some rows will have to get eliminated in the result.

(Refer Slide Time: 05:54)

Select Operation – selection of rows (tuples)

- Relation r

	A	B	C	D
α	α	1	7	\times
α	β	5	7	\times
β	β	12	3	\times
β	β	23	10	\times

- $\sigma_{A=B \wedge D > 5}(r)$

	A	B	C	D
α	α	1	7	
β	β	23	10	

I say that this is the selection is $D < 1$ this will fail, this will fail, this will fail. So, all of them will fail so, it is possible that the result of an operation could be either the whole relation as we saw last time or a null relation which has where none of the record with feature because, none of the record satisfy the condition. So, that is a basic select operation.

(Refer Slide Time: 06:38)

Project Operation – selection of columns (Attributes)

- Relation r :

	A	B	C
α	10	1	
α	20	1	
β	30	1	
β	40	2	

- $\Pi_{A,C}(r)$

	A	C
α	1	
α	1	
β	1	
β	2	

Let us move on look at the next one is called a projection operation. So, select chooses a subset of the rows. Projection necessarily chooses projects a set of columns from the

original relation. So, this is quite straightforward to see, it is written in terms of this notation Π and then you write the columns that you want in the result of projection. So, we say this is $\Pi_{A,C}$. So, which means basically the column, which is not selected in the projection, you can simply forget about that simply erase it. If you erase that you get this relation and once you get that, please recall that a relation is a set and in a set every element has to be distinct. So, after erasing B this first and the second row have become identical.

So, naturally both of them cannot be there, it will have to be made distinct by erasing any one of them. And hence, they become one row in the result; obviously, I can project on any of the individual single or all the fields also I can do a $\Pi_{A,B,C}(r)$ of course, that means, in this case that will mean that it is a set which is equal to r anyway, but; obviously, I must have at least 1 column at least one attribute to project on. I cannot project on a null set of attributes because that does not give me a schema.

So, there will have to be some attribute one or more attribute on which I project. So, selection and projection, selection has given me the set of rows to retain and projection has given me what are the columns to retain in the result. And combining them I can do several different operations in a database table which can give me several interesting results. Before proceeding further,

Let us look into some of the typical other operations that relational algebra allows. The next one is U given two relations I can take a U this is nothing but a set theoretic union. The two relations r and s must have the same set of columns A and B because, if the columns are not same, then the U does not make sense. They because certainly if the columns are different attributes are different their types and type of data values would be different. So, they cannot be put to a same table so when two relations have the same set of attributes then their instances can be taken a union of.

(Refer Slide Time: 09:49)

The slide title is "Set difference of two relations". It shows two relations, r and s , represented as tables:

- Relations r, s :
 - r (left):

A	B
α	1
α	2
β	1
 - s (right):

A	B
α	2
β	3

$r - s$:

A	B
α	1
β	1

Arrows indicate the records unique to r (α1, β1) are highlighted.

SWAYAM: NPTEL MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur, Jan-Apr. 2018
Database System Concepts - 8th Edition
04.8 ©Silberschatz, Korth and Sudarshan

So, all records that exist in both these relations will be put together into a single table.
So, here $\alpha 1$ is coming here $\beta 1$ is coming here.

(Refer Slide Time: 10:09)

The slide title is "Union of two relations". It shows two relations, r and s , represented as tables:

- Relations r, s :
 - r (left):

A	B
α	1
α	2
β	1
 - s (right):

A	B
α	2
β	3

$r \cup s$:

A	B
α	1
α	2
β	1
β	3

Arrows indicate the records unique to each relation are combined into a single table.

SWAYAM: NPTEL MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur, Jan-Apr. 2018
Database System Concepts - 8th Edition
04.8 ©Silberschatz, Korth and Sudarshan

In terms of relation, $\alpha 1$ is coming here, $\alpha 2$ is coming here, $\beta 1$ is coming here, $\beta 3$ is coming here and $\alpha 2$ is coming here. You can see that this record $\alpha 2$ exists in both the relations. And by the set theoretic notion of uniqueness, in the \cup they have to be unified. So, one of them will be removed, it does not matter because they are identical

anyway. So, for relations having the same set of attributes we can simply make a \cup of all its records.

(Refer Slide Time: 10:57)

The slide is titled "Set difference of two relations". It shows two relations, r and s , represented as tables:

- Relations r, s :**

A	B
α	1
α	2
β	1

A	B
α	2
β	3
- $r - s$:** A third table representing the result of the set difference operation, which contains tuples from r that are not in s .

A	B
α	1
β	1

On the right, there is a Venn diagram with two overlapping circles labeled r and s . The region of circle r that does not overlap with circle s is shaded, visually representing the set difference $r - s$.

Vertical text on the left edge of the slide reads: "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Date, IIT Kharagpur, Jan-Apr., 2018".

At the bottom left is a small video thumbnail of a professor. At the bottom center is the text "Database System Concepts - 8th Edition" and "04.9". At the bottom right is the copyright notice "©Silberschatz, Korth and Sudarshan".

So, other the 3rd operation is the a 4th operation is doing a set difference it is works simply as set theoretic difference. Again, the two relations must have the same set of attributes and I can do a difference of $r-s$ which mean that all tuples which exist in r , but do not exist in s will be included. So, this is included because, this is not here, but this is not included because it is in s this is included.

Because, this is this does not exist in the set s so it is the set, so you take the set r and then erase all the records like this which exist in s and you get $r-s$. So, if I look into just to recap if I look into the Venn diagram, then this is these are set $r-s \in r$, but $\notin s$. So, there is a 4th operation that one can do with the in relational algebra.

(Refer Slide Time: 12:28)

The slide is titled "Set intersection of two relations". It shows two relations, r and s , represented as tables:

- Relation r :

A	B
α	1
α	2
β	1
- Relation s :

A	B
α	2
β	3

The intersection $r \cap s$ is shown as:

A	B
α	2

A hand-drawn Venn diagram illustrates the intersection as the overlap of two circles labeled r and s . A note at the bottom left says "Note: $r \cap s = r - (r - s)$ ".

SWAYAM-NIPEL-MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apri. 2018
Database System Concepts - 8th Edition
04.10
©Silberschatz, Korth and Sudarshan

Fifth is a set intersection(\cap) of 2 relations. So, again the two relations need to have the same set of attributes. You can take their intersection, which is the record which belongs to both. It is the record that belongs to both of them and as you are aware now, set intersection actually is not a new operation. It is not a fundamental operation because, if I have r , if I have s , then this is $r - s$.

Now, if I subtract this $r - s$ which is this set from r which is this bigger set, then what will be remaining? This is what will be remaining? So, if I subtract $r - s$ from r then what will remain? Is necessarily the intersection of this is $r \cap s$. So, set intersection is not a fundamental operation of relational algebra. But, can be used because it can be expressed in terms of set difference.

(Refer Slide Time: 14:03)

joining two relations -- Cartesian-product

- Relations r, s :

A	B
α	1
β	2

 r

C	D	E
α	10	a
β	10	a
β	20	b
γ	10	b

 s

- $r \times s$:

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. Das, IIT Kharagpur. Jam-Apr. 2018
Database System Concepts - 8th Edition
04.11 ©Silberschatz, Korth and Sudarshan

Next comes how can we join two different relations which have different set of attributes? So, relation r has A B and relation s has C D E. So, we can take a Cartesian product(\times). So, taking Cartesian product is making all possible combinations. So, necessarily if since this has two relations and this has 3 relations. So, this will have 1 2 3 4 5 6 7 8 this has 4 relations. So, these are 8, 8 total all possible pairing of relations of r or records of r and records of s are included. So, that is the Cartesian product all possible combinations. This is this is how we can join two relations, but certainly what is a important is something which we will discuss shortly.

(Refer Slide Time: 15:00)

Cartesian-product – naming issue

- Relations r, s :

A	B
α	1
β	2

 r

C	D	E
α	10	a
β	10	a
β	20	b
γ	10	b

 s

- $r \times s$:

A	r.B	s.B	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. Das, IIT Kharagpur. Jam-Apr. 2018
Database System Concepts - 8th Edition
04.12 ©Silberschatz, Korth and Sudarshan

Now, in the Cartesian product then issue may happen because there could be attributes which are common between two relations. So, if you if two attributes are common when you take Cartesian product how do you put their name? Because as with the example show here between r and s the attribute b is common so, how do you take care of that? So, when the, such common names happen then we actually change the name of the attribute with the name of the relation. So, b coming from r will be called **r.b** and b coming from s will be called **s.b**.

(Refer Slide Time: 15:46)

Renaming a Table

- Allows us to refer to a relation, (say E) by more than one name.
 $\rho_x(E)$
returns the expression E under the name X
- Relations r

	A	B
α	1	
β	2	

r

$r \times \rho_s(r)$

	$r.A$	$r.B$	$s.A$	$s.B$
α	1	a	1	
α	1	β	2	
β	2	a	1	
β	2	β	2	

$r \times r$

$r.A$ $r.B$ $s.A$ $s.B$

©Silberschatz, Korth and Sudarshan

And accordingly, the relational algebra, gives you a way to rename a table and put it is name differently. So, this is given a this is given by this relation is given by this symbol ρ . So, you can using a relation r you can actually give it a different name s and do that in terms of. So, with that you can actually because otherwise you cannot compute $r \times r$. Because, if you try to do compute $r \times r$ then you will have $r.a r.b$ and you will again have $r.a r.b$. So, you are using this to rename r to s and then compute this. So, renaming a table is another feature which is provided of course, it is not a fundamental operation of the algebra, but this is what makes the any kind of Cartesian product possible.

(Refer Slide Time: 16:55)

The slide features a sailboat icon in the top left corner. The title 'Composition of Operations' is centered at the top in a red font. Below the title, there is a bulleted list of points:

- Can build expressions using multiple operations
- Example: $\sigma_{A=C}(r \times s)$

Below the list, there are two tables representing relations r and s. Relation r has attributes A, B, C, D, E and contains the following tuples:

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

Relation s has attributes A, B, C, D, E and contains the following tuples:

A	B	C	D	E
α	1	α	10	a
β	2	β	10	a
β	2	β	20	b

At the bottom left is a video thumbnail showing a man speaking. The footer includes the text 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018', 'Database System Concepts - 8th Edition', '04.14', and '©Silberschatz, Korth and Sudarshan'.

Finally, we can make composition of operations that if for example, what we show here is we have two relations r and s you have taken a Cartesian product and then we have taken a selection. So, taken a Cartesian product of r and s to produce the table as you can see the process($r \times s$) and then $\sigma_{A=C}(r \times s)$ based on that condition. So, all these operations can be combined in multiple different ways to give you really complex relational algebra operations.

(Refer Slide Time: 17:33)

The slide features a sailboat icon in the top left corner. The title 'Joining two relations – Natural Join' is centered at the top in a red font. Below the title, there is a bulleted list of points:

- Let r and s be relations on schemas R and S respectively. Then, the "natural join" of relations R and S is a relation on schema $R \cup S$ obtained as follows:
 - Consider each pair of tuples t_r from r and t_s from s .
 - If t_r and t_s have the same value on each of the attributes in $R \cap S$, add a tuple t to the result, where
 - t has the same value as t_r on r
 - t has the same value as t_s on s

At the bottom left is a video thumbnail showing a man speaking. The footer includes the text 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018', 'Database System Concepts - 8th Edition', '04.15', and '©Silberschatz, Korth and Sudarshan'.

There is a nice operation which is a derived one which can be written in terms of other operations which is called a natural join(\bowtie) which we will use very heavily.

(Refer Slide Time: 17:48)

Natural Join Example

- Relations r, s:

	A	B	C	D
r	α	1	α	α
	β	2	γ	α
	γ	4	β	b
	α	1	γ	a
	δ	2	β	b

	B	D	E
s	1	α	α
	3	α	β
	1	α	γ
	2	b	δ
	3	b	ϵ

- Natural Join
- $r \bowtie s$

	A	B	C	D	E
	α	1	α	α	α
	α	1	α	α	γ
	α	1	γ	α	α
	α	1	γ	α	γ
	δ	2	β	b	δ

$\Pi_{A, r.B, C, r.D, E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$

SHIVAM: NITEL NOC MOOC Instructor: Prof. P. P. Date, IIT Kanpur, Jan-Apr. 2014
Database System Concepts - 6th Edition

04.16 ©Silberschatz, Korth and Sudarshan

Let me first show you an example of that, suppose I have two relations r and s and what is important is there are some attributes which are common between them. Now, we saw earlier that in Cartesian product, in terms of common attributes. We basically the attributes got renamed in terms of the table name, but this is not what we are looking at in relation natural joint.

What we want to say is, if an attribute is common between two tables then, while you join them, the records from two fields can be joined if their value on that common attribute is same. So, it is what it tries to do is it tries to make a Cartesian product of this two tables ($r \times s$) first take all possible combinations, but then you select only those rows where the values are identical between columns having the same name. So, for example, if you if you look into this row and this row. So, what will happen in the Cartesian product I will have this **$\alpha 1 \alpha a 1 a \alpha$** .

Let me write it in smaller. So, I am doing a Cartesian product I am looking at this row I am looking at this row.

(Refer Slide Time: 19:39)

Natural Join Example

- Relations r, s :

	A	B	C	D
α	1	α	a	
β	2	γ	a	
γ	4	β	b	
α	1	γ	a	
δ	2	β	b	

r

	B	D	E
1	a	α	
3	a	β	
1	a	γ	
2	b	δ	
3	b	ϵ	

s

- Natural Join
- $r \times s$

	A	B	C	D	E
α	1	α	a	α	
α	1	α	a	γ	
α	1	γ	a	α	
α	1	γ	a	γ	
δ	2	β	b	δ	

$\Pi_{A, r.B, C, r.D, E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apri- 2018

Database System Concepts - 8th Edition 04.16 ©Silberschatz, Korth and Sudarshan

So, I have A B C D I have B D E and $\alpha \alpha a 1 a \alpha$. Now, they match on B they match on D. So, I will say this is this will get retained, but in the Cartesian product I will also have the first row of r going with the second row of s $\alpha 1 \alpha a 3 a \beta$. Here, the B does not match D does not does match, but the B does not match. So, this particular entry will not go in the final result. To take the Cartesian product and you only retain those rows where the values match for the identically named attribute. That is why, so you take the Cartesian product now look at the expression the attributes common attributes are B and C, B and D.

So, the B attribute is $r.B$ and $s.B$. So, we say that in the Cartesian product $r.B = s.B$. The name is common; the value will have to be same similarly D is a common attribute. So, $r.D$ value in the $r.D$ and the value in the $s.D$ has to be same. So, based on the Cartesian product you do a selection for equality of values on attributes which are identical between the two column between the two relations between the two tables.

This is a final selection, as you do that you get a table where there are two Bs $r.B$ $s.B$ there are two Ds $r.D$ $s.D$, but according to this selection for all at all records for all rows the value on $r.B$ and value on $s.B$ are same value on $r.D$ and value on $r.s.B$ are same. Because that is how we have done the selection. So, there is no reason to keep two B columns or two D columns.

(Refer Slide Time: 22:33)

Natural Join Example

- Relations r, s :

	A	B	C	D
α	1	α	a	
β	2	γ	a	
γ	4	β	b	
α	1	γ	a	
δ	2	β	b	

r

	B	D	E
1	a	α	
3	a	β	
1	a	γ	
2	b	δ	
3	b	ϵ	

s

- Natural Join
- $r \bowtie s$

	A	B	C	D	E
α	1	α	a	α	
α	1	α	a	γ	
α	1	γ	a	α	
α	1	γ	a	γ	
δ	2	β	b	δ	

$\Pi_{A, r.B, C, r.D, E}(\sigma_{r.B = s.B \wedge r.D = s.D}(r \times s))$

Handwritten notes: \cancel{B}, \cancel{D}

SWAYAM, NPTEL-NOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apr. 2018
Database System Concepts - 8th Edition
04.16 ©Silberschatz, Korth and Sudarshan

So now, you project based on **A r .B C r.D and E** which means that s.B and s.D are left out you do not project them. So, after you have done this projection, you get the final result of the natural join, which has a union of all the attributes that the two relations at A B C D and B D E union is A B C D E and you have all those records whose values matched on the common attributes between relation r and relation D. So, you can say that if I now do a selection, if I now do a projection on A B C or rather A B C D.

I will get a subset of r if I do a projection on B D E I will get a subset of s. So, this is the natural join operation in relational algebra as you can see this is a derived operation because we could use the Cartesian product selection and projection to get this, but as I tell you we will see more of this when we look at the look at all these different aj query coding, but natural join is one of the most widely used most fundamental relational algebra operation that you will often need beyond selection and projection. So, these were the 6 operations and the important derived operations of relational algebra. Besides that relational algebra has some aggregation operators.

(Refer Slide Time: 24:13)

The slide is titled "Aggregate Operators" in red at the top center. In the top left corner, there is a small image of a sailboat on water. The top right corner has the letters "PPD". On the left side, there is vertical text: "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apri- 2018". At the bottom left, it says "Database System Concepts - 8th Edition". At the bottom right, it says "04.18 ©Silberschatz, Korth and Sudarshan". The footer also contains several small icons for navigating through the presentation.

- Can we compute:
 - SUM
 - AVG
 - MAX
 - MIN

For example, given a table we could compute the sum of values on a column we could compute average of values, max of values, mean of values. So, these somehow aggregate values of multiple rows on a particular column. And therefore, these are called aggregate operators we will see when we talk about SQL. We will see how this really can be coded in SQL and used, but these are this become very convenient to use because often we will need to know.

If this is I mean these are the instructors. And so, let us see which instructor has a maximum load of courses? How many based on hours or which instructor has? What is the average load on the different instructors and so on? So, in every possible context different aggregate operators are frequently required and they are also available as operators in most of the pure as well as commercial query languages.

(Refer Slide Time: 25:18)

The slide has a header 'Notes about Relational Languages' with a sailboat icon. It lists four bullet points: 'Each Query input is a table (or set of tables)', 'Each query output is a table', 'All data in the output table appears in one of the input tables', and 'Relational Algebra is not Turing complete'. The footer includes 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur', 'Database System Concepts - 8th Edition', '04.19', and '©Silberschatz, Korth and Sudarshan'.

- Each Query input is a table (or set of tables)
- Each query output is a table
- All data in the output table appears in one of the input tables
- Relational Algebra is not Turing complete

Finally, to note that relational algebra in relational average every query input is a table. And the output is also a table, it is always manipulating one or more tables into a single table that is what we are, I mean in very simple terms that is the way you can look at it. And all data in the output table appears in one of the input tables that is no new data gets generated. It is basically taking, selecting, combining data from different input tables. It does not generate a new data that is that has to be that is if I see that a, we an attribute for a particular row has a value 15.

Then there must be some input table where there is a record where in that field as an attribute value 50. Otherwise, this cannot happen again relational algebra is not turing complete in the sense that, there are algorithms which cannot be coded in relational algebra. We mentioned this earlier too and that is the reason that is a foundational reason of why the SQL language commercial SQL language is based on relational algebra is not turing complete there.

So, we might need to use other programming languages along with the relational algebra coding for solving some of the application problems.

(Refer Slide Time: 26:50)



Summary of Relational Algebra Operators

Symbol (Name)	Example of Use
σ (Selection)	$\sigma \text{ salary} \geq 85000 \text{ (instructor)}$ Return rows of the input relation that satisfy the predicate.
Π (Projection)	$\Pi ID, \text{salary} \text{ (instructor)}$ Output specified attributes from all rows of the input relation. Remove duplicate tuples from the output.
\times (Cartesian Product)	$\text{instructor} \times \text{department}$ Output pairs of rows from the two input relations that have the same value on all attributes that have the same name.
\cup (Union)	$\Pi name \text{ (instructor)} \cup \Pi name \text{ (student)}$ Output the union of tuples from the two input relations.
$-$ (Set Difference)	$\Pi name \text{ (instructor)} - \Pi name \text{ (student)}$ Output the set difference of tuples from the two input relations.
\bowtie (Natural Join)	$\text{instructor} \bowtie \text{department}$ Output pairs of rows from the two input relations that have the same value on all attributes that have the same name.



 Database System Concepts - 8th Edition 04.20 ©Silberschatz, Korth and Sudarshan

To summarize this is the, so this table is what you not only should remember, but you should become a expert of in terms of the operators of relational algebra which we will start using very heavily as we start doing the query coding and processing. So, we talked about selection, which takes rows selectively we talked about projection which takes out certain columns of a table. We talked about Cartesian product of two relations which make all possible combined relations.

We talked about union of records from two tables having identical set of attributes. We talked about set difference which again is the difference of records of one relation from another, given that they have identical set of attributes. We have shown that set difference can be used to also compute set intersection, it is not in the fundamental operation, but is the derived 1 and we have shown a very interesting operation based on Cartesian product selection and projection called natural join where two tables can be joined based on 1 or more common attributes they have.

Now, I am sure you have already noted that if I am doing a natural join between two tables. Which do not have any common attribute then the result is merely the Cartesian product because the selection around the Cartesian product has no condition to select any of the you know any of the fields any of the rows separately. So, it merely turns out to be a Cartesian product.

(Refer Slide Time: 28:38)

Module Summary

- Introduced relational algebra
- Familiarized with the operators of relational algebra

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur. Jam-Apr. 2018

Database System Concepts - 8th Edition

04.21

©Silberschatz, Korth and Sudarshan

So, in this module, we have introduced the relational algebra and we have familiarized ourselves with the fundamental and derived operators of relational algebra. Going forward in the next module, will take a deeper look into the relational model and start progressing towards the query design and database design.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science and Engineering
Indian Institute of Technology, Kharagpur

Lecture-06
Introduction to SQL/1

Welcome to module 6 of database management systems, this is the starting of week 2. In week 1 we have done 5 modules, after the overview of the course, we have primarily introduced the basic notions of DBMS and we have discussed about the relational module, the fundamentals of it.

(Refer Slide Time: 00:28)

Week 01 Recap

- **Module 01: Course Overview**
 - Why Databases?
 - KYC: Know Your Course
- **Module 02: Introduction to DBMS/1**
 - Levels of Abstraction
 - Schema & Instance
 - Data Models
 - DDL & DML
 - SQL
 - Database Design
- **Module 03: Introduction to DBMS/2**
 - Database Design
 - Database Engine, Users, Architecture
 - History of DBMS
- **Module 04: Introduction to Relational Model/1**
 - Attribute Types
 - Relation Schema and Instance
 - Keys
 - Relational Query Languages
- **Module 05: Introduction to Relational Model/2**
 - Relational Operations
 - Aggregate Operations

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jun-Aug' 2018

Database System Concepts - 6th Edition

06.2

©Silberschatz, Korth and Sudarshan

(Refer Slide Time: 00:54)

Module Objectives

- To understand relational query language
- To understand data definition and basic query structure

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P P Das, IIT Kharagpur.. Date: June-August, 2018

Database System Concepts - 6th Edition

PPD

06.3

©Silberschatz, Korth and Sudarshan

In this background, in this week we are primarily focusing on the query language the structured query language SQL. So, all the 5 modules will relate to discussions on query language and this module 6 7 and 8, the 3 modules will introduce SQL at a first level and the last 2 modules 9 and 10, I am sorry 9 and 10 will discuss about intermediate, that is somewhat advanced level of features in the SQL.

The objective of the current module is to, understand the relational query language and particularly the data definition and the basic query structure, that will hold for all SQL queries, particularly this the modules this week would be important for writing any kind of database applications, squaring the database to find information from the existing data and to manipulate it. So, please put a lot of focus in the whole material of this week and practice them well, to understand the basic issues of database systems, in depth in a well oriented manner. In this module we will first talk about the history and then we will see how to define data and start manipulating them.

(Refer Slide Time: 02:25)

Module Outline

- History of SQL
- Data Definition Language (DDL)
- Basic Query Structure (DML)

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur.. Date-Apr-2018
Database System Concepts - 6th Edition 06.4 ©Giberschatz, Korth and Sudarshan

(Refer Slide Time: 02:37)

History of Query Language

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
 - SQL-86
 - SQL-89
 - SQL-92
 - SQL:1999 (language name became Y2K compliant!)
 - SQL:2003
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
 - Not all examples here may work on your particular system

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur.. Date-Apr-2018
Database System Concepts - 6th Edition 06.6 ©Giberschatz, Korth and Sudarshan

SQL was originally called IBM sequel language and was a part of system R, it was subsequently renamed as structured query language and like any other good programming language that we have, SQL also gets standardized by ANSI and ISO and there have been several standards of SQL, that has come up with SQL 92 being the most popular one, and the commercial system, most of them try to provide support for SQL 92 features, but they do vary between themselves. So, it is possible that the examples that we show here, may or may not all of them execute in the system that you are using. So, you will have to look at what standard your system is actually following.

So, the first what we will talk about is a DDL data definition language, as we had discussed earlier this is the features to create that schema, the tables in a database management system.

(Refer Slide Time: 03:59)

The slide has a decorative header featuring a sailboat on water. The title 'Data Definition Language' is centered in red. Below the title is a text block and a bulleted list. In the bottom left corner, there is a video player showing a person speaking, with the text 'Database System Concepts - 8th Edition' above it. The bottom right corner contains a navigation bar with icons and the text 'Übersicht, Korth und Sudarshan'.

The SQL data-definition language (DDL) allows the specification of information about relations, including:

- The schema for each relation
- The domain of values associated with each attribute
- Integrity constraints
- And as we will see later, also other information such as
 - The set of indices to be maintained for each relations
 - Security and authorization information for each relation
 - The physical storage structure of each relation on disk

So, it allows for the creation or definition of the schema, for each relation that we have in the database it specifies the domain of values associated with each attribute of the schema and it also defines a variety of integrity constraints, later in the course we will see that, it also has to specify other related information like indexing, security authorization, physical storage and so on.

(Refer Slide Time: 04:30)

The slide has a header 'Domain Types in SQL' with a sailboat icon. The content lists various SQL domain types with descriptions:

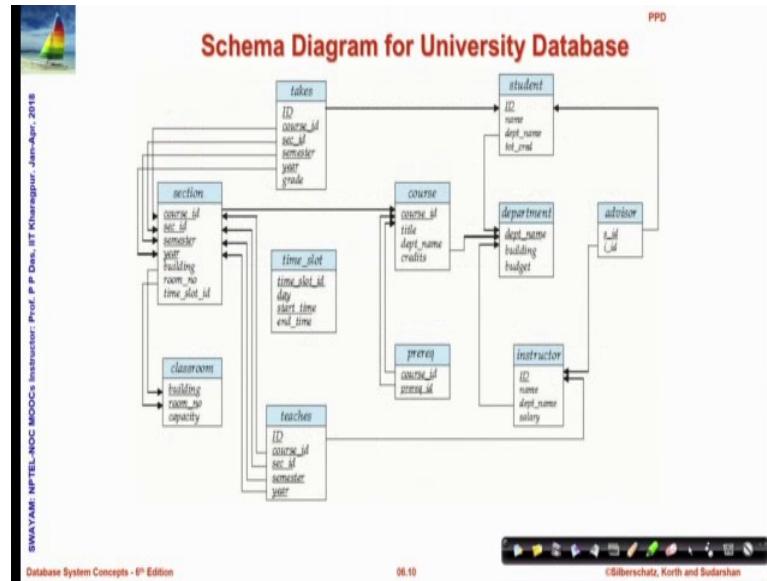
- **char(n).** Fixed length character string, with user-specified length n
- **varchar(n).** Variable length character strings, with user-specified maximum length n
- **int.** Integer (a finite subset of the integers that is machine-dependent)
- **smallint.** Small integer (a machine-dependent subset of the integer domain type)
- **numeric(p,d).** Fixed point number, with user-specified precision of p digits, with d digits to the right of decimal point. (ex., **numeric(3,1)**, allows 44.5 to be stored exactly, but not 444.5 or 0.32)
- **real, double precision.** Floating point and double-precision floating point numbers, with machine-dependent precision
- **float(n).** Floating point number, with user-specified precision of at least n digits
- More are covered in Chapter 4

Navigation icons and slide information are visible at the bottom.

So, first the domain of possible values. So, we have already specified that, every domain in SQL is more of an atomic nature. So, they are more like the primitive or built in data types of languages like c, c++, java. So, the common domain types are character which are basically strings of character, having a certain length then you can have variable character string which means that the length here specifies that, the maximum length that the string can take, but a string could be shorter than that integer; obviously, then small integer, which in a system may give a smaller range of integer values.

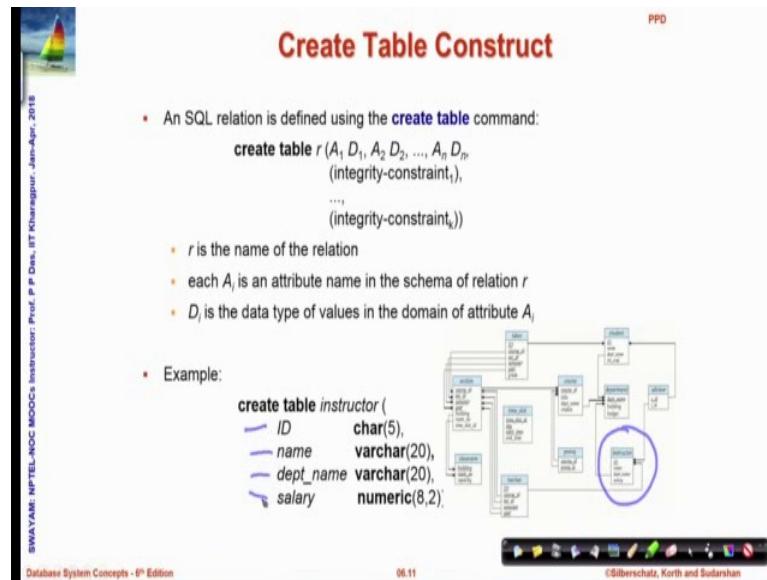
Then there is a numeric type which is often very important, which says what is the precision of the numbers that are to be written in this format stored in this format. So, d basically gives that precision value and p gives the size, then you can have real and double precision numbers you can have floating point numbers and so on, and there are some more data types, which we will discuss later in the course.

(Refer Slide Time: 05:52)



Given all these domain types; so, what we will try to do is, here is a schema for the university database, which has multiple different relations designed in that showing the attributes and marking out, what are the keys? And what are the foreign keys? So, we would take examples of some of these and try to code them in the SQL.

(Refer Slide Time: 06:19)

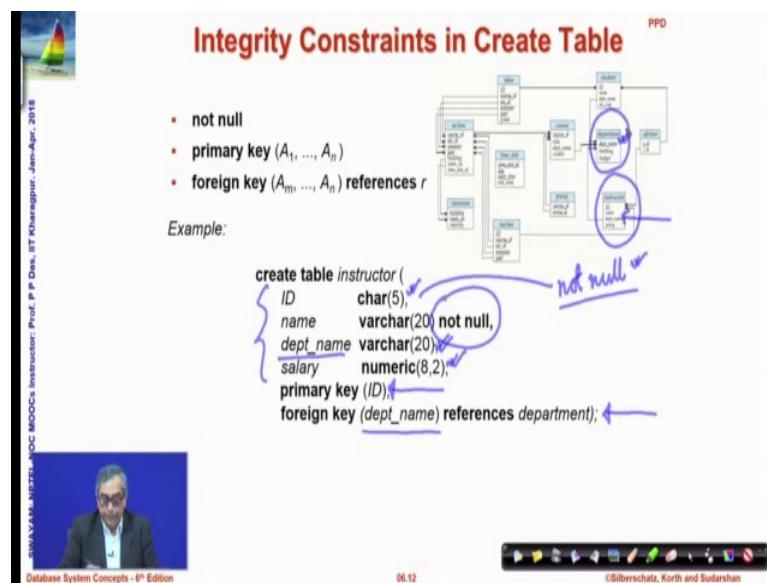


Now, to create a table this is how you go about, the SQL keywords CREATE TABLE is the basic command. So, with the CREATE TABLE you have to specify a name, here the name that is given is in terms of this name r, which is the name of the relation and then

you provide a series of attributes, separating by comma A_i are different attributes and for every attribute, there is a corresponding type domain type specified. So, it says that A_1 is of domain type D_1 , A_2 is of domain type D_2 and so on. And all of these attribute descriptions, are then followed by a series of integrity constraints. It is possible that a CREATE TABLE may not provide any constraint, but often you will have a number of constraints to work with. So, here is one example.

So, in this we are trying to code the creation of this instructor table, as you can see it has 4 different fields ID, name, department name and salary and for each one we have specified the domain type. So, ID is char 5, this means that the identity of this table instructor will be strings of length 5 whereas, the name or the department name are strings, but they have a maximum length 20, but they could have be of variable length where a salary is of numeric type having specification (8, 2). So, it can have 2 decimal places and be of size 8 maximum. So, this is the basic form of definition that we have for creating a table, defining a table or defining a schema.

(Refer Slide Time: 08:28)



Now, we can add a number of integrity constraints to the CREATE TABLE, 3 integrity constraints we will discuss here, one is not null, one is primary key and other third is foreign key. So, not null we specify whether a field can be null or not, primary key as we have seen will specify the attributes which form the primary key, and the foreign key will specify the attributes which reference some other table and our key in that table. So,

here is an example, in the instructor relation here, we had seen this part, the attribute what we have added here is, this not null. So, we say the name is not null which means that, in the instructor table it is not possible to insert a record, where the name of the instructor is null that is unknown, but it is possible, if the same thing is not said about department name same thing is not said about salary. So, it is possible that these could be null.

Now, we additionally say that primary key is ID. So, this field ID is a primary key and it is a property of SQL CREATE TABLE command that, if an attribute is referred as a primary key, then it cannot be not null. So, you do not need to specify that is here you do not need to write not null, because it is a primary key it will be known to be not null, because certainly we have discussed that key is the distinguishing attribute in a database table. So, it cannot be null.

So, it will not be able to distinguish similarly, we have finally we have the third integrity constant which is foreign key which says that, it is referencing this table department and the foreign key of this is here, the dep_name this particular field is a foreign key, which is a key of the department table. And so, we will be able to refer this, from this table as a foreign key and we know that it will be key in the department table. So, these are the ways to specify the integrity constraint.

(Refer Slide Time: 11:24)

And a Few More Relation Definitions

```

CREATE TABLE student (
    ID          VARCHAR(5),
    name        VARCHAR(20) NOT NULL,
    dept_name   VARCHAR(20),
    tot_cred    NUMERIC(3,0),
    PRIMARY KEY (ID),
    FOREIGN KEY (dept_name)
        REFERENCES department;
)

CREATE TABLE takes (
    ID          VARCHAR(5),
    course_id   VARCHAR(8),
    sec_id      VARCHAR(8),
    semester    VARCHAR(6),
    year        NUMERIC(4,0),
    grade       VARCHAR(2),
    PRIMARY KEY (ID, course_id, sec_id, semester, year),
    FOREIGN KEY (ID) REFERENCES student,
    FOREIGN KEY (course_id, sec_id, semester, year) REFERENCES section;
)

Note: sec_id can be dropped from primary key above, to ensure a student cannot be registered for two sections of the same course in the same semester

```

SWAYAM: NPTEL-NCOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Aug-2015

PPD

Database System Concepts - 6th Edition 06.13 ©Silberschatz, Korth and Sudarshan

So, here are couple of more examples. So, I will not go through them in detail, I will request you to take time and carefully understand them, again these are about different relations that exist here, about the student and about the courses that the students take, and in every case we have specified the set of fields, that you have in the table in the design of the schema are listed, in the CREATE TABLE the ID information about the primary key is provided and also the information about the foreign key.

Here department name is the foreign key which is mapping to this point, similar things can be observed about the takes relationship which show how students are actually taking courses. So, it relates different it has a set of fields but it has 2 kinds of foreign keys, one that relate to the student through the ID and this combination of attributes which refer to the section.

So, this is how different tables can be created using the data definition language, here is a note that you should observe that if you consider this section ID, the section ID is a part of the primary key which means that 2 records cannot be same, if they are different in the section ID, then such records are allowed. So which means that it is possible that, a student can attend or take a course in the same semester, in the same year with 2 different section IDs because they are primary keys.

So, they can be different. So, if we drop this from the primary key then we will enforce the condition that, no student will be able to take a course, in 2 sections in the same semester and the same year. So, these are the different design choices that we have and we will move on, here is one more example trying to show you the CREATE TABLE command for the course relation, that we have in the university database.

(Refer Slide Time: 13:59)

The slide shows a database schema diagram with the following tables and their relationships:

- course**: course_id (varchar(8)), title (varchar(50)), dept_name (varchar(20)), credits (numeric(2,0)). Primary key is course_id, foreign key dept_name references department.
- student**: ID (int), name (varchar(20)), dept_id (int).
- department**: dept_name (varchar(20)), building (varchar(20)), budget (float).
- instructor**: ID (int), name (varchar(20)), dept_name (varchar(20)), salary (float).
- proj**: course_id (int), name (varchar(20)), instr_id (int).
- section**: course_id (int), room_no (int), time_slot_id (int).
- time_slot**: section_id (int), day (char(1)), start_time (time), end_time (time).
- classroom**: section_id (int), building (varchar(20)), room_no (int), capacity (int).
- teacher**: section_id (int), instr_id (int), name (varchar(20)).
- enroll**: student_id (int), course_id (int), grade (float).

Relationships include: student referencing department; course referencing department; section referencing time_slot; section referencing classroom; teacher referencing section; enroll referencing student and course.

PPD
And more still
SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur.. Date: Apr-2018
Database System Concepts - 6th Edition
06.14
©Giberschatz, Korth and Sudarshan

(Refer Slide Time: 14:11)

The slide lists several database update operations:

- Insert**: insert into instructor values ('10211', 'Smith', 'Biology', 66000);
- Delete**: Remove all tuples from the student relation
delete from student
- Drop Table**: drop table r
- Alter**:
 - alter table r add A D
where A is the name of the attribute to be added to relation r and D is the domain of A
 - All existing tuples in the relation are assigned null as the value for the new attribute
 - alter table r drop A
where A is the name of an attribute of relation r
 - Dropping of attributes not supported by many databases

PPD
Updates to tables
SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur.. Date: Apr-2018
Database System Concepts - 6th Edition
06.15
©Giberschatz, Korth and Sudarshan

Moving on, let us look at how to update or actually put in different records, in a table which has already been created, the basic command is insert, and the keywords for that is INSERT INTO and values, in between you write the name of the relation, where the record will have to be inserted and then the values, will have to be put as a tuple in the same order in which you have defined the attributes of that relation, and certainly each of the values like this is ID value, next is the name value, the department, the salary, each one of them should be from the same domain type, as has been specified during the CREATE TABLE come on.

So, these things will have to remember. And so, every record will get inserted through one insert command, similarly a deletion can be done by DELETE FROM students, if you do DELETE FROM students without specifying which record you want to delete, basically all records will get deleted. We will see how selective deletion will happen, that will come on later. DROP TABLE is a command to remove a table, a table that has been created can be removed from the database altogether, by doing DROP TABLE and the relation name you can also change the schema of a table by using ALTER TABLE.

So, the form is alter table is a key words, you can add a new attribute to relation are by writing the name of the attribute and the domain of the attribute one after the other. Similarly, it is possible also to drop an attribute that already exists and the syntax for that will be, ALTER TABLE the relation name drop is a keyword and the name of the attribute, mind you all database systems may not allow you to drop an attribute to ALTER TABLE to remove attributes. And so, it works in some and it does not work in the rest.

(Refer Slide Time: 16:54)

Basic Query Structure

- A typical SQL query has the form:

```
select A1, A2, ..., An
from R1, R2, ..., Rm
where P
```

- A_i represents an attribute
- R_j represents a relation
- P is a predicate

A ₁	A ₂	A _n

UNIVARIATE METRIC-MOOC Instructor: Prof. P. P. Datta, IIT Kharagpur - Jain-Agarwal, 2015
Database System Concepts - 6th Edition
06.17
©Übersetzung, Korith und Sudarshan

Now, that was about the definition of the table and the basic definition of the data. So, now we will get into the basics query structure which is with tables with existing data, how do I query and find out different information.

So, the structure of an SQL query and this you should observe very carefully is normally said to be, SELECT FROM WHERE colloquially will often say, let us have a SELECT

FROM WHERE. So, it has 3 keywords select which is followed by a set of; this is a set of attributes. So, this specifies that, when a select query runs it will finally give us a new relation and in that relation, the attributes that will be available are the attributes that feature in the select list. The next clause or the next keyword in this is from, which specifies a set of existing relations.

So, r_1, r_2, r_m represent different relations and these are the relations, which will be used to actually find the information extract the information. Finally, the where clause has a predicate as a condition, which specify that what condition has to be satisfied. So, that certain peoples from the relations r_1 to r_m , will be chosen and put in this new selected result, table in terms of the attributes A_1 to A_n .

(Refer Slide Time: 18:37)

The select Clause

- The **select** clause lists the attributes desired in the result of a query
 - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:

```
select name
from instructor
```

- NOTE: SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)
 - E.g., `Name` \equiv `NAME` \equiv `name`
 - Some people use upper case wherever we use bold font

SWAYAMPRABHA MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018
Database System Concepts - 8th Edition
Silberschatz, Korth and Sudarshan

So, this is the basic understanding of the structure of the SQL query and naturally as I mentioned that will result in a relation. Now, we will go over each and every clause carefully, the select clause as I said will list all the attributes. So, it is like a projection (**Π**) in terms of the relational algebra that we have done. So, if we write `select name from instructor` then this will result in finding the names of all instructors from the instructor table, because this is, this you know is a relation, because it is happening it is a featuring in the from clause and in select, we are saying that the attribute that we want to select is the attribute name.

So, it will the instructor table has 4 attributes ID, name, dept. name and salary from that it will simply take the name of the instructor and list that in the output table. So, the basic form of selection that happens and this point you may also note, that in SQL everything is case insensitive, it does not matter whether you write in upper case or lower case. So, you can choose the style that you prefer to use.

(Refer Slide Time: 19:51)

The slide is titled "The select Clause (Cont.)" in red. It contains the following text and code snippets:

- SQL allows duplicates in relations as well as in query results
- To force the elimination of duplicates, insert the keyword **distinct** after select
- Find the department names of all instructors, and remove duplicates


```
select distinct dept_name
from instructor
```
- The keyword **all** specifies that duplicates should not be removed


```
select all dept_name
from instructor
```

At the bottom left, there is a video thumbnail of a man speaking, and at the bottom right, there is a navigation bar with icons.

This is a very important factor, that you should keep in mind that we said, while introducing relational algebra, that in the relational algebra every relation is a set and which means that according to set theory. We cannot have 2 tuples in the same relation, which are identical in all its values, because set theory does not allow that, but please keep in mind that SQL actually allows duplicates in relations. So, it is possible that in the same relation in the same table, I may have more than one record which are identical in all the fields, in all the attributes of that table and this will have lot of consequences and we will see how often, this property will have to be used.

So, if you want a typical set theoretic kind of output, that is if you want the result to be distinct, all records to be distinct, then you have to explicitly say that you want distinct values to be selected. So, all that you are doing is, you are after select and before the attribute name, you introduce another keyword distinct.

So, select distinct dept name from instructor will actually, select the departments of all instructors and quite why if it just selects department name of all instructor, then it is

quite possible that the same department name will appear number of times, because every department has multiple instructors. But when we use distinct, then every name will feature only once in that selection, then you can also specify another keyword all, which ensures that the duplicates are not removed. So, if you do select all depth name, then all the names will feature with duplicates, if some department had 3 instructors the name of that department will feature thrice.

(Refer Slide Time: 22:05)

The slide title is "The select Clause (Cont.)". It features a small sailboat icon in the top left corner and a video camera icon in the bottom right corner. The main content is a bulleted list of points about the SELECT clause:

- An asterisk in the select clause denotes "all attributes"
`select *
from instructor`
- An attribute can be a literal with no `from` clause
`select '437'`
 - Results is a table with one column and a single row with value "437"
 - Can give the column a name using:
`select '437' as FOO`
- An attribute can be a literal with `from` clause
`select 'A'
from instructor`
 - Result is a table with one column and N rows (number of tuples in the `instructor` table), each row with value "A"

On the left side of the slide, there is a vertical sidebar with the text "SWAYAM-NIDHI-NOOC-MOOCs Instructor: Prof. P. Datta, IIT Kanpur - Jan-Apr-2018". At the bottom left, there is a small video thumbnail showing a person speaking. At the bottom right, it says "Database System Concepts - 6th Edition" and "06.20". The footer of the slide contains the copyright notice "©Giberschutz, Korth and Sudarshan".

You can use an asterisk after select, to specify that you are interested in all the attributes that the relation or the collection of relations in the from clause has. You can also specify a select with a literal and without a from clause, if you do that then it will simply return you a table with a single row having that literal value and you can also rename that table, using what is known as clause as a command. So, this will give you a table Foo, where there is only one row and that row has an entry 437. You can use that, for other purposes also you can do a select of a literal, from a table with using a from clause where in, you will get a single column table, where as many is as there are records in the instructor will be produced.

(Refer Slide Time: 23:16)

The slide title is "The select Clause (Cont.)". The content discusses the use of arithmetic expressions in the select clause, mentioning division by 12 and renaming the result. A query example is shown:

```
select ID, name, salary/12
from instructor
```

The slide also includes a small video player showing a speaker and navigation icons.

Select clause can also use basic arithmetic operations for example, here we are showing a select where the third attribute as you can see, the third attribute is salary by 12, assuming that the instructor table has a salary number which is annual salary by 12 naturally you give you the monthly salary. So, those such arithmetic choices can also be made. you can also rename that field that particular salary by 12 field by a new name, as I said as can be used to rename. So, if we if you use that then, when you get the output you will get the column names are ID, name and monthly salary and in monthly salary you will actually have a computation, which is salary by 12 and in the same way, you can use multiple different kinds of arithmetic operators.

(Refer Slide Time: 24:16)

The slide has a header 'The where Clause' with a sailboat icon. The content includes a bulleted list, code snippets, and a footer with navigation icons.

- The **where** clause specifies conditions that the result must satisfy
 - Corresponds to the selection predicate of the relational algebra
- To find all instructors in Comp. Sci. dept

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```
- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
 - To find all instructors in Comp. Sci. dept with salary > 80000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and salary > 80000
```
- Comparisons can be applied to results of arithmetic expressions

Navigation icons at the bottom include arrows, a magnifying glass, and other symbols.

Now, we come to the where clause, where clause specifies the condition is a predicate which corresponds to the selection predicate of relational algebra. So, it will specify some condition, here is an example if we want to find all instructors from the instructor table, who are associated with computer science department, then you can say select name from instructor and to specify that they are from the computer science department, you will specify department name is equal to computer science.

So, this will ensure that you select the records only when this condition is satisfied. So, all records for which department name is different from computer science will not be included here. You can also write predicates using the different logical connectives and or not and so on. So, here is an example, where you are finding all instructors in computer science with salary greater than 80000. So, here we have used and clause. So, only records where the department name is computer science and salary is greater than 8000 will be chosen in the result. So, and then the projection (**Π**) will be done on the name of those instructors, you can apply comparisons of arithmetic expression. So, where clause can really write different kind of things.

(Refer Slide Time: 26:04)

The slide has a header 'The from Clause' with a sailboat icon. It contains a bulleted list of points about the 'from' clause, followed by a SQL query example, and ends with a footer and navigation icons.

- The **from** clause lists the relations involved in the query
 - Corresponds to the Cartesian product operation of the relational algebra
- Find the Cartesian product *instructor* *X* *teaches*

```
select *  
from instructor, teaches
```

 - generates every possible *instructor* – *teaches* pair, with all attributes from both relations
 - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra)

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P P Das, IIT Kharagpur.. Date-Age: 2018
Database System Concepts - 6th Edition 06.23 ©Giberschatz, Korth and Sudarshan

Finally, the from clause is sets all the different relations from where you are actually looking for the records. So, it kind of corresponds to the Cartesian product of the relational algebra. So, if we want to say compute instructor Cartesian product teachers, then you can say SELECT * instructor one table comma teachers. So, this will choose records from instructed relation, as well as from teacher's relation and in all possible combined way, it will put them in the output we have used a star.

So, all fields of in instructor and all fields of teaches will be there in the output, and since some fields may have identical name like ID in instructor and there is ID in teachers, they will be qualified by the name of the relation, from can have 1 relation, 2 relation any number of relations as you require. So, this will cause the Cartesian product to be computed which may not be very useful as an independent feature, but we will see in the next module how it can give very important computations, in terms of computing joints and so on.

(Refer Slide Time: 27:37)

Cartesian Product

instructor				teaches				
ID	name	dept_name	salary	ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2009
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2010
32343	El Said	History	60000	15151	MU-199	1	Spring	2010
...	22222	PHY-101	1	Fall	2009

inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...
...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2009
...
...

Database System Concepts - 6th Edition 06.24 ©Silberschatz, Korth and Sudarshan

So, here is an example of the Cartesian product that we talked of. So, here is a instructor relation, the teacher's relation and as you can see when we have done this cartesian product, that is SELECT * from instructor comma teachers, then all fields this is an ID of the instructor, there is an ID in teachers.

So, that is also specified here qualified by the name of the relation whereas, name comes in directly because there is no attribute called name in teachers, the department name comes in directly, salary comes in, directly course ID comes in, section ID comes in, semester comes in, year comes in and so on. And the combination of all tuples in their instructor relation, against all tuples of the teacher's relation all possible combinations have come in this result, which eventually is a cartesian product of the relational algebra.

(Refer Slide Time: 28:50)

Module Summary

- Introduced relational query language
- Familiarized with data definition and basic query structure

SWAYAM - NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kanpur Date: June-August, 2018

Database System Concepts - 6th Edition

06.25

©Silberschatz, Korth and Sudarshan

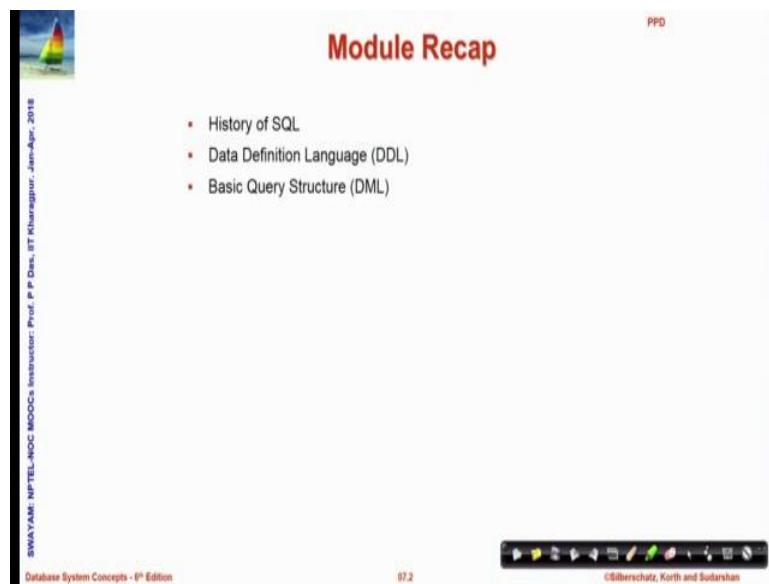
So, this is what we have to summarize, we have introduced the relational query language and particularly familiarize ourselves with the data definition that is creation of the table creation, of the schema with the attribute names, domain types and constraints. And the updates to the table in terms of insertion and deletion of values or addition or deletion of attributes or removing a table altogether and then, we have given the basic structure of the SELECT FROM WHERE query of SQL, which will be the key language feature of a query language, that we will continue to discuss all through this course.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science and Engineering
Indian Institute of Technology, Kharagpur

Lecture - 07
Introduction to SQL/2

Welcome to module 7 of database management systems, this is a second part out of total 3 of introduction to SQL.

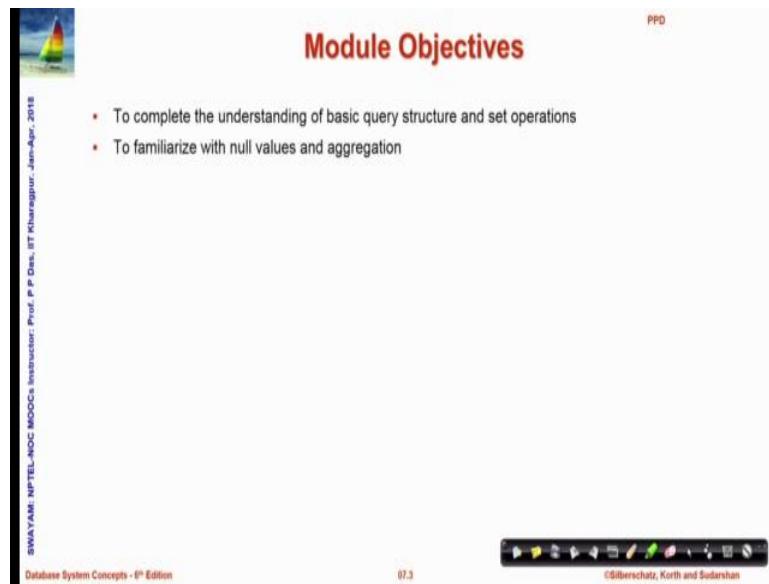
(Refer Slide Time: 00:31)



The slide is titled "Module Recap" in red at the top right. It features a small sailboat icon in the top left corner. A vertical sidebar on the left contains the text "SWAYAM: NPTEL-AHOC MOOC Instructor: Prof. P P Das, IIT Kharagpur, Jan-Apr., 2018". The main content area lists three bullet points: "History of SQL", "Data Definition Language (DDL)", and "Basic Query Structure (DML)". At the bottom, there is footer text: "Database System Concepts - 6th Edition", "07.2", and "©Silberschatz, Korth and Sudarshan". A navigation bar with various icons is at the very bottom.

In the last module, we have discussed about the evolution of SQL the data definition language part of it and the basic structure of queries.

(Refer Slide Time: 00:42)



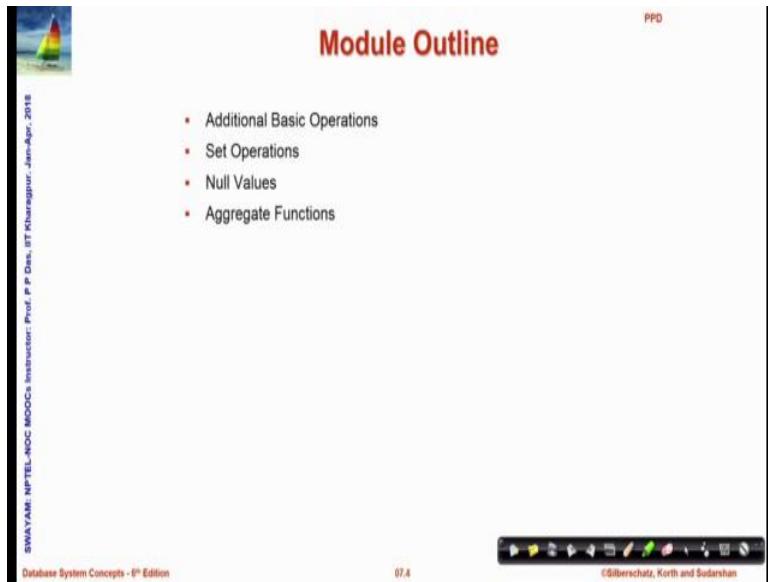
The slide is titled "Module Objectives" in red. It features a small sailboat icon in the top left corner and a progress bar at the bottom right. The content includes a list of objectives:

- To complete the understanding of basic query structure and set operations
- To familiarize with null values and aggregation

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P P Date, IIT Kharagpur, Jan-Apr., 2018
Database System Concepts - 6th Edition
07.3 ©Silberschatz, Korth and Sudarshan

In the current module we will complete the understanding of the basics queries structure and will see how, common set theoretic operations can be performed in terms of queries. We will familiarize ourselves with the handling of null values and aggregation operation that will be frequently required for forming queries.

(Refer Slide Time: 01:05)



The slide is titled "Module Outline" in red. It features a small sailboat icon in the top left corner and a progress bar at the bottom right. The content includes a list of topics:

- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P P Date, IIT Kharagpur, Jan-Apr., 2018
Database System Concepts - 6th Edition
07.4 ©Silberschatz, Korth and Sudarshan

So, this is a module outline, these are the topics, that we will discuss. So, we started the discussion of some more basic operations in the query.

(Refer Slide Time: 01:15)

The slide has a header 'Cartesian Product' with a sailboat icon. The content includes a bulleted list of points and a code snippet:

- Find the Cartesian product *instructor* X *teaches*
- ```
select *
from instructor, teaches
```
- generates every possible *instructor* – *teaches* pair, with all attributes from both relations
- For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra)

At the bottom, there is a video player showing a person speaking, the text 'Database System Concepts - 6th Edition', the number '07.6', and the copyright '©Übersetzung, Korth und Sudarshan'.

We have already discussed this that, if we do SELECT \* FROM 2 tables then it results in a Cartesian product we have seen this result earlier..

(Refer Slide Time: 01:26)

The slide has a header 'Cartesian Product' with a sailboat icon. It displays two tables side-by-side:

| instructor |             |                  |               | teaches   |                  |               |                 |             |
|------------|-------------|------------------|---------------|-----------|------------------|---------------|-----------------|-------------|
| <i>ID</i>  | <i>name</i> | <i>dept_name</i> | <i>salary</i> | <i>ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> |
| 10101      | Srinivasan  | Comp. Sci.       | 65000         | 10101     | CS-101           | 1             | Fall            | 2009        |
| 12121      | Wu          | Finance          | 90000         | 10101     | CS-315           | 1             | Spring          | 2010        |
| 15151      | Mozart      | Music            | 40000         | 10101     | CS-347           | 1             | Fall            | 2009        |
| 22222      | Einstein    | Physics          | 95000         | 12121     | FIN-201          | 1             | Spring          | 2010        |
| 32343      | El Said     | History          | 60000         | 15151     | MU-199           | 1             | Spring          | 2010        |
| ...        | ...         | ...              | ...           | 22222     | PHY-101          | 1             | Fall            | 2009        |

| <i>inst_ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> | <i>teaches.ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> |
|----------------|-------------|------------------|---------------|-------------------|------------------|---------------|-----------------|-------------|
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 10101             | CS-101           | 1             | Fall            | 2009        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 10101             | CS-315           | 1             | Spring          | 2010        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 10101             | CS-347           | 1             | Fall            | 2009        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 12121             | FIN-201          | 1             | Spring          | 2010        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 15151             | MU-199           | 1             | Spring          | 2010        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 22222             | PHY-101          | 1             | Fall            | 2009        |
| ...            | ...         | ...              | ...           | ...               | ...              | ...           | ...             | ...         |
| ...            | ...         | ...              | ...           | ...               | ...              | ...           | ...             | ...         |
| 12121          | Wu          | Finance          | 90000         | 10101             | CS-101           | 1             | Fall            | 2009        |
| 12121          | Wu          | Finance          | 90000         | 10101             | CS-315           | 1             | Spring          | 2010        |
| 12121          | Wu          | Finance          | 90000         | 10101             | CS-347           | 1             | Fall            | 2009        |
| 12121          | Wu          | Finance          | 90000         | 12121             | FIN-201          | 1             | Spring          | 2010        |
| 12121          | Wu          | Finance          | 90000         | 15151             | MU-199           | 1             | Spring          | 2010        |
| 12121          | Wu          | Finance          | 90000         | 22222             | PHY-101          | 1             | Fall            | 2009        |
| ...            | ...         | ...              | ...           | ...               | ...              | ...           | ...             | ...         |
| ...            | ...         | ...              | ...           | ...               | ...              | ...           | ...             | ...         |

At the bottom, there is a video player showing a person speaking, the text 'Database System Concepts - 6th Edition', the number '07.7', and the copyright '©Übersetzung, Korth und Sudarshan'.

(Refer Slide Time: 01:32)

The slide shows a SQL query being executed:

```

 * Find the names of all instructors who have taught some course and the course_id
 select name, course_id
 from instructor, teaches
 where instructor.ID = teaches.ID

```

Annotations highlight the 'instructor' and 'teaches' tables in the query, and the 'ID' column in both tables.

Below the query, the results are shown as two tables:

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 12121 | Wu         | Finance    | 90000  |
| 15151 | Mozart     | Music      | 40000  |
| 22222 | Einstein   | Physics    | 95000  |
| 32343 | El Said    | History    | 60000  |

| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2009 |
| 10101 | CS-315    | 1      | Spring   | 2010 |
| 10101 | CS-347    | 1      | Fall     | 2009 |
| 12121 | FIN-201   | 1      | Spring   | 2010 |
| 15151 | MU-199    | 1      | Spring   | 2010 |
| 22222 | PHY-101   | 1      | Fall     | 2009 |

Annotations show arrows pointing from the first table to the second, indicating an 'Equi-Join, Natural Join'.

At the bottom, there is a watermark: 'BINAYAK NPTEL-NOC MOOCs Instructor: Prof. P. Dini, IIT Kharagpur - Jan-2017 - 2018'.

Know by itself as we had said that, by itself the Cartesian product may not be very useful, but suppose we want to answer this kind of a query, that find all names of all instructors who have taught some course and for those you also write the course\_id. So, what we are interested in is, the name of the instructor and course\_id.

So, we put that on the select course these are the 2 tables that are required because, the name of the instructor is there in the instructor table relation and then the relationship between which instructors teach which course is in the teachers table. So, in the instructor table we have the name in the teachers table, we have the course\_id and also this relationship as to which course is taught by which teacher and here, we have the relationship between which id of the instructor has which name. So, what we do is, when we do this Cartesian product will get something like this, as we have already seen, but we want to qualify it by with a predicate which say that, we will out of all these combinations will choose only those where, the instructor.ID equals the teaches.ID. So, if we look at say in the first row the instructor.ID is this, teachers.ID is this, which is id of the teachers table they are same. So, which says the this instructors Srinivasan actually teaches the course CS-101 whereas, if you look into this row, it says that instructor.ID is this and teaches.ID is this and we are really not interested in this combination, because this combination does not convey anything meaningful. So, by the use of this where clause, we will try to choose only those records where, these 2 ids are same which will tell us that, this particular instructor actually teaches that particular course.

So, if we do that, then you will find that majority of the records that, actually came about in the Cartesian product are eliminated from this result. So, I have struck them out. So, you can currently see in this part you can see only 4 records the 3 courses taught by Srinivasan. So, and the one course taught by you. So, in the output you will have this name, this name part and this course\_id part because, you are projecting on these 2. So, this will be the output table which will get generated and just to remind you, this is the notion of natural join that, we had discussed in relational algebra in this case we will actually call it equi join because, we are using a equality condition after the Cartesian product to join these 2 relationship. So, this is a very critical operation in many of cases in our database query system.

(Refer Slide Time: 05:23)

The screenshot shows a presentation slide with the following details:

- Title:** Examples
- Content:**
  - A bullet point: "Find the names of all instructors in the Art department who have taught some course and the course\_id"
  - Below the bullet point is a SQL query:

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID and instructor.dept_name = 'Art'
```
- Navigation:** A toolbar at the bottom includes icons for back, forward, search, and other presentation controls.
- Page Number:** 07.8
- Credit:** ©Silberschatz, Korth and Sudarshan
- Source:** SWAYAM-NPTEL-NOC's INSTRUCTOR: Prof. P. R. Dand, IIT Kharagpur. -Jan-Apr- 2018

This is another extension of this similar example. So, here we have added another predicate in the where clause, specifying that instructed or department name is Art.

So, which means this will now, give the names of all instructors in the art department only who have taught some course and specify their course\_id. So, in different such ways, you can manipulate and create queries.

(Refer Slide Time: 05:57)

The slide has a header 'The Rename Operation' with a sailboat icon. It contains the following text and code:

- \* The SQL allows renaming relations and attributes using the **as** clause:  
 $old-name \text{ as } new-name$
- \* Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.  

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```
- \* Keyword **as** is optional and may be omitted  
 $instructor \text{ as } T \equiv \text{instructor } T$

At the bottom left is a video thumbnail of a speaker, and at the bottom right are navigation icons.

It is possible to rename, we have already seen examples you can rename a relation you can rename an attribute and the style is to use AS. So, here you can see that, in the SELECT query we have said from instructor as T. So, the name of this relation can be treated as T and we again see that as instructor as S. So, actually what we are doing, we are doing a join between the same relation; instructor and instructor.

And we are trying to find out all instructor who have higher salary than some instructor in computer science. So, the some instructor in computer science is specified by this condition, because department has to be computer science and the fact that salary is higher. So, as if you treat that, though it is actually a join between a product between instructor and instructor the same relation, but by renaming you treat them as if they are 2 different tables having name T and S, and then it becomes easier to write this kind of query. So, otherwise it is quite difficult to write this query to find out, because you need to actually create a product of one relation with itself.

(Refer Slide Time: 07:30)

The slide features a sailboat icon in the top left corner. The title 'Cartesian Product Example\*' is centered at the top in red. Below the title is a table with two columns: 'person' and 'supervisor'. The data is as follows:

| person | supervisor |
|--------|------------|
| Bob    | Alice      |
| Mary   | Susan      |
| Alice  | David      |
| David  | Mary       |

Below the table is a bulleted list of three items:

- Find the supervisor of "Bob"
- Find the supervisor of the supervisor of "Bob"
- Find ALL the supervisors (direct and indirect) of "Bob"

At the bottom left is a small video thumbnail of a man speaking. The bottom right shows a navigation bar with icons and the text 'Database System Concepts - 6<sup>th</sup> Edition'.

Keyword AS is optional you can just write instructor, and then the name the new name that you want to give and that itself will work here. Is another Cartesian product example, here given a relation which lists a person and his or her supervisor, we want to find out all supervisors direct or indirect of that person.

So, I leave this as an exercise to you, to think over as to how we can actually compute this query?

(Refer Slide Time: 07:56)

The slide features a sailboat icon in the top left corner. The title 'String Operations' is centered at the top in red. Below the title is a bulleted list of four items:

- SQL includes a string-matching operator for comparisons on character strings. The operator `like` uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring
  - underscore ( \_ ). The \_ character matches any character
- Find the names of all instructors whose name includes the substring 'dar'

Below the list is a handwritten note in blue ink:

select name  
from instructor  
where name like "%dar%"

At the bottom left is a small video thumbnail of a man speaking. The bottom right shows a navigation bar with icons and the text 'Database System Concepts - 6<sup>th</sup> Edition'.

SQL supports several string operations and of particular interest at 2 specific symbols, characters which allow us doing certain match, percentage is used to match any substring and underscore is used to match any particular character and we use a keyword LIKE, to find out different string patterns that can be matched. So, here we want to find the names of all instructors, whose name includes the substring “dar” and by writing this so saying the predicate is formed is named like this.

So, what is there is a percentage before there is a percentage after. So, anywhere “dar” will feature in the name, this predicate will turn out to be true otherwise if there is no “dar” in the name the predicate will turn out to be false and that particular record will not get selected. So, in this way we can using LIKE, we can actually do different kinds string operations as conditions in the weight loss.

(Refer Slide Time: 09:07)

The slide has a title 'String Operations' in red at the top right. On the left, there is a small image of a sailboat on water. The main content area contains a bulleted list of SQL string operations and a sample query. At the bottom, there is footer information and a navigation bar.

**String Operations**

- SQL includes a string-matching operator for comparisons on character strings. The operator `like` uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring
  - underscore ( \_ ). The \_ character matches any character
- Find the names of all instructors whose name includes the substring "dar"

```
select name
from instructor
where name like '%dar%'
```
- Match the string "100%"

```
like '100 \%' escape '\'
```
- in that above we use backslash (\) as the escape character

BRUNNEN-NPTEL-NOOC MOOCs Instructor: Prof. P. P. Dini, IIT Kharagpur - 2018

Database System Concepts - 6<sup>th</sup> Edition 07.12 ©Güterschütz, Korth and Sudarshan

So now, naturally this brings in an issue of for what if my string itself has a percentage or an underscore character. So, the rule followed is you will need to escape that, with the escape character that you define. This is a style which you have seen in C programming as well.

(Refer Slide Time: 09:29)

The slide features a title 'String Operations (Cont.)' at the top right. On the left, there is a small video window showing a man speaking. The main content area contains a bulleted list of string matching patterns:

- Patterns are case sensitive
- Pattern matching examples:
  - 'Intro%' matches any string beginning with "Intro"
  - '%Comp%' matches any string containing "Comp" as a substring
  - '\_\_\_\_\_' matches any string of exactly three characters
  - '\_\_\_ %' matches any string of at least three characters

At the bottom left, it says 'Database System Concepts - 8<sup>th</sup> Edition'. At the bottom center, it shows '07.13'. At the bottom right, it says '©Übersetzung, Korth und Sudarshan'.

Patterns are certainly case sensitive. So, it depends it will distinguish between upper case as well as lower case, and these are different examples of string matching that you can do where, you can match at this beginning of a string, end of a string, anywhere in the string, specific number of characters in a string and so on. SQL supports concatenation, conversion of lower to upper case and vice versa and different other common string operations, those are available as functions in SQL and can be used for convenience.

(Refer Slide Time: 10:09)

The slide features a title 'Ordering the Display of Tuples' at the top right. On the left, there is a small video window showing a man speaking. The main content area contains a bulleted list of ordering rules:

- List in alphabetic order the names of all instructors

```
select distinct name
from instructor
order by name
```
- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
  - Example: **order by name desc**
- Can sort on multiple attributes
  - Example: **order by dept\_name, name**

At the bottom left, it says 'Database System Concepts - 8<sup>th</sup> Edition'. At the bottom center, it shows '07.14'. At the bottom right, it says '©Übersetzung, Korth und Sudarshan'.

Now, let us address a different question. Let us say, we have computed a query and then often we would want that the result be ordered in according to certain order particularly the value of certain field if we want the result to be ordered, then SQL allows you to do that by another clause that you add to the query, which is called order by. So, what this will do, we have already seen this query this will find out the names of all the instructors and the names will occur in a distinct manner because, distinct is specified, but then the output will be in terms ordered by the name.

And the ordering can be descending or ascending, you can control that, by specifying whether you want descending or ascending by default the ordering is ascending. So, that makes the presentation of the result of and very easy and you can certainly sort on multiple fields as well, so it can be ordered based on combination of fields.

(Refer Slide Time: 11:15)

## Where Clause Predicates

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq \$90,000$  and  $\leq \$100,000$ )
 

```
select name
from instructor
where salary between 90000 and 100000
```
- Tuple comparison
 

```
select name, course_id
from instructor, teaches
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```

SQL where clause also allows between as a comparison parameter. So, between can specify 2 values. So that, whenever the field value will be between these 2 given values the condition will be predicated will be taken to be true otherwise is taken to be false. You can compare based on people as well. So, in this case you could have written, you could have checked for equality of instructor.ID with teachers.ID and department name with the literal biology, but you can compact it by writing a tuple notation as is shown here. So, these are common convenient ways of writing different where clause.

(Refer Slide Time: 12:02)

**Duplicates\***

- In relations with duplicates, SQL can define how many copies of tuples appear in the result!
- Multiset versions of some of the relational algebra operators – given multiset relations  $r_1$  and  $r_2$ :
  - $\sigma_\theta(r_1)$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$ , and  $t_1$  satisfies selections  $\sigma_\theta$ , then there are  $c_1$  copies of  $t_1$  in  $\sigma_\theta(r_1)$
  - $\Pi_A(r)$ : For each copy of tuple  $t_1$  in  $r_1$ , there is a copy of tuple  $\Pi_A(t_1)$  in  $\Pi_A(r_1)$  where  $\Pi_A(t_1)$  denotes the projection of the single tuple  $t_1$
  - $r_1 \times r_2$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$  and  $c_2$  copies of tuple  $t_2$  in  $r_2$ , there are  $c_1 \times c_2$  copies of the tuple  $t_1, t_2$  in  $r_1 \times r_2$

SWAMY: NPTEL-NOC MOOC Instructor Prof. P P Das, IIT Kharagpur Date: Jan-August, 2018  
Database System Concepts - 6<sup>th</sup> Edition

07.16 ©Übersetzung, Korth und Sudarshan

Now, we have specified that SQL does carry duplicates. So, unlike relational algebra which set theoretically specify that, their duplicates should not be there if in SQL there could be duplicate entries in the same relation. So, there is a; this is called when duplicates are allowed in set theory, then such sets where duplicates are allowed and known as multi sets. So, there are multi set versions of the SQL queries or so to say the relational algebra operations.

So, you have a selection, which can be multi set selection, which means that, if there are certain  $c_1$  number of copies of a tuple in the relation, which satisfy the condition theta then all of them will feature in the result. And all those copies can be seen simultaneously, because it is a multi-set condition, similar definitions are hold for projection, as well as for Cartesian product. So, I will leave it to you to go through the details and convince yourself that, these multi set relations really extend the traditional single set distinct definition of the relational algebra.

(Refer Slide Time: 13:27)

The slide has a decorative background of a sailboat on water. The title 'Duplicates (Cont.)\*' is at the top. A list of bullet points discusses multiset relations  $r_1$  and  $r_2$ , their Cartesian product, and SQL semantics. Handwritten annotations show sets  $r_1 = \{(1, a), (2, a)\}$  and  $r_2 = \{(2, 3), (3, 3)\}$ , and the result of  $\Pi_B(r_1) \times r_2$  as  $\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3)\}$ . Below is an SQL snippet and its multiset equivalent.

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Datta, IIT Kharagpur - Jan-Aug - 2018  
Database System Concepts - 6<sup>th</sup> Edition

Duplicates (Cont.)\*

- Example: Suppose multiset relations  $r_1(A, B)$  and  $r_2(C)$  are as follows:  
 $r_1 = \{(1, a), (2, a)\}$        $r_2 = \{(2, 3), (3, 3)\}$
- Then  $\Pi_B(r_1)$  would be  $\{(a), (a)\}$ , while  $\Pi_B(r_1) \times r_2$  would be  $\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3)\}$
- SQL duplicate semantics:  

```
select A1, A2, ..., An
 from r1, r2, ..., rm
 where P
```

is equivalent to the multiset version of the expression:

$$\prod_{A_1, A_2, \dots, A_n} (\sigma_P (r_1 \times r_2 \times \dots \times r_m))$$

07.17 ©Übersetzung, Korth und Sudarshan

So, here is an example where, there are 2 multi set relations as you can see particularly this one, which has identical duplicate entries. So, using that you can define a projection of  $r_1$  on  $B$ , which will certainly give you it is you are doing projection on  $B$ . So, it will give you  $A$  only. So, you will have this result itself will be a multi set because you will get 2 as. So, this result will be like  $a, a$ , and then you have  $r_2$  with which you are doing the Cartesian product. So, you will have all possible combinations all these 6 are the result in the SQL whereas, in set theory typically the result should have been only these 2 tuples.

(Refer Slide Time: 14:38)

The slide has a decorative background of a sailboat on water. The title 'SET OPERATIONS' is at the top. To the right is a list of topics: Additional Basic Operations, Set Operations, Null Values, and Aggregate Functions. Below is a video player showing a professor speaking.

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Datta, IIT Kharagpur - Jan-Aug - 2018  
Database System Concepts - 6<sup>th</sup> Edition

SET OPERATIONS

- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions

07.18 ©Übersetzung, Korth und Sudarshan

(Refer Slide Time: 14:43)

The slide has a header 'Set Operations' in red. On the left, there is a small sailboat icon. The main content consists of three bullet points with corresponding SQL queries:

- Find courses that ran in Fall 2009 or in Spring 2010  
(`select course_id from section where sem = 'Fall' and year = 2009`)  
`union`  
(`select course_id from section where sem = 'Spring' and year = 2010`)
- Find courses that ran in Fall 2009 and in Spring 2010  
(`select course_id from section where sem = 'Fall' and year = 2009`)  
`intersect`  
(`select course_id from section where sem = 'Spring' and year = 2010`)
- Find courses that ran in Fall 2009 but not in Spring 2010  
(`select course_id from section where sem = 'Fall' and year = 2009`)  
`except`  
(`select course_id from section where sem = 'Spring' and year = 2010`)

At the bottom left, it says 'SWAYAM: NPTEL-NOC MOOCs Initiator: Prof. P. P. Desai, IIT Kanpur Date: June-August, 2018'. At the bottom center, it says 'Database System Concepts - 6<sup>th</sup> Edition' and '07.19'. At the bottom right, there is a navigation bar with icons.

Now, we take a quick look into the common set operations. So, it is possible to do union, intersection, difference kind of operations very easily with SQL queries. So, suppose we want to find all courses, that ran in fall 2009 or in spring 2010. So, certainly the first part of the query is simple. This will give you all courses that run. So, you are taking out the course\_id from section is where, the course running information is provided and you part putting 2 conditions, which say that they actually this courses ran in fall 2009.

So, this is the first query the second query says the courses, that run in spring 2010 and you are you have an or condition in the statement of what you are looking for. So, you do a union, union is another keyword. So, this will simply give you a relation of the course\_id attribute as the only attribute, which has records from the first as well as the second query. Similarly, you can find out the courses that run, both in fall 2009 and spring 2010 by using intersect, which basically give you the intersection of the result of the first and the second query.

You could also do difference set difference by doing find courses, that ran in fall 2009, but not in 2010. So, what will that mean, that will mean that the result of the first query from the results of the second query be subtracted be done a difference from. So, those that, had run in the fall 2009 and then was again run in spring 2010 will get removed. So, that is done through the accept keyword.

So, in this way you can very easily do set operations, whenever that is easy to conceive; obviously, you can write these queries in several other different forms, but this is just to show you how set theoretic operations can be easily written.

(Refer Slide Time: 17:00)

The slide has a title 'Set Operations (Cont.)' at the top right. On the left, there is a small image of a sailboat on water. Below the title, there are three bullet points with corresponding SQL code:

- Find the salaries of all instructors that are less than the largest salary  

```
select distinct T.salary
from instructor as T, instructor as S
where T.salary < S.salary
```
- Find all the salaries of all instructors  

```
select distinct salary
from instructor
```
- Find the largest salary of all instructors  

```
(select "second query")
except
(select "first query")
```

At the bottom left, there is a video thumbnail of a man speaking. At the bottom center, it says '07:20'. At the bottom right, it says '©Übersetzung, Korth und Sudarshan'.

You can do set operations like this. in terms of fine salaries of all instructions that are less than a largest salary. So, again we are using renaming to think of the same relation as 2, and then as if from the relation T you are trying to look at relation S and finding out what are the salaries, which are smaller than that and certainly whatever comes out is the one which is not the largest because, certainly the largest will not satisfy this particular condition, because it will get compared with itself. You can find salaries of all instructors, and then you can find the largest salary.

So, this is all salaries which are less than largest, this is all salaries including the largest. So, what happens, if you subtract that is from this if you subtract this from all salaries, if you remove the salaries, that are not largest naturally what you get is a largest salary. So, this is a interesting way to find the largest salary we will see later on that there could be several other ways particularly the use of aggregate function, which make these computations easier to perform, but these are the typical ways to use set theoretic operations.

(Refer Slide Time: 18:29)

The slide is titled "Set Operations (Cont.)" in red. It features a small sailboat icon in the top left corner. In the bottom left, there is a video feed of a man speaking. The bottom right contains a navigation bar with icons for back, forward, search, and other presentation controls. The main content area lists points about set operations:

- Set operations **union**, **intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.
- Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:
  - $m + n$  times in  $r \text{ union all } s$
  - $\min(m,n)$  times in  $r \text{ intersect all } s$
  - $\max(0, m - n)$  times in  $r \text{ except all } s$

The set operations, so we have seen 3 of them union, intersect and except they automatically these operations are set theoretic. So, each of them automatically, eliminate the duplicate, unlike what SQL by default, SQL by default does what, allows duplicates, but set operations will eliminate duplicates, because they are set operations. So, if you want the SQL type of behaviour, if you want the duplicates to be preserved, then you can have a multi set version of these set operations, which are known as union all, intersect all, except all like that.

And naturally if you do these operations then, here is the simple formula of the number of tuples, that will get computed in different cases. you can study and convince yourself that these are the correct numbers. Let us, go to the treatment of we talked about null values, that we said that it is possible that certain records in a relation may have one or more attributes where, the value is not known and to represent, that the value is not known we are putting a placeholder called null.

(Refer Slide Time: 19:52)

The slide has a header 'Null Values' in red. To the left is a small sailboat icon. On the right is a vertical bar with a black background and white text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Doshi, IIT Kanpur Date: June-August, 2015'. Below the title is a bulleted list of points about null values. A blue arrow points from the word 'null' in the list to the word 'null' in the SQL code. Another blue arrow points from the word 'is' in the predicate 'is null' to the word 'is' in the list. At the bottom, there is a snippet of SQL code: 'select name from instructor where salary is null'. The footer of the slide includes the text 'Database System Concepts - 6<sup>th</sup> Edition', the date '07.23', and copyright information '©Übersetzung, Korth und Sudarshan'.

So, let us see what is the consequence of that null value in terms of doing these query operations. So, the null signifies an unknown value. So, if I do 5 plus null then naturally the result is null. So, what you are saying that I am adding an unknown quantity to 5. So, then what would you say is the result is unknown. So, that is the basic semantics of adding null to a number. So, it is possible to check, if particularly a field is a null for a record, and that is done by a predicate “is null”.

So, in this particular query we are trying to find all instructors, whose salary is null that is not known. So, this is a predicate. So, for a particular record for which salaries null, this will become true and that will get included in the result, but for all records for which, there is some value for the salary. So, salary is known it is not null those will not get included in the result.

(Refer Slide Time: 20:58)

## Null Values and Three Valued Logic

- Three values – *true, false, unknown*
- Any comparison with *null* returns *unknown*
  - Example:  $5 < \text{null}$  or  $\text{null} > \text{null}$  or  $\text{null} = \text{null}$
- Three-valued logic using the value *unknown*:
  - OR:  $(\text{unknown or true}) = \text{true}$ ,  
 $(\text{unknown or false}) = \text{unknown}$   
 $(\text{unknown or unknown}) = \text{unknown}$
  - AND:  $(\text{true and unknown}) = \text{unknown}$ ,  
 $(\text{false and unknown}) = \text{false}$ ,  
 $(\text{unknown and unknown}) = \text{unknown}$
  - NOT:  $(\text{not unknown}) = \text{unknown}$
  - "P is unknown" evaluates to true if predicate P evaluates to unknown

So, the basic semantics of null is then, combined with the truth values because, we know our basic predicate logic is 2 values true and false. But now, you have a third value unknown that is, you may not know the value of a predicate. So, how does it play around with the true and false values, you can reason through that quite easily if you are comparing with a null in whatever way, then naturally the result is unknown. So, it returns a null, if you are doing any connectives for example, if you are doing or of null or true, then the result should be true because, in or we say that if any of the components is true then the result is true.

So, here you do not need to know what is that unknown well you can say it is true, but if you do or with false or of unknown with false the second row or of unknown with false if you do this, then naturally this is unknown because, since this is false the result would be true only if unknown value is true and the result would be false if the unknown value is false, you do not know what that unknown value is. So, you have to say that you have result is unknown. So, using that same logic you could see verify, I would ask you to verify offline at home you please verify, that all these combinations of true false with unknown are valid. So, if P is unknown as a predicate will evaluate to true if P is not known..

(Refer Slide Time: 23:05)



## Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value
  - avg: average value
  - min: minimum value
  - max: maximum value
  - sum: sum of values
  - count: number of values

Database System Concepts - 6<sup>th</sup> Edition      07.26      ©Büterschutz, Korth and Sudarshan

Now, we come to the aggregate functions, there are several aggregate functions they can be used for convenience and these are the common months, that operate on the multi set values naturally aggregate functions operate on a particular column they try to aggregate in a particular column and return a single value for example, average would be meaning, that you are trying to find average of the values of a particular column.

(Refer Slide Time: 23:28)



## Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department

```
select avg (salary)
from instructor
where dept_name= 'Comp. Sci.';
```
- Find the total number of instructors who teach a course in the Spring 2010 semester

```
select count (distinct ID)
from teaches
where semester = 'Spring' and year = 2010;
```
- Find the number of tuples in the course relation

```
select count (*)
from course;
```



Database System Concepts - 6<sup>th</sup> Edition      07.27      ©Büterschutz, Korth and Sudarshan

So, here is an example. So, we are trying to find the average salary of instructors in computer science department. So, naturally what you output is average salary. So, mind you this will output relation will have one attribute, which is average salary and since, average salary is a single quantity it will only have one, and here I have made use of this

aggregate function average. So, it says you do average on the attribute salary and where do you get that attribute, from you get it from the table instructor and then we are saying that we are not interested to find average of salary of all instructors. We are interested to find the average salary of those instructors who work for computer science. So, you put this back clause. So, this will ensure, that you find the average salary of instructors in computer science department.

So, in similar way you can use other aggregate functions LIKE, if you want to know the total number of instructors who teach a course in this semester. So, you first put the where clause, naturally you where will you find this information you will find this information in teachers, teachers is the relation which tells you which instructor is teaching what course. So, that comes in from, then you have to specify that teaching a course in spring 2010 semester. So, the where clause specifies, that the semester is spring and the year is 2010. So, this will give you all records, which show that some instructor is teaching the course in spring 2010 semester.

Now, naturally there could be multiple the same instructor could happen multiple times, because an instructor may be teaching more than one course. So, you make the instructor.ID, instructor.ID that you have here, you make that distinct. So, that you get only those instructors, every instructor who is teaching one course or more than one course will feature only once in this total list, and then you simply count it. use aggregate function count on that.

So, that will tell you how many instructors are teaching some course in spring 2010 mind you, this is critical to use this key word distinct, because unless you use that, then all that you will eventually find out is not the number of instructors who are teaching the course you will find out the number of courses, that are being offered in spring 2010 because, there could be the same instructor teaching more than one course. If you just want to count the number of tuples you can do count on star. because, what is star is all the attributes. So, from if we want to find out the number of course, you suggest to count \* on course.

(Refer Slide Time: 27:03)

**Aggregate Functions – Group By**

- Find the average salary of instructors in each department

```
select dept_name, avg(salary) as avg_salary
from instructor
group by dept_name;
```

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 76766 | Crick      | Biology    | 72000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 12121 | Wu         | Finance    | 90000  |
| 76543 | Singh      | Finance    | 80000  |
| 32343 | El Said    | History    | 60000  |
| 58583 | Califieri  | History    | 62000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 22222 | Einstein   | Physics    | 95000  |

| dept_name  | avg_salary |
|------------|------------|
| Biology    | 72000      |
| Comp. Sci. | 77333      |
| Elec. Eng. | 80000      |
| Finance    | 85000      |
| History    | 61000      |
| Music      | 40000      |
| Physics    | 91000      |

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P P Das, IIT Kharagpur... Date: Aug-2015  
Database System Concepts - 6<sup>th</sup> Edition 07.28 ©Übersetzung, Korth und Sudarshan

So, this is showing you the computation of average salary of instructors in each department.

So now, what do you want to do is earlier you try to find out the average salary in one department. Now, you want that for all the departments for each department I want. So, my result now, is not a single row, it is not a single value it is a pair where, I show the department and the average salary in that department. So, this is what I want visible and this is what I have.

So, naturally the information comes from instructor, that is from what I want is a department name and the average salary and I want to give it a nice name avg\_salary. So, I have done a rename. So, I get avg\_salary here, but then what I want is I do not want an average d1 over this whole set of fields. I want separate average to be done here, to be done here, to be done on this to be on this. So, these are these are basically groupings by the department as you can see, that this particular relation has been sorted according to the department name.

So, when I want to do apply an aggregate function on certain subgroups of records, I use this particular clause GROUP BY, and use a name of a field. So, what it does is if the values in the GROUP BY field in this case department name are identical those records are put together and over those records and average is computed. So, the average, that is computed over these records are put in here, average that is computed in terms of these

records are put in here, there is only one record, so average that is computed in terms of that is put in here. So, GROUP BY is a very useful feature along with the aggregation functions and it allows you to club information based on certain attribute and then compute the aggregation on some other field.

(Refer Slide Time: 29:27)

The slide has a header 'Aggregation (Cont.)' with a sailboat icon. It contains a bullet point: 'Attributes in **select** clause outside of aggregate functions must appear in **group by** list'. Below this is a code snippet:

```
/* erroneous query */
select dept_name, ID, avg (salary)
from instructor
group by dept_name;
```

At the bottom left, it says 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P P Das, IIT Kharagpur - Jan-Apr-2018 Database System Concepts - 6<sup>th</sup> Edition'. At the bottom right, it shows a navigation bar with icons and the text 'Übersicht, Korth und Sudarshan'.

Mind you, you will have to do GROUP BY and create the result id result table you will have to make sure that, all your result table attribute are used in the GROUP BY, which is not an aggregate function. So, here ID is not used. So, this is not a query, that SQL would support.

(Refer Slide Time: 29:53)

The slide has a header 'Aggregate Functions – Having Clause' with a sailboat icon. It contains a bullet point: 'Find the names and average salaries of all departments whose average salary is greater than 42000'. Below it is a SQL query:

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur.. Date: Aug-18  
Database System Concepts - 6<sup>th</sup> Edition  
07:30  
©Übersetzung, Korth und Sudarshan

You can further refine your result by saying that, find names and average salary of all departments; this much you have already done.

Now, you are qualifying that, whose average salary is greater than 42000. So, of all that we have for example, if we look here for example, in this music department average salary is less than 42000. So, you do not want that in the result you want only those where, the average salary is greater than 42000 and the way to do that is to add another clause called HAVING they say that, the average salary is greater than 42000. So, you are adding another predicate for actually qualifying the aggregated value. Now, the HAVING clause actually applies after along with the GROUP BY because, naturally the HAVING relates to the grouping.

So, once the grouping has happened groups have been formed then the HAVING clause will be evaluated on that, in contrast where clause also has a predicate, but the where clause is applied before forming the groups. So, this point this note has to be understood carefully because, if you have a where clause to choose the records it will first apply, then out of those records chosen the grouping will happen and once the grouping has happened, then the aggregate function will evaluate and the HAVING clause will get evaluated, the predicate of HAVING clause will get evaluated.

(Refer Slide Time: 31:45)

The slide has a header 'Null Values and Aggregates' in red. On the left, there is a small image of a sailboat on water. The main content includes a bulleted list of points and a corresponding SQL query:

- Total all salaries

```
select sum (salary)
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount

- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
  - count returns 0
  - all other aggregates return null

At the bottom left, there is a small video thumbnail of a man speaking. At the bottom right, there is a navigation bar with icons and the text '©Übersetzung, Korth und Sudarshan'.

Certainly, if there are null values in terms of aggregates, then there is a question of what will happen. So, the general strategy is that, whenever you perform aggregation then the null values are all ignored. So, if on that field there is no value which is not null, that is if all values are null then the result is null otherwise the result is computed by ignoring the null values.

So, these are what do you have of course, if you count then, if the collection has only null values the count will return you 0, but all other aggregates will return you simply null.

(Refer Slide Time: 32:29)

**Module Summary**

- Completed the understanding of basic query structure and set operations
- Familiarized with null values and aggregation

SWAYAM - NPTEL-NOC MOOCs Instructor: Prof. P P Das, IIT Kharagpur, Date: June-August, 2018  
Database System Concepts - 6<sup>th</sup> Edition  
07.32  
©Silberschatz, Korth and Sudarshan

So, to summarize we have started the basic understanding of the basics query structure in the last module. Now, we have completed that with some more additional operations, we have understood the set theoretic operations and very importantly we have familiarized with the treatment of null values and aggregation functions particularly the GROUP BY and HAVING clauses and how do null values and aggregation interact in terms of an SQL query.

**Database Management System**  
**Prof. Partha Pratim Das**  
**Department of Computer Science and Engineering**  
**Indian Institute of Technology, Kharagpur**

**Lecture - 08**  
**Introduction to SQL/3**

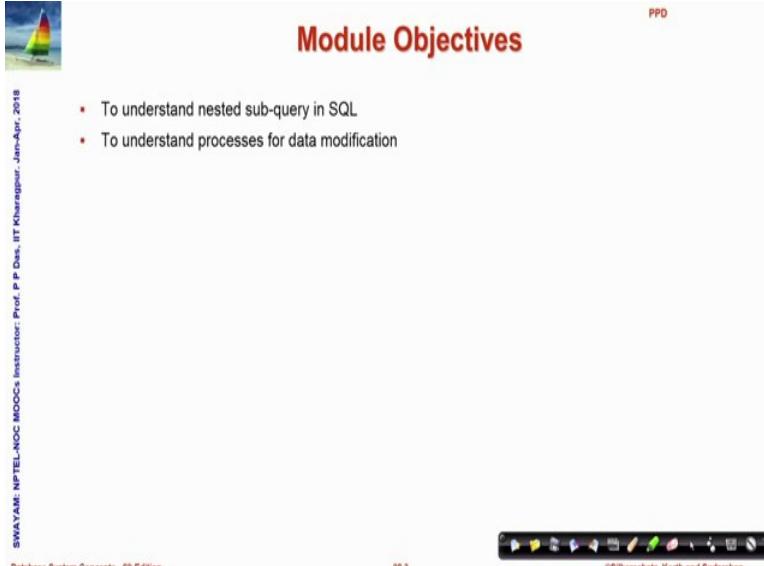
Welcome to module 8 of Database Management Systems. We have been discussing basic SQL queries and this is the third and closing part of that introductory discussion that we started in the 6th module.

(Refer Slide Time: 00:34)

The slide is titled "Module Recap" in red at the top right. It features a small sailboat icon in the top left corner. On the far left, there is vertical text: "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr., 2018". At the bottom left, it says "Database System Concepts - 8<sup>th</sup> Edition". In the center, there is a bulleted list of topics: "Additional Basic Operations", "Set Operations", "Null Values", and "Aggregate Functions". The bottom right contains a navigation bar with icons for back, forward, search, and other presentation controls, along with the text "©Silberschatz, Korth and Sudarshan". The slide is numbered "08.2" at the bottom center.

So, just to quickly recap, these are the things that we did in the last module completing the understanding of basic operations the null values and aggregate functions.

(Refer Slide Time: 00:44)



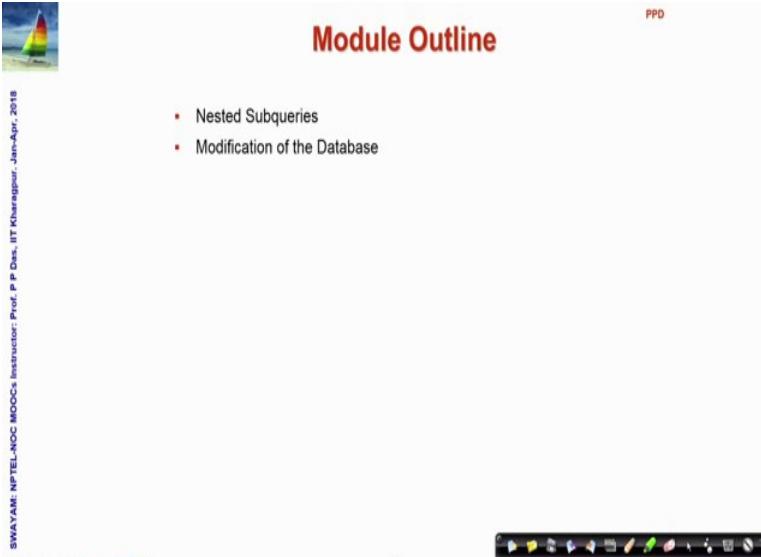
The slide is titled "Module Objectives" in red. It features a small sailboat icon in the top left corner. On the right side, there is a "PPD" watermark. The content includes a list of objectives:

- To understand nested sub-query in SQL
- To understand processes for data modification

At the bottom, there is a footer with the text "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Agr., 2018", "Database System Concepts - 8<sup>th</sup> Edition", "08.3", and "©Silberschatz, Korth and Sudarshan". A navigation bar is also present at the bottom.

In the current module, we want to understand a feature which is very very important in SQL query forming it is called the nested query or more formally nested sub query. In SQL and we would like to understand the process of data modification.

(Refer Slide Time: 01:05)



The slide is titled "Module Outline" in red. It features a small sailboat icon in the top left corner. On the right side, there is a "PPD" watermark. The content includes a list of topics:

- Nested Subqueries
- Modification of the Database

At the bottom, there is a footer with the text "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Agr., 2018", "Database System Concepts - 8<sup>th</sup> Edition", "08.4", and "©Silberschatz, Korth and Sudarshan". A navigation bar is also present at the bottom.

And those are the two things that are outlined here.

So, let us start with nested sub queries.

(Refer Slide Time: 01:12)

The slide has a header 'Nested Subqueries' with a sailboat icon. On the left, there's a vertical sidebar with text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018'. The main content area contains a bulleted list, a SQL query example, and a detailed explanation of subquery components.

- SQL provides a mechanism for the nesting of subqueries
- A **subquery** is a **select-from-where** expression that is nested within another query
- The nesting can be done in the following SQL query

```
select A1, A2, ..., An
 from r1, r2, ..., rm
 where P
```

as follows:

- A<sub>i</sub> can be replaced by a subquery that generates a single value
- r<sub>i</sub> can be replaced by any valid subquery
- P can be replaced with an expression of the form:  
B <operation> (subquery)  
where B is an attribute and <operation> to be defined later

At the bottom, there's a video player showing a man speaking, the text 'Database System Concepts - 8th Edition', the number '08.6', and the copyright '©Silberschatz, Korth and Sudarshan'.

A sub query is necessarily a SELECT FROM WHERE expression that is nested within another query, this is. So, these are the key things where expression. So, it is nothing new over; what we have already learned? But it is a part of another query it is nested within another query and that is the reason it is called a sub query. So, it by itself is not the result. this will be used, this is a SELECT FROM WHERE expression that will be used in a nested form in another some other queries to actually generate the result.

So, that is a nested sub query. So, the nesting can be done in one or more of the three clauses that a SELECT FROM WHERE SQL query has. An attribute can be replaced a relation can be any valid sub query or a sub query could form the part of a predicate in the where clause all of these are possible. So, we will discuss them one by one. So, first we will start by discussing how sub queries work in the where clause.

(Refer Slide Time: 02:44)

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018

## Subqueries in the Where Clause

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality

Database System Concepts - 8<sup>th</sup> Edition 08.8 ©Silberschatz, Korth and Sudarshan

The common use for you having sub queries in the where clause is to perform different kind of tests for membership comparison cardinality and so, on.

(Refer Slide Time: 02:58)

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018

## Set Membership

- Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

Set Member ship

Database System Concepts - 8<sup>th</sup> Edition 08.8 ©Silberschatz, Korth and Sudarshan

So, let us look at this; you have already seen this query before find courses offered in fall 2009, and in fall in spring 2010. Earlier we have shown different ways of coding this, now we are showing yet another way to be able to code this in SQL.

So, first from the beginning certainly we need the courses. So, the SELECT clause is trivial it has to be the distinct course id is certainly the information will come from

section which has information about offering of courses as we have also seen already. So, those two are no brainer. Now let us see what how do I find courses that are offered in fall 2009 and in spring 2010. So, again the first part, the courses offered in fall 2009 is coded in this part; in part of that where clause predicate when he says semester has to be fall and here is 2009. So, this part is also done.

Now, we need to ensure that whatever I mean, if I assume that after this this part were not there, then this will only give me courses which are offered in fall 2009, but we want the courses that are in fall 2009 and in spring 2010. So, we do something interesting what I do is, we write a separate query here which is courses that are offered in spring 2010. SELECT course\_Id just in the same way, SELECT course\_Id, section, semester year. So, this particular query will give me the courses offered in spring 2009. So, what do we have in one part? I have so, if you look at this part this courses that happened in fall 2009; if we look at this part courses that happened in spring 2010 good.

Now, what I want, I need that the; it I am interested in the courses that happened in both. So, for a course that exists here I want to specify that that course Id must be present here. So, what I am checking for? I am checking for a set membership; this is a set right. So, I am trying to check, whether that course Id which is being selected in the first part exists in, in is another keyword. In this particular, this particular relation that is specified by the second part of this query which is courses offered in spring 2010.

If it is, if the course Id is present, then that course must have been offered in both the semesters if it is not present, then it is offered only in fall 2009, and not in spring 2010 and certainly the courses that were not offered in fall 2009, and only offered in spring 2010 will feature here, but they do not exist here. So, they will never come up for tests. So, as a result what I get finally, is the effect of computing courses that are offered in fall 2009 and in spring 2010. this part of the query which I used as a part of the where clause is my nested sub query.

And in this case as we have seen it is used for set membership and this is a basic idea of using nested sub queries; that is a nested sub query will always give you a relation. So, you try to put that relation in the right context of the where clause from clause or SELECT clause, and then make use of it in building up your logical.

(Refer Slide Time: 07:14)

The slide features a small sailboat icon in the top left corner. The title 'Set Membership' is centered at the top in a red font. Below the title, there are two bullet points:

- Find courses offered in Fall 2009 and in Spring 2010
- Find courses offered in Fall 2009 but not in Spring 2010

For each bullet point, there is a corresponding SQL query:

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id not in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

At the bottom of the slide, there is footer text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018', 'Database System Concepts - 8<sup>th</sup> Edition', '08.9', and '©Silberschatz, Korth and Sudarshan'.

So, let us run through some examples this is, what you are saying is, earlier one was the courses offered in both here we are trying to do kind of the difference saying the courses offered in fall 2009, but not in spring 2010 certainly we easily get that by changing the membership to negation of the membership earlier it was in now you do not in you will simply get that it is up to you to take some examples and convince yourself that this kind of a nesting will work.

(Refer Slide Time: 07:43)

The slide features a small sailboat icon in the top left corner. The title 'Set Membership (Cont.)' is centered at the top in a red font. Below the title, there is one bullet point:

- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

Below the bullet point is a complex SQL query:

```
select count (distinct ID)
from takes
where (course_id, sec_id, semester, year) in
(select course_id, sec_id, semester, year
from teaches
where teaches.ID= 10101);
```

Following the query, there is a note:

▪ Note: Above query can be written in a much simpler manner.  
The formulation above is simply to illustrate SQL features.

At the bottom of the slide, there is footer text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018', 'Database System Concepts - 8<sup>th</sup> Edition', '08.10', and '©Silberschatz, Korth and Sudarshan'.

We find the set of the find the total number of distinct students, we have taken the course section taught by the instructor Id, some Id is given. So, again we form a nested query here is a nested query, which tells me the courses taught by this particular teacher 10101, and then we check set membership in terms of this course Id, section Id that is fields of the takes relation to see that, whether that particular tuple can exist in the course offered by the specific teacher; if it does then take out that I d which is which will turn out to be the student Id. In this case because that is the text relation has the student I d take out the student I d and count it as distinct. So, this can simply give you the answer to this squared; obviously, we agree that this can be written in a simpler form also, but we are including it here just for the sake of illustrating the feature.

(Refer Slide Time: 09:03)

## Set Comparison – “some” Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department
 

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept name = 'Biology';
```
- Same query using > **some** clause
 

```
select name
from instructor
where salary > some (select salary
from instructor
where dept name = 'Biology');
```

**SWAYAM: NPTEL-NOC MOOCs Instruction:** Prof. P. P. Das, IIT Kharagpur, Jan-Apr, 2018  
**Database System Concepts - 5<sup>th</sup> Edition**

There is another clause called the some clause look at this fine names of instructor; salary is greater than that of some which means at least one instructor in biology department and we have already seen this coding before.

Now, we can do this in terms of the nested query by using, again this is certainly the salary of instructors in biology department and we are doing greater than sum; that means, that the salary here being checked must find at least one record here. So that it is greater than that salary value. So, it is greater than some is a nice way to find existential records using the nested sub query.

(Refer Slide Time: 10:02)



## Definition of “some” Clause\*

▪  $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F <\text{comp}> t)$   
Where  $<\text{comp}>$  can be:  $<$ ,  $\leq$ ,  $>$ ,  $=$ ,  $\neq$

|                                                                                                                        |                                        |
|------------------------------------------------------------------------------------------------------------------------|----------------------------------------|
| $(5 < \text{some} \begin{array}{ c } \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$            | (read: 5 < some tuple in the relation) |
| $(5 < \text{some} \begin{array}{ c } \hline 0 \\ \hline 5 \\ \hline 0 \\ \hline \end{array}) = \text{false}$           |                                        |
| $(5 = \text{some} \begin{array}{ c } \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$                        |                                        |
| $(5 \neq \text{some} \begin{array}{ c } \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$ (since $0 \neq 5$ ) |                                        |

$(= \text{some}) = \text{in}$   
However,  $(\neq \text{some}) \neq \text{not in}$

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018  
Database System Concepts - 8<sup>th</sup> Edition 08.12 ©Silberschatz, Korth and Sudarshan

The logic of some clause, I have detailed out here. So, we will not go through each one of them in this discussion.

(Refer Slide Time: 10:09)



## Set Comparison – “all” Clause

▪ Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department

```
select name
from instructor
where salary > all (select salary
 from instructor
 where dept name = 'Biology');
```

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018  
Database System Concepts - 8<sup>th</sup> Edition 08.13 ©Silberschatz, Korth and Sudarshan

I leave it unto you to study and convince yourself, that you understand the semantics of some. So, similarly we have an all clause which say that if we want to say, they find the names of all instructors; whose salary is greater than the salary of all instructors in the biology department in case of sum we can write or we will write all and it will check every salary will check with the whole set of salaries in this sub query. And only if that is

greater than that particular record, that particular name will be included in the result. Otherwise it will be excluded from the result.

(Refer Slide Time: 10:52)



## Definition of "all" Clause\*

- $F <\text{comp}> \text{all } r \Leftrightarrow \forall t \in r (F <\text{comp}> t)$

|   |
|---|
| 0 |
| 5 |
| 6 |

$(5 <\text{all } \boxed{5} ) = \text{false}$

|    |
|----|
| 6  |
| 10 |

$(5 <\text{all } \boxed{10} ) = \text{true}$

|   |
|---|
| 4 |
| 5 |

$(5 = \text{all } \boxed{5} ) = \text{false}$

|   |
|---|
| 4 |
| 6 |

$(5 \neq \text{all } \boxed{6} ) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6\text{)}$

$(\neq \text{all}) = \text{not in}$   
However,  $(= \text{all}) \neq \text{in}$

↳

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018

Database System Concepts - 8<sup>th</sup> Edition      08.14      ©Silberschatz, Korth and Sudarshan

Similar to some there is a basic semantics of all which is also worked out here and I leave that to your study at home.

(Refer Slide Time: 11:03)



## Use of "exists" Clause

- Yet another way of specifying the query "Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester"
 

```
select course_id
 from section as S
 where semester = 'Fall' and year = 2009 and
 exists (select *
 from section as T
 where semester = 'Spring' and year= 2010
 and S.course_id = T.course_id);
```
- Correlation name – variable S in the outer query
- Correlated subquery – the inner query



↳

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018

Database System Concepts - 8<sup>th</sup> Edition      08.16      ©Silberschatz, Korth and Sudarshan

You can test for empty relations by using the exists construct. So, if you say exists r, then that is a predicate which mean that r is not empty; if r is empty then exist is false and not exist is the negation of exist. So, it can be used to specify the query like find all courses

taught both in fall 2009 and spring 2010. So, all that you have to do earlier you did it by set membership.

So, now you are trying to do this by this exists. So, you are saying that this is again the same query which gives you the courses that are in spring 2010 and also in this fall 2009 and you check whether this relation, whether this particular nested query is an empty one or not. if it is an empty one, then exist will fill and the whole where clause will fail, if it was not an empty one then you have found such an entry it was offered and therefore, it will get included.

So, these are just different ways of expressing similar things, but what you should note is the nested sub query is a very convenient way to frame up the logic in multiple different ways that you would like to do. So, these are the different names given the correlation name and the correlated sub query. incidentally; the nested query is often referred to as the inner query and the query in which the nesting has happened, is known as the outer query.

(Refer Slide Time: 13:04)

The slide features a title 'Use of “not exists” Clause' in red font at the top right. To the left of the title is a small icon of a sailboat on water. Below the title, there is a list of bullet points and a block of SQL code. At the bottom of the slide, there is footer information and a navigation bar.

**Use of “not exists” Clause**

▪ Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name
from student as S
where not exists ((select course_id
 from course
 where dept_name = 'Biology')
 except
 (select T.course_id
 from takes as T
 where S.ID = T.ID));
```

▪ First nested query lists all courses offered in Biology  
▪ Second nested query lists all courses a particular student took

▪ Note:  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$

▪ Note: Cannot write this query using = all and its variants

SWAYAM: NPTEL-NOCOCS Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Astr. 2018

Database System Concepts - 8<sup>th</sup> Edition 08.17 ©Silberschatz, Korth and Sudarshan

Here is another example which illustrate the use of not exist; so which I leave it for your own study.

(Refer Slide Time: 13:15)

The slide has a header 'Test for Absence of Duplicate Tuples' with a sailboat icon. It contains a bulleted list of three items, followed by a SQL query with a nested subquery highlighted in a blue box.

- The **unique** construct tests whether a subquery has any duplicate tuples in its result
- The **unique** construct evaluates to "true" if a given subquery contains no duplicates
- Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique
 select R.course_id
 from section as R
 where T.course_id= R.course_id
 and R.year = 2009);
```

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr., 2018

Database System Concepts - 8<sup>th</sup> Edition 08.18 ©Silberschatz, Korth and Sudarshan

We can check for uniqueness that is; test for absence of duplicate tuples by using the unique keyword. So, we can see here that here is a nested query and using unique to find out all courses that were offered at most once in 2009. So, if a course was offered more than once, then naturally multiple records will feature and the result the unique will fail. unique will be true only if; there is only one entry which shows that it is offered at most once in that semester.

Now, I move on. So, we have been discussing about sub queries in the where clause; now we move on to sub queries in the from clause.

(Refer Slide Time: 14:09)

The slide has a header 'Subqueries in the From Clause'. It contains a bulleted list:

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000

A SQL query is shown:

```
select dept_name, avg_salary
from (select dept_name, avg(salary) as avg_salary
 from instructor
 group by dept_name)
 where avg_salary > 42000;
```

Handwritten annotations include:

- A blue box encloses the subquery part of the query.
- A blue bracket on the right side groups the entire subquery with the outer query, labeled '(dept\_name, avg\_salary)'.

At the bottom left, there is a small video player showing a person speaking, with the text 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018'. At the bottom right, it says 'Database System Concepts - 8<sup>th</sup> Edition', '08.20', and '©Silberschatz, Korth and Sudarshan'.

So, as we have already seen a nested sub query is a relation. So, it can naturally be used in the place of any relation that we have in the from clause. So, we are trying to find out average instructor salaries of those departments where the average salary is greater than 42,000. So, look at. this is a nested sub query. So, where what is been found here this will compute the average salary department wise average salary which we have already seen group by department name and then you do average on the salary and you give it that field a new name.

So which means that; this is equivalent to having a relation which has two attributes depth name and avg\_salary. So, from that you are now trying to do the selection and what is the condition that the average salary has to be written. So, you already have that as a part of the field the average salary. So, you just need to put that in the where clause and you have only those coming out of this particular relation where average salary is greater than 42,000 to be selected in this SELECT query. So, these will feature in the output of this selection.

(Refer Slide Time: 15:56)

The slide features a sailboat icon in the top left corner. The title 'Subqueries in the From Clause' is centered at the top in red. Below the title is a bulleted list of points. Two SQL code snippets are shown, each with a blue box highlighting a specific part of the query. Handwritten-style arrows point from the highlighted parts to the explanatory text in the list. The bottom right corner contains a navigation bar with icons.

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apr., 2018

Database System Concepts - 8<sup>th</sup> Edition

08.20

©Silberschatz, Korth and Sudarshan

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000
  - ```
select dept_name, avg_salary
      from (select dept_name, avg(salary) as avg_salary
            from instructor
            group by dept_name)
           where avg_salary > 42000;
```
- Note that we do not need to use the **having** clause
- Another way to write above query
 - ```
select dept_name, avg_salary
 from (select dept_name, avg(salary)
 from instructor
 group by dept_name) as dept_avg(dept_name, avg_salary)
 where avg_salary > 42000;
```

So, that is how you can very easily use a nested sub query in the; from clause and for this we did earlier. We solve this problem using the having clause, but they here we will not need we did not need the having clause to do this there is this is another way. So, here what we have done is the same; if we if we look into the nested sub query. This is actually the same. all that we have done we have given it a new name by the renaming feature, and then this as if becomes a relation and on that the computation is done rest of it is similar.

(Refer Slide Time: 16:50)

The slide features a sailboat icon in the top left corner. The title 'With Clause' is centered at the top in red. Below the title is a bulleted list of points. A SQL code snippet is shown with a blue box highlighting a part of the query. Handwritten-style arrows point from the highlighted part to the explanatory text in the list. The bottom right corner contains a navigation bar with icons.

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apr., 2018

Database System Concepts - 8<sup>th</sup> Edition

08.21

©Silberschatz, Korth and Sudarshan

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs
- Find all departments with the maximum budget
  - ```
with max_budget(value) as
          (select max(budget)
           from department)
      select department.name
        from department, max_budget
       where department.budget = max_budget.value;
```

There is a with clause that provides a way of computing a temporary relation and that can be subsequently used.

So, let us look at example. we are trying to find all departments with maximum budget. So, this is my basic query, we want to find department name, department dot name, that is; a departments name from the department table and the budget must be same as maximum budget. So, for that I need to know the maximum budget that exists across the department. So, look at what has been done, here we have a nested query which aggregates the maximum budget from the departments.

So, this gives you the value of the maximum budget. We make that into a temporary table max budget with an attribute value. So, this is renaming. So, you cannot see the renaming is being used in very interesting ways. So, this is my nested query that gives me a relation and this is my definition of the relation. So, max budget now is a temporary relation a relation that I used subsequently in my from clause and with allows me to do that this relation will not be available.

Otherwise after this query this relation will not exist this is just a temporary one computed for this query. So, this gives me the budget value, this gives me the department and department specific budget, and this condition tells me that I can choose all the departments which has the maximum budget very nice way of using this.

(Refer Slide Time: 19:06)

The slide title is "Complex Queries using With Clause*" in red. On the left, there is a small image of a sailboat on water. Below the title, there is a bulleted list:

- Find all departments where the total salary is greater than the average of the total salary at all departments

Below the list is a block of SQL code:

```
with dept_total (dept_name, value) as
  (select dept_name, sum(salary)
   from instructor
   group by dept_name),
dept_total_avg(value) as
  (select avg(value)
   from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value > dept_total_avg.value;
```

At the bottom of the slide, there is a video player showing a person speaking, and some navigation icons.

So, with clause can be used in even more involved way again this is an example which is more complex use and I leave it to you to practice study and understand. we move on to sub queries in the SELECT clause finally.

(Refer Slide Time: 19:26)

The slide has a decorative header with a sailboat icon and the title 'Scalar Subquery' in red. The content includes a list of bullet points and a SQL query. The footer contains a photo of the speaker, the course name 'Database System Concepts - 8th Edition', the date 'Sat Apr 14 2018', and the copyright notice '©Silberschatz, Korth and Sudarshan'.

- Scalar subquery is one which is used where a single value is expected
- List all departments along with the number of instructors in each department


```
select dept_name,
        (select count(*)
         from instructor
         where department.dept_name = instructor.dept_name)
        as num_instructors
        from department;
```
- Runtime error if subquery returns more than one result tuple

A scalar sub query is one; where there is a single value is expected. So, we can very easily use that in the SELECT. So, what if you look at this part which is the sub query? So, you are saying list all departments along with number of instructors each department has. So, this condition tells that, the from the instructor; we are taking out those that department name where the instructor works to count them and then you form that as a new attribute mind you in while we were using this in the from clause.

We were treating that as a relation because nested query will give a relation, but here in the SELECT clause the entities are attributes. So, this as is renaming of attribute which means, but this is a relation that is why this notion of scalar sub query is required, that is though this is a relation what does the relation compute it computes a single value. So, that value we are putting as an attribute named num instructors.

So, we have department name and the number of instructor, then for each and every department; that we actually have from the department list. So, it is a very interesting way of using this nested sub query in terms of the SELECT clause. naturally; since in the SELECT clause, I cannot have, I mean every entry in the SELECT clause has to be an attribute pure relations are not possible.

So, if the sub query returns more than one table which cannot be conceived as one or more attributes, then it will be runtime error that will not be allowed; because we do not know how to handle multiple tuples in terms of a select clause. ok. Next we move on to discussing the modifications to the database, how do we modify the database?

(Refer Slide Time: 22:12)

The slide features a small sailboat icon in the top left corner. The title 'Modification of the Database' is centered at the top in red. On the left edge, there is vertical text: 'SWAYAM/NFTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018'. In the center, there is a video frame showing a man with glasses and a dark suit. Below the video frame, the text 'Database System Concepts - 8th Edition' is visible. To the right of the video frame, there is a navigation bar with icons and the text '08.26' and '©Silberschatz, Korth and Sudarshan'.

- Deletion of tuples from a given relation
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation

So, we will look into some of the ways of changing the records or removing records from that earlier. We saw a delete of record which removed all records from a relation, but now we will see selective deletion insertion and update of values.

(Refer Slide Time: 22:33)

The slide features a small sailboat icon in the top left corner. The title 'Deletion' is centered at the top in red. On the left edge, there is vertical text: 'SWAYAM/NFTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018'. In the center, there is a video frame showing a man with glasses and a dark suit. Below the video frame, the text 'Database System Concepts - 8th Edition' is visible. To the right of the video frame, there is a navigation bar with icons and the text '08.27' and '©Silberschatz, Korth and Sudarshan'.

- Delete all instructors
`delete from instructor`
- Delete all instructors from the Finance department
`delete from instructor
where dept_name= 'Finance';`
- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building
`delete from instructor
where dept_name in (select dept_name
from department
where building = 'Watson')`

Now, deleting all instructors are easy DELETE FROM instructor all records sorry this and this becomes an empty table, but suppose we want to delete all instructors from the finance department, then like we do in the SELECT FROM WHERE we again use the where clause as a predicate and say that DELETE FROM instructor, but you do the deletion provided this condition is satisfied that is; department name is same as finance. So, it is very similar to the SELECT FROM WHERE, but the effect is unlike SELECT FROM WHERE, where no tables change in the database here.

The table is actually changing; because these instructor records are deleted whose department name was finance. The third example shows delete all tuples in the instructor relation for those instructors, associated with the department located in the Watson building. So, Watson building may have multiple departments. So, all instructors who worked on those departments which are located in the Watson building that will go. So, you do, this is again you are using nested query.

Now, you know how to use a nested query. So, you use nested query which will give you the names of departments, which gives you a relation with a single attribute with names of departments housed in the Watson building, then you use the set membership to check whether a particular department belongs to that set, if it does then it is in Watson building; otherwise, it is not in Watson building if it does belong to the Watson building then this where clause becomes true and the corresponding instructor record is deleted and that is of this whole different kinds of selective deletion can happen.

(Refer Slide Time: 24:34)

The slide has a header 'Deletion (Cont.)'. On the left, there is a small video window showing a professor. The main content area contains the following text and code:

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary)  
from instructor);
```

- **Problem:** as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
 1. First, compute `avg (salary)` and find all tuples to delete
 2. Next, delete all tuples found above (without recomputing `avg` or retesting the tuples)

At the bottom, it says 'Database System Concepts - 8th Edition' and '08.28'.

Delete all instructors whose salary is less than the average salary of instructor. Again this is; so, you compute the selection sub query which computes the average salary and check if the salary is less than the average salary and delete that. sounds straightforward, but just wait, just wait, I mean did we do it do a right thing an average salary is computed by taking the sum of all salaries in the relation. And then dividing it by the number of relations this has to be the average salary certainly if I remove a record then the average itself will change.

So, if I write the query in this manner then what I am saying on the face of it looks correct, but then actually can it be correct because the moment a condition is satisfied and that record is deleted this average value itself has changed. So, that is not. So, that will depend then the result will depend on the order in which the deletion is happening, but that is not what was meant what was meant is take all the records for the present find out the average find out all records which have a salary less than that average and remove them.

So, this initially you know easy trivial looking solution is not actually correct. So, you will have to do the solution in two stages: first compute the average salary, find all tuples to delete? Next delete all tuples found above without re-computing. The average in the present solution the average is recomputed, which is the wrong thing.

(Refer Slide Time: 26:32)

The slide has a header 'Insertion' in red. On the left, there is a small sailboat icon and some vertical text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018'. The main content includes a bulleted list and two SQL code snippets:

- Add a new tuple to *course*

```
insert into course  
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```
- or equivalently

```
insert into course (course_id, title, dept_name, credits)  
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```
- Add a new tuple to *student* with *tot_creds* set to null

```
insert into student  
values ('3003', 'Green', 'Finance', null);
```

At the bottom, there is footer text: 'Database System Concepts - 8th Edition', '08.29', and '©Silberschatz, Korth and Sudarshan'.

So, again I will leave that for you to solve. we move on to looking at modifications in terms of insertion. So, we had seen this earlier we can add a new tuple by inserting to then the relation names, and then you save values and the tuple of values. We can specify the; if we if we do not remember the order of attributes in the relation then we can also specify the order in which we are spaced actually giving the information.

So, you are saying INSERT INTO course and what we have done here is we have actually specified the order in which that tributes occurred and that order and the order of values must be the same. In the first case, this order of values is decided by the order of the attributes that exist in terms of the create table. Add a new tuple to student with total credit set to null that is I do not know; if we were adding a student initially does not have a credit right. The credit is a nullable field the credit will be earned after the student has gone through the courses and all that. So, if I do not know the value of a field then I can write null which is a special value designating unknown for at the time of insertion.

(Refer Slide Time: 28:00)



Insertion (Cont.)

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apr- 2018

- Add all instructors to the *student* relation with tot_creds set to 0

```
insert into student
select ID, name, dept_name, 0
from instructor
```
- The **select from where** statement is evaluated fully before any of its results are inserted into the relation
- Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem



Database System Concepts - 8th Edition 08.30 ©Silberschatz, Korth and Sudarshan

And all instructors to the student relation with total credit set to 0. So, I can also combine insert with select. So, we are taking the first part this part is selection which generates a whole lot of records having ID, name, department name and the total credit set to 0. From the instructor and insert them into the students. So, these will get inserted into the students SELECT FROM WHERE statement is evaluated fully. So, this first select from will be done before any of its results are inserted in the relation. So, that is the basic condition that SQL guarantees; because, if that were not the case then such situations will become circular and will cause problem.

(Refer Slide Time: 29:03)



Updates

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jam-Apr- 2018

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
 - Write two **update** statements:

```
update instructor
set salary = salary * 1.03
where salary > 100000;
update instructor
set salary = salary * 1.05
where salary <= 100000;
```
 - The order is important
 - Can be done better using the **case** statement (next slide)



Database System Concepts - 8th Edition 08.31 ©Silberschatz, Korth and Sudarshan

Updates can be done based on particular values. So, you can update a table and what it means that? It you could update the values of specific fields.

So, here in the first case; we are giving trying to give a 3 percent salary raise for salaries which are more than 100,000 and some 5 percent raise for salaries which are less than equal to 100,000 and mind you this order in which you do the update is important, because if you do the later update first then someone who was what qualified in the later part was less than 100,000 with the increase will become more than 100,000 and will also qualify for the second one. So, that will become wrong. So, update often is dependent on the order.

(Refer Slide Time: 29:58)

The slide has a title 'Case Statement for Conditional Updates' at the top right. On the left, there is a small image of a sailboat on water. Below the title, there is a bulleted list: 'Same query as before but with case statement'. To the right of the list is a SQL update statement:

```
update instructor  
set salary = case  
    when salary <= 100000 then salary * 1.05  
    else salary * 1.03  
end
```

At the bottom left, there is a small video thumbnail showing a person speaking. The video player interface shows '08.32' and 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur'. The bottom right corner of the slide has the text '©Silberschatz, Korth and Sudarshan'.

And therefore, you have yet another ah feature to take care of this when you have a specific order to do things it is called the case. So, you say when salary case is a new keyword, when is a key word when salary is less than equal to 100,000, then this is how you hike otherwise this is how you hike. So

(Refer Slide Time: 30:28)

The slide has a title 'Updates with Scalar Subqueries' at the top right. On the left, there is a small logo of a sailboat on water. The main content area contains a bulleted list of points and some SQL code. The footer includes the text 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. Des., IIT Kharagpur - Jam-Agr.', 'Database System Concepts - 8th Edition', '08.33', and '©Silberschatz, Korth and Sudarshan' along with a set of navigation icons.

- Recompute and update tot_creds value for all students

```
update student S
set tot_cred = (select sum(credits)
    from takes, course
    where takes.course_id = course.course_id and
        S.ID= takes.ID and
        takes.grade <> 'F' and
        takes.grade is not null);

```

- Sets tot_creds to null for students who have not taken any course
- Instead of `sum(credits)`, use:

```
case
when sum(credits) is not null then sum(credits)
else 0
end
```

It can it looks more like the; if statement of C, C++ you can do updates with scalar sub queries. We have seen scalar sub queries already. So, you can use a scalar sub query, again I would not go through that details please study and you will be able to understand.

(Refer Slide Time: 30:50)

The slide has a title 'Module Summary' at the top right. On the left, there is a small logo of a sailboat on water. The main content area contains a bulleted list of points. The footer includes the text 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. Des., IIT Kharagpur - Jam-Agr.', 'Database System Concepts - 8th Edition', '08.34', and '©Silberschatz, Korth and Sudarshan' along with a set of navigation icons.

- Introduced nested sub-query in SQL
- Introduced data modification

So, these are different examples. So to summarize, we have introduced a very powerful feature in SQL query known as nested sub query, where we can write a SELECT FROM WHERE expression as a part of the where predicate or as a relation in the from clause or as one or more collection of attributes in the select clause and it can be used in several

other places also. We have seen the ways to perform data modification in terms of deleting inserting and updating records. And we have also seen how nested sub query often may be very useful not only in terms of performing a query, but also in terms of performing certain data modifications.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science and Engineering
Indian Institute of Technology, Kharagpur

Lecture - 09
Intermediate SQL/1

Welcome to module 9 of database management systems. We have discussed about the introductory level of SQL the structured query language. in this module and the next we will take up some more intermediate level features of SQL. So, these modules are called intermediate SQL.

(Refer Slide Time: 00:42)

Module Recap

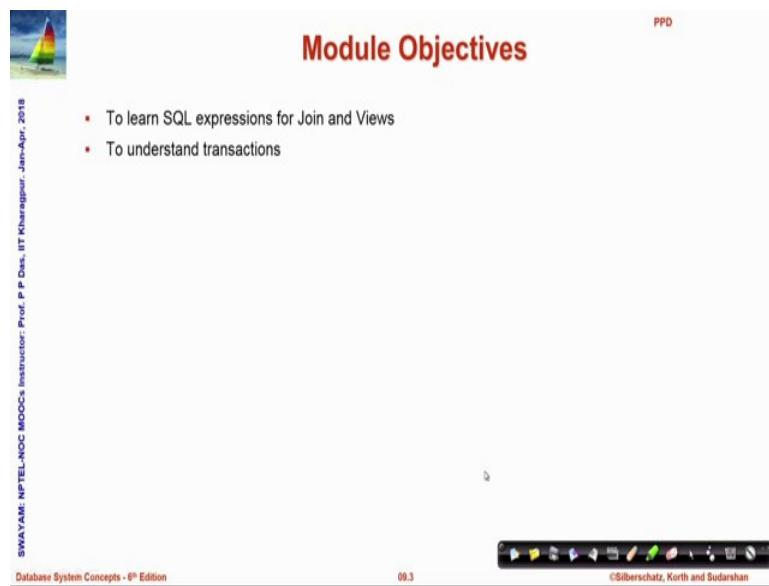
- Nested Subqueries
- Modification of the Database

PPD

Jun-Apr. 2018
Prof. P. P. Das, IIT Kharagpur - Jun-Apr. 2018
SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jun-Apr. 2018
Database System Concepts - 8th Edition
00.2
©Silberschatz, Korth and Sudarshan

So, in the last module this is what we have done which was part of the closing part of the introductory SQL. The nested sub queries and modifications to the database.

(Refer Slide Time: 00:55)



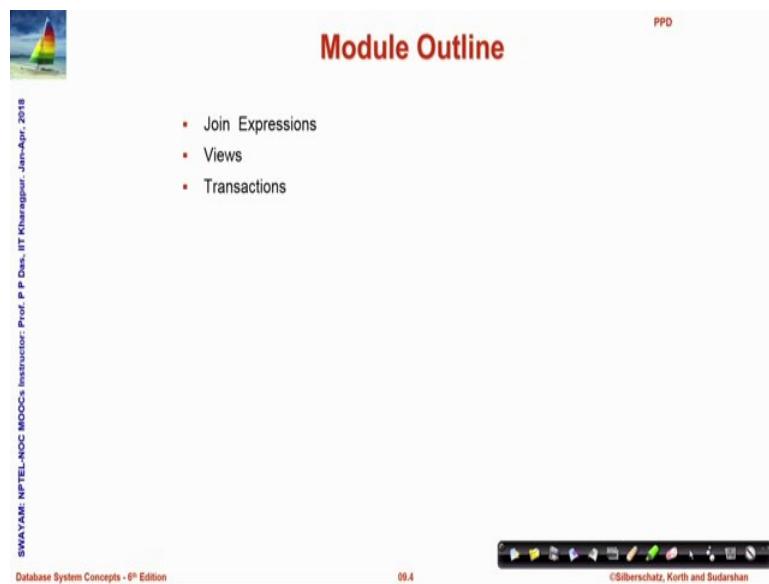
This slide is titled "Module Objectives" in red text at the top center. In the top right corner, there is a small red "PPD" label. On the left side, there is a small sailboat icon. The main content consists of a bulleted list:

- To learn SQL expressions for Join and Views
- To understand transactions

At the bottom left, it says "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018". At the bottom center, it says "Database System Concepts - 8th Edition" and "09.3". At the bottom right, it says "©Silberschatz, Korth and Sudarshan" and features a standard presentation navigation toolbar.

Today, we will, in this module learn about SQL expressions for join and views and we will take a quick look into understanding the transaction.

(Refer Slide Time: 01:11)



This slide is titled "Module Outline" in red text at the top center. In the top right corner, there is a small red "PPD" label. On the left side, there is a small sailboat icon. The main content consists of a bulleted list:

- Join Expressions
- Views
- Transactions

At the bottom left, it says "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018". At the bottom center, it says "Database System Concepts - 8th Edition" and "09.4". At the bottom right, it says "©Silberschatz, Korth and Sudarshan" and features a standard presentation navigation toolbar.

This is the module outline as it will span.

(Refer Slide Time: 01:15)

The slide features a small sailboat icon in the top left corner. In the top right, there is a red 'PPD' logo. Below it, a bulleted list includes 'Join Expressions', 'Views', and 'Transactions'. The main title 'JOIN EXPRESSIONS' is centered in large red capital letters. At the bottom, there is a navigation bar with icons for back, forward, search, and other presentation controls.

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr., 2018

JOIN EXPRESSIONS

PPD

- Join Expressions
- Views
- Transactions

Database System Concepts - 8th Edition

09.5

©Silberschatz, Korth and Sudarshan

So, we start with the join expressions in SQL join as we have already introduced.

(Refer Slide Time: 01:21)

The slide has a small sailboat icon in the top left. The title 'Joined Relations' is centered in red capital letters. Below the title is a bulleted list of four points about join operations. At the bottom, there is a video frame showing a man speaking, a navigation bar, and course information.

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr., 2018

Joined Relations

- **Join operations** take two relations and return as a result another relation
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition).
- It also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause

Database System Concepts - 8th Edition

09.6

©Silberschatz, Korth and Sudarshan

Takes two relations and returns a result as another relation. So, it is two different instances of two schemas and we try to connect them according to certain properties.

So, a join operation, is primarily a Cartesian product, which relates tuples in two relations under certain conditions of match. It also specifies that after the joining has been done, what are the tuples, which will be present in the output join. So, join

operation typically uses sub query, is used in a sub query, in the FROM clause we will see those usages later.

(Refer Slide Time: 02:19)

The slide has a header 'Types of Join between Relations' with a small sailboat icon. On the right, it says 'PPD'. The left margin contains vertical text: 'SWAYAM: NPTEL-NOOC MOOCs', 'Instructor: Prof. P. P. Das, IIT Kharagpur', and 'Date: Jan-Apr., 2018'. The main content is a bulleted list of join types:

- Cross join
- Inner join
 - Equi-join
 - Natural join
- Outer join
 - Left outer join
 - Right outer join
 - Full outer join
- Self-join

At the bottom, there's a video player showing a man speaking, the text 'Database System Concepts - 8th Edition', the number '09.7', and the copyright '©Silberschatz, Korth and Sudarshan'.

So, if we look into the different types of join that SQL support, these are the different classifications. So, we have CROSS JOIN, we have INNER JOIN, which specifically could be EQUI JOIN and even more specifically NATURAL JOIN and we will see there are variety of OUTER JOIN, that are possible. There could be a SELF JOIN also, where one relation is joined with itself.

(Refer Slide Time: 02:51)

The slide has a header 'Cross Join' with a small sailboat icon. On the right, it says 'PPD'. The left margin contains vertical text: 'SWAYAM: NPTEL-NOOC MOOCs', 'Instructor: Prof. P. P. Das, IIT Kharagpur', and 'Date: Jan-Apr., 2018'. The main content is a bulleted list explaining CROSS JOIN:

- CROSS JOIN returns the Cartesian product of rows from tables in the join
 - Explicit

```
select *  
from employee cross join department;
```
 - Implicit

```
select *  
from employee, department;
```

At the bottom, there's a video player showing a man speaking, the text 'Database System Concepts - 8th Edition', the number '09.8', and the copyright '©Silberschatz, Korth and Sudarshan'.

CROSS JOIN is just a formal join based name for Cartesian product of two rows. So, you could explicitly do a CROSS JOIN, which you can see here or you can implicitly also do a CROSS JOIN, but by specifying two or more relations in from clause, and taking all the attributes from there, we have seen these kind of Cartesian products earlier. So, CROSS JOIN here is more as placeholder. in the context of the join semantics, the pure Cartesian product is a CROSS JOIN, but what would be more interesting is, when we take different kinds of inner and OUTER JOINs.

(Refer Slide Time: 03:30)

Join operations – Example

- Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that
 - prereq information is missing for CS-315 and
 - course information is missing for CS-437

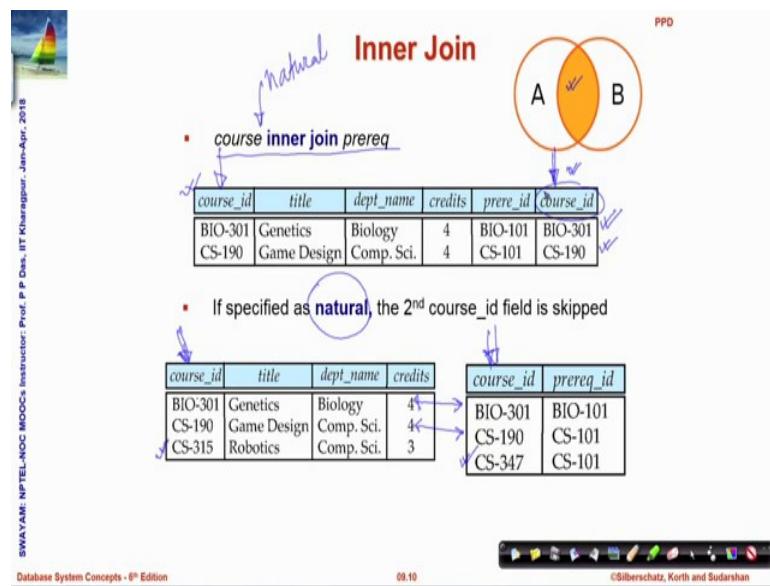
Navigation icons: back, forward, search, etc.

SWAYAM: NPTEL-NOCO MOOCs Instructor: Prof. P. P. Deshpande, IIT Kharagpur
 Database System Concepts - 8th Edition 09.9 ©Silberschatz, Korth and Sudarshan

So, let us start with a simple example to understand the issues. So, there is a relation *course*, which has four attributes and in this particular instance, it has three tuples; three rows and there is another relation *prereq*, which specifies the prerequisite for every course.

So, it has two attributes; the *course_id* and the corresponding *prereq_id*. It also has three rows; three tuples here, and if you look at the instances carefully, you will find that the three courses that are specified in the *course* relation all are not specified in the *prereq* relation. Bio-301 and CS-190 is present in *prereq*, but CS-315 is not present in at the same time, the *prereq* has one particular tuple, specifying the prerequisite of CS-347, which in turn is not present in the *course* relation. So, with this observation, let us start trying to see what different joint mean.

(Refer Slide Time: 04:54)



So, in a join is computed, then in terms of the two relations that we have, there is an attribute `course_id`, which is common. So, once we have taken the cross product, we will from the cross product, only retain those rows, where the `course_id` in relation `course` and the `course_id` in relation `prereq` are same. So, when we do this particular record, when it gets mapped with this corresponding record, it will generate the corresponding output record. Similarly, CS 190, when it is mapped to the CS 190, in the `prereq`, it will generate the second record, we have already understood this, the third record in the courses CS 315 has no match here in `prereq`.

So, that will not feature in the output. Similarly, in the `prereqs` C S 347 that exist has no match in `courses`. So, that also will not appear in the output and also in the output you find that the `course_id` has actually featured twice. This is the first column `course_id` comes from `course`. So, it should more formally be called **course.course_id** whereas, the second one comes from `prereq`. So, it should that should formally be called **prereq.course_id**.

Now, if in addition to saying that, this is an INNER JOIN, if I also specify the word natural, I can say natural here, if say natural, then this second duplicate attribute `course_id` will be dropped from the output that becomes a NATURAL JOIN, INNER JOIN as the name suggests, finds out the inner part of the two relations.

So, if we look at the two relations as A and B only those records, which are both have instance in A as well as B, in terms of equality of this course_id attribute will come in the output. So, this is the basic type of join, INNER JOIN which is most commonly used.

(Refer Slide Time: 07:34)

The slide has a title 'Outer Join' in red font. To the left of the title is a small icon of a sailboat on water. On the right side, there is a decorative footer bar with various icons. The slide content includes a bulleted list of four points:

- An extension of the join operation that avoids loss of information
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join
- Uses *null* values

At the bottom left, there is vertical text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Date: Apr. 2018'. At the bottom center, it says 'Database System Concepts - 8th Edition' and '09.11'. At the bottom right, it says '©Silberschatz, Korth and Sudarshan'.

Now, we can extend this into a different kind of join, known as OUTER JOIN in the INNER JOIN. As you have seen that courses that exist in the course relation, but are not there in the prereq or the ones that exist in the prereq and is not there in the course are not featuring in the final INNER JOIN output. So, there is some loss of information in terms of this. So, why we are doing this we can compute and add tuples from one relation that may not match with any tuple in the other relation and if we want to do that then naturally for the other attributes of that tuple in the target relation we will not know the values. So, we will use null values this is the basic idea of OUTER JOIN.

(Refer Slide Time: 08:34)

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

So, let us see what it specifically means? We first talked about left OUTER JOIN, left in the sense that we have, this is how it is written. Left OUTER JOIN is a sequence of commands that you give, you are also saying it is natural, which means that the common attribute will not feature twice in the output and this is the left relation and this is the right relation. So, left OUTER JOIN specifies that in the output all records of the left relation, in this case the course relation must feature.

So, naturally when we do the join, we will get these two records as we have got in terms of the INNER JOIN, in terms of course 315 the CS 315, the third course there is no instance in the prereq, we will still have that in the output, but since the prereq value for that, the prerequisite value is not known, the prerequisite id will be set to null here. So, left OUTER JOIN ensure that, all relations of the left relation, all tuples of the left relation will necessarily feature in the output and that is a reason. If you see in the Venn diagram, the whole of this set, A is shown whereas, this part certainly will not feature.

(Refer Slide Time: 10:05)

Right Outer Join

A B

course natural right outer join prereq

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

Now, similarly we can have a right OUTER JOIN, where the concept is the same except that. Now, we ensure that all records of the right relation, in this case the prereq relation will feature and therefore, C S 2347 for which there is no entry in the course relation will also come as a record and since we do not know the title department name and credits for these fields, we will put them as null and this again is a natural one. So, course_id is featuring only once. So, you will understand that since, we have a left version and we have a right version.

(Refer Slide Time: 10:47)

Joined Relations

- Join operations take two relations and return as a result another relation
- These additional operations are typically used as subquery expressions in the **from** clause
- Join condition – defines which tuples in the two relations match, and what attributes are present in the result of the join
- Join type – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated

Join types
inner join
left outer join
right outer join
full outer join

Join Conditions
natural
on <predicate>
using (A_1, A_1, \dots, A_n)

We can actually have a full version as well. So, if we look into the join relations in general, it takes two relations and returns a result and those additional operations are used in the sub query in from and there is a set of join conditions. So, these are the join conditions that we are specifying, whether it is natural and we will soon see that we can actually not depend only on the attributes that are common, we can actually specify that which attributes should be used in computing the joint. So, those are on condition and the using clause, we will just illustrate them soon and finally, there are four types of join that can happen, that is the INNER JOIN. We have seen the left OUTER JOIN, right OUTER JOIN and we will soon see the full OUTER JOIN.

(Refer Slide Time: 11:52)

Full Outer Join

PPD

- course natural full outer join prereq

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

SWAYAM NPTEL-MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur. Jam-Apr. 2015

Database System Concepts - 8th Edition 09.15 ©Silberschatz, Korth and Sudarshan

So, full OUTER JOIN as you must have guessed will ensure that you get certainly the tuples from the INNER JOIN, which is here, you will get the tuple from the left OUTER JOIN that is here, that is a tuple which exists in course and there is no corresponding matching tuple in the prereq and you will also get the tuple from the right OUTER JOIN, that is for tuple, which exists in the prereq relation, but there is no corresponding tuple in the course relation and corresponding missing values are all set to null. So, these three kinds of OUTER JOIN are possible.

(Refer Slide Time: 12:38)

The slide has a header 'Joined Relations – Examples' with a sailboat icon. On the left, vertical text reads 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018'. At the bottom, it says 'Database System Concepts - 8th Edition' and '09.16 ©Silberschatz, Korth and Sudarshan'.

course inner join prereq on
`course.course_id = prereq.course_id`

course_id	title	dept_name	credits	prere_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

▪ What is the difference between the above (equi_join), and a natural join?

▪ course left outer join prereq on
`course.course_id = prereq.course_id`

course_id	title	dept_name	credits	prere_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	null	null

So, you can also specify join by saying that explicitly saying what attribute we want to join on and if you specify that, then you are saying it is a course INNER JOIN prereq. This part was same then you are putting an on clause, saying in the on clause, you will have to provide a predicate that is, which field should equate or match with what field. So, you are saying `course.course_id = prereq.course_id`.

So, this result incidentally happens to be same as just doing the INNER JOIN, but we are illustrating that, on clause can explicitly use. For example, between the two relations, we have more than one common attribute, but we may want to actually do the INNER JOIN based on only one of them or equality on two of them and so on..

So, this kind of a join, where INNER JOIN, where you set two fields to be equal or two or more fields to be equal is also known as EQUI JOIN and since we have not specified natural, you can again observe, then the `course_id` attribute has occurred twice. If it was said natural then the second `course_id` attribute would not have come in the result, this is a showing. The left OUTER JOIN in terms of on clause and we have seen similar results. And now, this can be seen in terms of the on clause as well, and you can see in that second `course_id` field. This entry is null, because actually you do not have that in the prerequisite set and; obviously, this set will be null, this field will be null..

(Refer Slide Time: 14:35)

The slide title is 'Joined Relations – Examples'. It features two tables illustrating joins:

- course natural right outer join prereq**:
A table showing courses and their prerequisites. Courses BIO-301 and CS-190 have prerequisites, while CS-347 does not.

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101
- course full outer join prereq using (course_id)**:
A table showing courses and their prerequisites, using course_id for the join. Courses BIO-301, CS-190, CS-315, and CS-347 are listed with their respective prerequisites.

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101

Other slide details: SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. Des. IIT Kharagpur - Jam-Astr. 2018. Database System Concepts - 8th Edition. 09.17. ©Silberschatz, Korth and Sudarshan.

So, this is another example, showing you the natural right OUTER JOIN, this is you, showing you full OUTER JOIN and we are showing the use of the using clause. We can say using and put a set of attributes and the meaning is the join will be performed, based on those attributes. So, here in this case again the join will be based on course_id.

(Refer Slide Time: 15:01)

The slide title is 'VIEWS'. It lists three types of views:

- Join Expressions
- Views
- Transactions

Other slide details: SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. Des. IIT Kharagpur - Jam-Astr. 2018. Database System Concepts - 8th Edition. 09.18. ©Silberschatz, Korth and Sudarshan.

So, that was about different kinds of join that we can do, which we going forward. We will see that form, a very critical has a very critical place, in terms of query formulation. Now, we take you to a different concept known as views now.

(Refer Slide Time: 15:20)

Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know an instructor's name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users

relation that is not of the conceptual model but is made available to a user as a "virtual relation" is called a **view**

We have seen that. So, far we have been computing certain query results, based on one or more relations one or more instances. Now, in some cases, we may want the results to be restrictive in terms of based on the user or based on the context, in which the result should be used. So, we may not want all fields of a result to be visible to all the users or to the application. So, we may not expose the whole logical model and in those cases, we introduce a view. So, here we are showing one, where, from the instructor relation, we are only picking up three fields and we are not picking up the salary field.

Now, you would think that, well this is what we can do in terms of the normal query and certainly then, what is the point of using this? Now, what we can do is, we can create this not adjust as a query, but as a view one, once we create this as a view, it actually this query expression is treated as what is known as a view expression and every time you want to use that view. The actual tuples in that view are computed, but this is not actually a relation that exists in the database. So, it is a kind of, can be thought of as a kind of virtual relation, which exists, which can be seen only when you use that.

(Refer Slide Time: 17:22)

The slide has a header 'View Definition' in red. On the left, there is a small sailboat icon and a vertical sidebar with text: 'SWAYAM: NPTEL-NOC MOOCs', 'Instructor: Prof. P. P. Desai, IIT Kharagpur', and 'Date: Jan-Apr., 2018'. The main content area contains a bulleted list and a code snippet:

- A view is defined using the **create view** statement which has the form
`create view v as < query expression >`
- where **<query expression>** is any legal SQL expression
- The view name is represented by **v**
- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates
- View definition is not the same as creating a new relation by evaluating the query expression**
 - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view

At the bottom, there is a video frame showing a man speaking, the text 'Database System Concepts - 8th Edition', the time '09:20', and a navigation bar.

So, there is a subtle but very strong difference between, actually computing a result, through a select query and defining a view, based on a select query and then making use of the view, as if it were actually a relation that existed. So, to do this, this is how we go about. It is the send text, is very similar to the create table. So, you do a create view give a name and then you specify as is the connective and specify the query expression, which is an SQL query, which will let you compute the view, every time you actually need it..

So, this is a view name, once a view is defined, the view name can be used as a virtual relation, it can be used exactly as we use any of the really existing relation, the conceptual relations that we have created, through create table. So, it is the difference, this is what needs to be understood very well, the view definition is not the same as creating a new relation, once you create the new relation, the time you have created it, you get the result and that result is explicitly available as a set of tuples as a table rather a view is a definition, which you stored in the database as an expression. So, every time you make use of that view, at that time the set of tuples are computed. It is not existing in the database as stored like the real relations and based on that computation, all the rest of the query will actually be executed.

(Refer Slide Time: 18:58)

The slide has a header 'Example Views' with a sailboat icon. It contains three bullet points:

- A view of instructors without their salary
create view faculty as
select ID, name, dept_name
from instructor
- Find all instructors in the Biology department
select name
from faculty
where dept_name = 'Biology'
- Create a view of department salary totals
create view departments_total_salary(dept_name, total_salary) as
select dept_name, sum(salary)
from instructor
group by dept_name;

Annotations include blue circles around 'dept_name' in the first two queries and 'dept_name, sum(salary)' in the third query, with arrows pointing to the word 'salary'.

Navigation icons at the bottom include back, forward, search, and other presentation controls.

So, let us take a quick look, this is a create view, we have created the view of a, of view called faculty, from instructor relation. Instructor relation is the real one, I existing one and faculty is a view expression being created and in that, what we have done? Simply, we have taken a done, a projection, we have lived left out the salary field. Now, we can make use of that view, you can see that we are doing from faculty.

So, this actually is a view, but this behaves as if this is varying relation. So, from faculty we are trying to find out the name of all those faculties; who belong to the biology department. So, what will happen, when they want to execute this query? This will refer to this view. So, to execute this query, it will have to first execute this query, get the temporary virtual instance of the virtual relation, created and based on that, this query will be computed and the results will be given accordingly.

So, that is the basic purpose of the view that the whole thing, the whole view expression remains as an abstraction in the database and computed whenever it is used. So, this is showing you another view, which shows certain computed information. For example, we are creating a view for departmental total salary, which will show as two fields department name and total salary, which has been created by aggregation.

So, anytime we make use of this view in a from clause, we will get, we will feel as if such a relation really exists where the department name and the total salary of the instructors in that department are stored, but it really does not exist. It is computed

whenever it is needed, whenever it is used, you can actually use views to create other views.

(Refer Slide Time: 20:52)

The slide has a title 'Views Defined Using Other Views' at the top right. On the left, there is a small image of a sailboat on water. A vertical column of text on the left edge reads 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr. 2018'. The main content area contains two bullet points with SQL code:

- `create view physics_fall_2009 as
select course.course_id, sec_id, building, room_number
from course, section
where course.course_id = section.course_id
and course.dept_name = 'Physics'
and section.semester = 'Fall'
and section.year = '2009';`
- `create view physics_fall_2009_watson as
select course_id, room_number
from physics_fall_2009
where building='Watson';`

At the bottom left, there is a video thumbnail showing a man speaking. The bottom right corner shows a navigation bar with icons for back, forward, search, and other controls. The footer of the slide includes 'Database System Concepts - 6th Edition', '09.22', and '©Silberschatz, Korth and Sudarshan'.

For example, this is one view which is the view of `physics_fall_2009`, which are all courses that are offered in physics, from the physics department in the semester fall of year 2009 and using that we can create another view. See here again, we are in the, from clause we are using this view. So, creating this using this view, we are creating yet another view, which shows the courses that ran in the Watson building. So, views can be used as I have already said as any other.

(Refer Slide Time: 21:35)

The slide features a small sailboat icon in the top left corner. The title 'View Expansion' is centered at the top in a red font. Below the title is a bulleted list: '▪ Expand use of a view in a query/another view'. To the right of the list is a block of SQL code. A curly brace is placed after the 'where building = 'Watson'' line, grouping the entire 'create view' statement. At the bottom of the slide, there is a video frame showing a man speaking, the text 'Database System Concepts - 8th Edition', the number '09.23', and the copyright notice '©Silberschatz, Korth and Sudarshan'.

```
create view physics_fall_2009_watson as
(select course_id, room_number
from (select course.course_id, building, room_number
      from course, section
     where course.course_id = section.course_id
       and course.dept_name = 'Physics'
       and section.semester = 'Fall'
       and section.year = '2009')
   where building= 'Watson';
```

Actual relation, but they do not really exist. So, if you expand out, if you just put the physics fall 2009 expression, within the view definition of physics for 2009 Watson, this is your earlier view relation. So, this is known as view expansion. So, this is actually the query that, you are executing and that has a lot of value.

(Refer Slide Time: 22:04)

The slide features a small sailboat icon in the top left corner. The title 'Views Defined Using Other Views' is centered at the top in a red font. Below the title is a bulleted list: '▪ One view may be used in the expression defining another view', '▪ A view relation v_1 is said to depend directly on a view relation v_2 if v_2 is used in the expression defining v_1 ', '▪ A view relation v_1 is said to depend on view relation v_2 if either v_1 depends directly to v_2 or there is a path of dependencies from v_1 to v_2 ', and '▪ A view relation v is said to be recursive if it depends on itself'. At the bottom of the slide, there is a video frame showing a man speaking, the text 'Database System Concepts - 8th Edition', the number '09.24', and the copyright notice '©Silberschatz, Korth and Sudarshan'.

So, as we have said views can be defined indirectly from one relation. So, these are called direct dependence or they could be defined in terms of a chain of relations v_1 in

terms of v2, v2 in terms of v3 and so on. And a view relation can be recursive also, that a view could be in terms of itself.

(Refer Slide Time: 22:29)

The slide features a small sailboat icon in the top left corner. The title 'View Expansion*' is centered in red. On the left, vertical text reads 'SWAYAM: NPTEL-NOOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018'. The main content is a bulleted list:

- A way to define the meaning of views defined in terms of other views
- Let view v_i be defined by an expression e_1 that may itself contain uses of view relations
- View expansion of an expression repeats the following replacement step:

```
repeat
    Find any view relation  $v_i$  in  $e_1$ 
    Replace the view relation  $v_i$  by the expression defining  $v_i$ 
    until no more view relations are present in  $e_1$ 
```
- As long as the view definitions are not recursive, this loop will terminate

At the bottom, there's a video player showing a man speaking, with controls for volume, brightness, and navigation. The text 'Database System Concepts - 8th Edition' and '09.25' are at the bottom left, and '©Silberschatz, Korth and Sudarshan' is at the bottom right.

A lot of power view expansion is the process that SQL uses to evaluate a view. So, I would request you to study this and understand that this process works. This is pretty much like pseudo code C program.

(Refer Slide Time: 22:47)

The slide features a small sailboat icon in the top left corner. The title 'Recursive View' is centered in red. On the left, vertical text reads 'SWAYAM: NPTEL-NOOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018'. The main content is a bulleted list:

- In SQL, recursive queries are typically built using these components:
 - A non-recursive seed statement
 - A recursive statement
 - A connection operator
 - The only valid set connection operator in a recursive view definition is UNION ALL
 - A terminal condition to prevent infinite recursion

At the bottom, there's a video player showing a man speaking, with controls for volume, brightness, and navigation. The text 'Database System Concepts - 8th Edition' and '09.26' are at the bottom left, and '©Silberschatz, Korth and Sudarshan' is at the bottom right.

Now, moving to recursive views, the views where the same relation can be used in the view to define another view, we need like every other recursive structure. We need first a

non-recursive statement, which is called the seed statement. we need a recursive statement, which can recur, we need a connection operator, which can connect the non-recursive and the recursive results together put them together, the only connective that is valid is union, all that is multi set union and we also need some kind of a terminal condition to guarantee that the recursion really a terminus, it does not go on forever.

(Refer Slide Time: 23:36)

Recursive View – Example

- In the context of a relation *flights*:

```
create table flights (
    source  varchar(40),
    destination varchar(40),
    carrier  varchar(40),
    cost     decimal(5,0));
```

source	destination	carrier	cost
Paris	Detroit	KLM	7
Paris	New York	KLM	6
Paris	Boston	American Airlines	8
New York	Chicago	American Airlines	2
Boston	Chicago	American Airlines	6
Detroit	San Jose	American Airlines	4
Chicago	San Jose	American Airlines	2

- Find all the destinations that can be reached from 'Paris'

SWAYAM: NPTEL-NOCs Instructor: Prof. P. P. Desai, IIT Kharagpur Date: 2018

Database System Concepts - 8th Edition

09.27

©Silberschatz, Korth and Sudarshan

So, let us take an example. So, this is in context of a relation *flights*, where the four fields are as specified and there is an instance shown, which show that different source destination of different carriers, carrying people from one source to the other destination and what do you want to find is all destinations that can be reached from Paris. Now, you can see that from Paris, if I can reach Detroit and from Detroit I can reach San Jose, then I can actually reach San Jose from Paris. So, that is the basic reachability. So, that will necessarily, if I want to compute that, then I will be able to compute this very easily by doing NATURAL JOIN of flights with flights provided, I take say source.

(Refer Slide Time: 24:38)

Recursive View – Example

* In the context of a relation *flights*:

```
create table flights (
    source      varchar(40),
    destination varchar(40),
    carrier     varchar(40),
    cost        decimal(5,0));
```

source	destination	carrier	cost
Paris	Detroit	KLM	7
Paris	New York	KLM	6
Paris	Boston	American Airlines	8
New York	Chicago	American Airlines	2
Boston	Chicago	American Airlines	6
Detroit	San Jose	American Airlines	4
Chicago	San Jose	American Airlines	2

* Find all the destinations that can be reached from 'Paris'

flights f1 \bowtie flights f2
 $f_1.destination = f_2.source$

Let us compute it like this, flights f1 join, flights f2 and I will have **f1.destination = f2.sources**. So, the idea is, if something goes from Paris to Detroit that is in f1 and if some flight goes from Detroit to San Jose that is in f2, then the destination in f1 and the source in f2 have to be equated. So, if we do this kind of a self EQUI JOIN, then we will be able to find out all flights that go from Paris to San Jose or all places that you can reach from Paris in one hop.

Naturally, once you reach, once you do that then you may be able to go to another destination in two hops and once you do that then, you may be able to reach another, yet another destination in three hops and so on. So, we do not really know how many hops, maximum would be required to compute this reachability information. So, that is the reason we need to make use of the recursion and so, this is how we express it.

(Refer Slide Time: 25:54)

Recursive View – Example

```

create recursive view reachable_from (source,destination,depth) as (
    select root.source, root.destination, 0 as depth
    from flights as root
    where root.source = 'Paris'
union all
    select in1.source, out1.destination, in1.depth + 1
    from reachable_from as in1, flights as out1
    where in1.destination = out1.source and
          in1.depth <= 100);

```

- A non-recursive seed statement
- A recursive statement
- A connection operator
- A terminal condition to prevent infinite recursion

Get the result by simple selection on the view:

```

select distinct source, destination
from reachable_from;

```

source	destination	carrier	cost
Paris	Detroit	KLM	7
Paris	New York	KLM	6
Paris	Boston	American Airlines	8
New York	Chicago	American Airlines	2
Boston	Chicago	American Airlines	6
Detroit	San Jose	American Airlines	4
Chicago	San Jose	American Airlines	2
Paris	Detroit		
Paris	New York		
Paris	Boston		
Paris	Chicago		
Paris	San Jose		

This example view, `reachable_from`, is called the *transitive closure* of the `flights` relation

Source: https://info.teradata.com/HTMLPubs/DB_TTU_16_00/index.html#page/SQL_Reference%2FB035-1184-160K%2Fname1472241335807.html%23wwID0EJ23T

So, if we look into this, we are specifying that is a recursive view. It will happen with itself, this is the name and this is what we want to compute source destination and we take another dummy attribute kind of which specify the depth of recursion. So, the present instance of the relation is at depth 0.

So, which defines your non recursive seat part. so, it says select. So, you have renamed it as flights as route, you are specified that it has to start from Paris and you can find out the source destination pair at depth 0, then you specify the recursive part that is the second hop has to be defined. So, we are saying that, if you had the reachability then call, let us call it in one. This reachability maybe in one hop that is a depth 0 maybe in 2 hops that is a depth 1 maybe at 3 hops; that is a depth 2 and you take another instance of flight as out 1 and what you need is the destination in the first in one has to be same as the source in the other so that they get connected and then you output the source from the first one destination from the second one and naturally the depth has got incremented, because you have done added one more hop and so, this is the and you add another condition saying that **in1.depth** should be \leq to 100. This is as I mentioned is a terminal condition which makes sure that, you do not get into infinite recursion. So, this view, recursive view cannot be used to compute any reachability, which has more than 101 hops.

So, that is to be noted and finally, we need to connect these two results, which is the initial start seed and the recursive one. So, this is the connection operator.

So, this is basically the idea of the recursive view, those of you who are more familiar with discrete structure, would have known among relations in some more depth, you would know that, we can define a transitive closure of a binary relation. So, this recursive view is necessarily computing the transitive closure from the flights relation. So, this is the instance of the flights and on the final computation, this is what you get. This gives you all the destinations that can eventually be reached from the source Paris.

(Refer Slide Time: 28:41)

The slide has a title 'The Power of Recursion' in red at the top right. To its left is a small icon of a sailboat on water. On the far left edge, there is vertical text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - Jan-Apr. 2018'. At the bottom left is the text 'Database System Concepts - 8th Edition'. In the center, there is a bulleted list:

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration
 - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *flights* with itself
 - This can give only a fixed number of levels of reachable destinations
 - Given a fixed non-recursive query, we can construct a database with a greater number of levels of reachable destinations on which the query will not work

At the bottom right are navigation icons and the text '09.29 ©Silberschatz, Korth and Sudarshan'.

So, the recursive is very powerful in the sense that without recursion a non-recursive version can only find flights up to a certain number of hops and whatever ma query you write it is always possible to write out a database instance which will have more hops and your query will necessarily fail.

(Refer Slide Time: 29:05)

The slide title is "The Power of Recursion". The text on the slide includes:

- Computing transitive closure using iteration, adding successive tuples to `reachable_from`
 - The next slide shows a `flights` relation
 - Each step of the iterative process constructs an extended version of `reachable_from` from its recursive definition
 - The final result is called the *fixed point* of the recursive view definition.
- Recursive views are required to be **monotonic**. That is, if we add tuples to `flights` the view `reachable_from` contains all of the tuples it contained before, plus possibly more

A hand-drawn diagram is present on the slide, showing a ladder leaning against a wall, with the word "union all" written above it.

At the bottom of the slide, there is a navigation bar with icons and text indicating the slide number (09.30) and copyright information (©Silberschatz, Korth and Sudarshan).

So, we make use of the recursion here to make sure that you can actually extend this to whatever depth you want and to compute this we keep on computing till no changes are possible and in that sense these recursive views are said to be monotonic in that every time you compute your result necessarily becomes larger and that is the reason you for the purpose of being monotonic you are actually making use of the union all. So, that makes it all inclusive

(Refer Slide Time: 29:49)

The slide title is "Example of Fixed-Point Computation". It features a table of flight data and a table showing the progression of tuples in closure through three iterations.

Flight Data Table:

source	destination	carrier	cost
Paris	Detroit	KLM	7
Paris	New York	KLM	6
Paris	Boston	American Airlines	8
New York	Chicago	American Airlines	2
Boston	Chicago	American Airlines	6
Detroit	San Jose	American Airlines	4
Chicago	San Jose	American Airlines	2

Iterations Table:

Iteration #	Tuples in Closure
0	Detroit, New York, Boston
1	Detroit, New York, Boston, San Jose, Chicago
2	Detroit, New York, Boston, San Jose, Chicago

At the bottom of the slide, there is a video player showing a professor speaking, and a navigation bar with icons and text indicating the slide number (09.31) and copyright information (©Silberschatz, Korth and Sudarshan).

So, now, if we go and this is the instance and this you can. Here, I have shown that how the iteration actually happens in the iteration 0, in the flights itself, you had three destinations, then you add two more in iteration 1, in iteration 2, you do not add anything else. So, your result henceforth will not change. So, you have reached a fixed point and computations, are over.

(Refer Slide Time: 30:10)

The slide has a header 'Update of a View' with a sailboat icon. The main content is a bulleted list:

- Add a new tuple to *faculty* view which we defined earlier

SQL code:

```
insert into faculty values ('30765', 'Green', 'Music');
```

This insertion must be represented by the insertion of the tuple

```
('30765', 'Green', 'Music', null)
```

into the *instructor* relation

At the bottom left is a video frame showing a man speaking, with text 'SWAYAM: NPTEL-NOC Instructor: Prof. P. P. Deshpande'. At the bottom right is a navigation bar with icons and the text 'Database System Concepts - 8th Edition' and '09.32'.

You can also update a view you can insert a record directly into a view, but since view only is partial information on the relation, when you INSERT INTO a view since, view is virtual. There will have to be an insertion, in the real relation and in the real relation you may not know certain fields. So, if you are doing this insertion into faculty, which is a view of instructor, then the salary field is not known. So, in the actual instructor a null will have to get inserted in the salary field. So, the salary field needs to be null able kind of field so updates on views have certain restrictions.

(Refer Slide Time: 30:47)

The slide features a small sailboat icon in the top-left corner. The title 'Some Updates cannot be Translated Uniquely' is centered at the top in a red font. On the left side, there is vertical text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Aut., 2018'. At the bottom left is the text 'Database System Concepts - 8th Edition'. In the bottom right corner, there is a navigation bar with icons for back, forward, search, and other presentation controls. The main content consists of a bulleted list of SQL statements and associated questions:

- `create view instructor_info as
 select ID, name, building
 from instructor, department
 where instructor.dept_name= department.dept_name;`
- `insert into instructor_info values ('69987', 'White', 'Taylor');`
 - which department, if multiple departments in Taylor?
 - what if no department is in Taylor?
- Most SQL implementations allow updates only on simple views
 - The `from` clause has only one database relation
 - The `select` clause contains only attribute names of the relation, and does not have any expressions, aggregates, or `distinct` specification
 - Any attribute not listed in the `select` clause can be set to null
 - The query does not have a `group by` or `having` clause

So, there are some more instances that I have given, which you can study and try to understand that what are the difficulties of updating on the view. So, it can be done, but it has to be done in a restrictive sense. So, these are the different conditions that has to happen for views to be updated.

(Refer Slide Time: 31:08)

The slide features a small sailboat icon in the top-left corner. The title 'And Some Not at All' is centered at the top in a red font. On the left side, there is vertical text: 'SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jam-Aut., 2018'. At the bottom left is the text 'Database System Concepts - 8th Edition'. In the bottom right corner, there is a navigation bar with icons for back, forward, search, and other presentation controls. The main content consists of a SQL statement and a question:

```
create view history_instructors as  
select *  
from instructor  
where dept_name= 'History';
```

- What happens if we insert ('25566', 'Brown', 'Biology', 100000) into `history_instructors`?

(Refer Slide Time: 31:11)

So, please go through these slides to understand what are there in terms of the views? Finally, view is a virtual relation, but it can be materialized also, that is materializing is basically computing a physical relation at the instance of the view, but naturally if you materialized then there is a certain point of time, where you have materialized where you have made it into a physical relation and hence, if your original source data in the view changes in future the materialized view also need to be updated otherwise, your data will get bad.

(Refer Slide Time: 31:43)

(Refer Slide Time: 31:49)

Transactions

- Unit of work
- Atomic transaction
 - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions
- Transactions begin implicitly
 - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
 - Can turn off auto commit for a session (e.g. using API)
 - In SQL:1999, can use: **begin atomic end**
 - Not supported on most databases

SWAYAM: NPTEL-NOOC Instructor: Prof. P. Desai, IIT Kharagpur - Jan-Apr - 2015

Database System Concepts - 8th Edition 09.37 ©Silberschatz, Korth and Sudarshan

Finally, in this module, we mentioned that, there is something called transactions, which we will take up at a later stage, in much depth. This is just to get you familiar with atom a transaction is a unit of work, which is usually atomic, which is either fully executed or if it fails, it will be rolled back as if it never occurred and this is required for isolation in concurrent transactions. So, we will talk about this lot more when we take up concurrency and related issues.

So, come transactions implicitly begin and they end by either committing the work that they have successfully finished or rolling back that, this cannot be done. So, there are some features in the SQL for doing transactions and, but usually you can transactions, commit by default and the only it is exceptions, when the rollback is happening and we will see more of that later.

(Refer Slide Time: 32:49)

Module Summary

- Learned SQL expressions for Join and Views
- Introduced transactions

SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Desai, IIT Kharagpur - Jan-Apr., 2018

Database System Concepts - 8th Edition

09.38

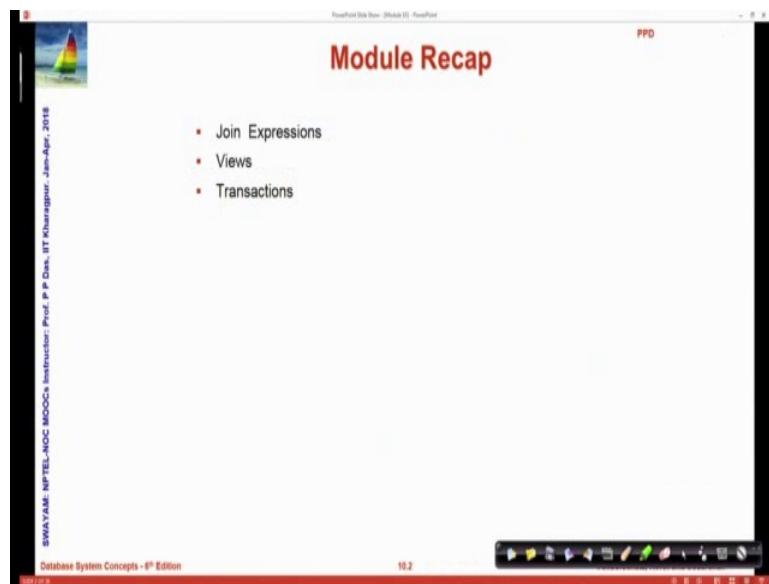
©Silberschatz, Korth and Sudarshan

So, to summarize in this module, we have learnt about two important SQL features in terms of join and views and we just introduced the basic notion of committing transactions.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science and Engineering
Indian Institute of Technology, Kharagpur

Lecture – 10
Intermediate SQL/2

(Refer Slide Time: 07:37)



Module Recap

- Join Expressions
- Views
- Transactions

SMAV YAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Das, IIT Kharagpur - June-April - 2018
Database System Concepts - 6th Edition

Welcome to module ten of database management systems. We have been discussing about intermediate level features in SQL; and this is a second and closing module on that. We have in the last module; talked about join expressions, views and transaction in a bit.

(Refer Slide Time: 07:46)

The slide is titled "Module Objectives" in a large, bold, red font at the top center. Below the title is a bulleted list of three items, also in red font. The footer of the slide includes a red bar with the text "Database System Concepts - 6th Edition" and "103". On the left side of the slide, there is a vertical sidebar with some text and a logo.

- To learn SQL expressions for integrity constraints
- To understand more data types in SQL
- To understand authorization in SQL

In this module, we will try to learn SQL expressions that are responsible for maintaining in the integrity of the database. We have talked about integrity a little. We will now see how explicitly integrity can be checked and how different kinds of integrity can be ensured through SQL. We will also talk about more data types.

We have seen the basic primitive data types, and we had promised that we will talk more about the data types including user defined data types here. And finally, we will talk about a very important aspect of authorisation as to who can do what in a database in through SQL.

(Refer Slide Time: 08:27)

The screenshot shows a Microsoft PowerPoint slide titled "Module Outline". The slide content is as follows:

- Integrity Constraints
- SQL Data Types and Schemas
- Authorization

At the bottom left of the slide, there is vertical text: "SWAYAM: NPTEL-NOC: MOOCs Instructor: Prof. P. P. Chakrabarti - IIT Kharagpur - Jan-Apr - 2016". At the bottom center, it says "Database System Concepts - 6th Edition". The bottom right corner shows the slide number "10.4". The top right corner has the word "PPD". The top bar of the slide indicates "PowerPoint Slide Show - Module 10 - PowerPoint".

So, this is a module out line.

(Refer Slide Time: 08:30)

The screenshot shows a Microsoft PowerPoint slide with the title "INTEGRITY CONSTRAINTS" in large red capital letters. Below the title, the same list of topics from the previous slide is displayed:

- Integrity Constraints
- SQL Data Types and Schemas
- Authorization

At the bottom left of the slide, there is vertical text: "SWAYAM: NPTEL-NOC: MOOCs Instructor: Prof. P. P. Chakrabarti - IIT Kharagpur - Jan-Apr - 2016". At the bottom center, it says "Database System Concepts - 6th Edition". The bottom right corner shows the slide number "10.5". The top right corner has the word "PPD". The top bar of the slide indicates "PowerPoint Slide Show - Module 10 - PowerPoint".

We start with integrity constraints.

(Refer Slide Time: 08:31)

The slide is titled "Integrity Constraints" in red. It contains a bulleted list of integrity constraints:

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency
 - A checking account must have a balance greater than \$10,000.00
 - A salary of a bank employee must be at least \$4.00 an hour
 - A customer must have a (non-null) phone number

On the left margin, there is vertical text: "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Ch. BT Kumar aggrawal - Jatin Agarwal". At the bottom, it says "Database System Concepts - 8th Edition" and "10.6".

Integrity constraints guard against accidental damage to the database that is there are certain real world facts that must be ensured in the database all the time. For example, in bank accounts, we have a minimum balance that need to be maintained, for a particular customer we might want that the customer's phone number must be present. We may have certain age bar in terms of entering into certain memberships or certain employment and so on. All these kinds of real world constraints need to be represented and maintained in the database and that is the purpose of the integrity constraint.

(Refer Slide Time: 09:16)

The slide is titled "Integrity Constraints on a Single Relation" in red. It contains a bulleted list of four types of integrity constraints:

- not null
- primary key
- unique
- check (P), where P is a predicate

On the left margin, there is vertical text: "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. P. Ch. BT Kumar aggrawal - Jatin Agarwal". At the bottom, it says "Database System Concepts - 8th Edition" and "10.7".

And we will first look at the issues of integrity constraint for a single relation, and we have seen the use of NOT NULL and PRIMARY KEY. We will also talk about what is UNIQUE and we will see how actually general constraints can be checked in terms of a CHECK clause with a predicate P.

(Refer Slide Time: 09:43)

The slide has a title 'Not Null and Unique Constraints' in red. On the left, there is a vertical sidebar with text: 'SWAYAM: NPTEL-NOC-MOOCs Instructional IP Course: ET-Kharagpur: Jain-Apr-2016'. Below this is a video window showing a man speaking. The main content area has two columns of bullet points:

- not null
 - Declare *name* and *budget* to be not null


```
name varchar(20) not null
budget numeric(12,2) not null
```
- unique (A_1, A_2, \dots, A_m)
 - The unique specification states that the attributes A_1, A_2, \dots, A_m form a candidate key
 - Candidate keys are permitted to be null (in contrast to primary keys).

To the right of the text is a hand-drawn diagram of a table with rows labeled t_1 and t_2 . The columns are labeled A_1, A_2, \dots, A_m . Arrows point from the labels to the corresponding columns in the table rows.

So, NOT NULL. we had seen this before we can while creating the database table. Create table we can specify a field to be NOT NULL and then in that case null values will not be allowed in those fields. So, they will must be some value given. You can say one or more attributes to be UNIQUE. If you specify them to be UNIQUE; that means, that in any instance of the table in future that cannot be two tuples which match in all of those attribute.

So, if you say A_1 to A_m are UNIQUE; that means, that if we have two different tuples in the table anytime in future. Two different rows in the table t_1 and t_2 then across A_1, A_2, \dots, A_m they must differ in at least one attribute value. So, uniqueness is a basic requirement for being a CANDIDATE KEY, but they are, but still permitted to be null which in contrast to what is the true for PRIMARY KEY already you know that PRIMARY KEY values cannot be null, but uniqueness allows a null value, but they have to be different.

(Refer Slide Time: 10:58)

The slide is titled "The check clause". It contains the following text and code:

- check (P)
where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (
    course_id varchar (8),
    sec_id varchar (8),
    semester varchar (6),
    year numeric (4,0),
    building varchar (15),
    room_number varchar (7),
    time_slot_id varchar (4),
    primary key (course_id, sec_id, semester, year),
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))
);
```

A blue arrow points from the word "P" in the first bullet point to the "check" keyword in the SQL code.

The CHECK clause is where you say CHECK and then you put a predicate. So, the idea is like this, suppose I know that I have specified an attribute called semester. It is a varchar 6 which means that it can have a string maximum of length 6, but naturally I can write anything there I can write morning in that field. I can write welcome in that field and so on, but those are not valid names of semester. So, I want that in my design semester must have any of these values only.

So, we say semester in. So, I have listed the values that are allowed in is a set membership. So, it says the semester in so which means the value of the semester is be one of this fall. And this whole thing now becomes the predicate P on which I give a check which means that, whenever I am creating once, you have created the table. when I want to insert or update the values in the table, the records in the table; the value of semester has to be always within this. Otherwise, the CHECK integrity constraint will fail, and the update or insert will not happen and an exception will be raised. This is the basic idea of CHECK constraints.

(Refer Slide Time: 12:25)

The slide is titled "Referential Integrity". It includes the following text and a diagram:

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation
- Example: If "Biology" is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for "Biology"

Hand-drawn diagram illustrating referential integrity:

- A diagram showing two tables: "Instructor" and "Department".
- The "Instructor" table has columns labeled "InstID" and "DeptName".
- The "Department" table has columns labeled "DeptID" and "DeptName".
- A blue arrow points from the "DeptName" column in the "Instructor" table to the "DeptName" column in the "Department" table.
- A handwritten note "InstID" is written above the first column of the "Instructor" table.
- A handwritten note "DeptName" is written above the second column of the "Instructor" table.
- A handwritten note "DeptID" is written above the first column of the "Department" table.
- A handwritten note "DeptName" is written above the second column of the "Department" table.

Now, let us move on to more involved integrity CHECK which goes beyond one table. So, let us suppose that we are talking about the instructor table we have the instructor table. An instructor table has a department name. Similarly, we have a department table which naturally has a department name. Now, we know that this is the key in the department table and therefore, here it is the foreign key.

Now, while we are inserting records in the instructor table how do we guarantee that the record that we insert has a corresponding entry in the foreign key table that is difference table. So, it is when we are inserting and in a faculty in the saying that the faculty belongs to biology department there needs to be a biology entry in the department table as well. So, this is known as a referential integrity that is once you refer from one table to the other that reference must also be a valid one; otherwise, all your computations will go wrong.

(Refer Slide Time: 13:48)

The slide has a title 'Referential Integrity' in red. Below it is a bulleted list:

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation
 - Example: If "Biology" is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for "Biology"
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S

Navigation icons and a progress bar are visible at the bottom.

So, this is a for saying it, formally there are two relations and one relation as a PRIMARY KEY which is used in the other relation as a foreign key then there is a referential integrity that needs to be maintained.

(Refer Slide Time: 14:02)

The slide has a title 'Cascading Actions in Referential Integrity'. Below it is a bulleted list of SQL code snippets:

- ```
* create table course (
 course_id char(5) primary key,
 title varchar(20),
 dept_name varchar(20) references department
)
```
- ```
* create table course (
    ...
    dept_name varchar(20),
    foreign key (dept_name) references department
        on delete cascade
        on update cascade,
    ...
)
```
- alternative actions to cascade: **no action, set null, set default**

Navigation icons and a progress bar are visible at the bottom.

So, here we are just showing the effect of that. So, we have created a table this is the first one is what you have seen earlier creating the table course and that table course needs the name of the department. So, we are specifying that it references the department table. Now, if it references the department table, it must ensure the referential integrity. So, this

just says that this refers to the department table, but I can be more specific to say what will happen if the integrity gets violated. For example, I have created this and the course table has an entry, which has a department name say biology. Naturally, biology department should have the entry in the department table with this department name biology for this to be valid.

Now, say for some reason the biology department is abolished and that particular record from the department table is removed, naturally, the course which is referring to biology in terms of its department name that particular record will become invalid. So, we can say that on delete, what you should be doing, one most common action that we specify in referential integrity is cascade that if the referred entity is deleted then the referring entity should also be deleted. So, if you delete the biology entry from the department table, then all courses which have biology through references to department as their field value should also get deleted.

Similar thing can be there on update also. For example, biology department say tomorrow changes the name to bioscience. Now, if I have a referential integrity put on the course table as on update cascade, then as I change the bioscience the name to bioscience. All records in the table course, which had the department name as biology will necessarily get updated. So, this is the way to maintain referential integrity. Cascading is one of the most common way to handle this, but there could be other ways to take action also that could be no action that you say ok, I do not care. Let that happen in that case, because of the violation that could be some exceptions thrown or even say that if this happens then I will set that field to null or I will set that field to some default value and so on.

(Refer Slide Time: 17:01)

The slide is titled "Integrity Constraint Violation During Transactions". It contains the following SQL code:

```
create table person (
ID char(10),
name char(40),
mother char(10),
father char(10),
primary key ID,
foreign key father references person,
foreign key mother references person)
```

Below the code is a bulleted list:

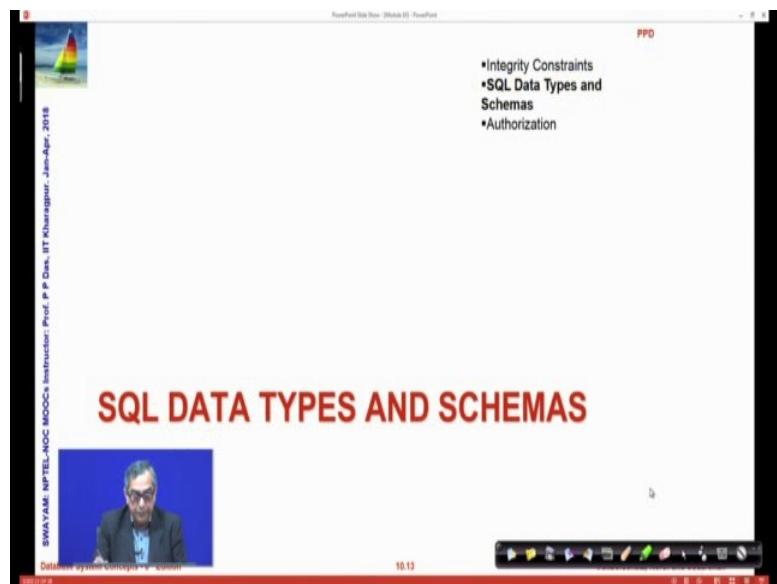
- How to insert a tuple without causing constraint violation ?
 - insert father and mother of a person before inserting person
 - OR, set father and mother to null initially, update after inserting all persons (not possible if father and mother attributes declared to be **not null**)
 - OR defer constraint checking (will discuss later)

At the bottom left, it says "DRAFT: NPTEL-NOCO-MOCOCs Instructional IP-Draft, IIT-Kharagpur - Jain-Agarwal" and "Database System Concepts - 6th Edition". At the bottom right, it shows "10.12" and some navigation icons.

So, this is how the referential integrity has to be handled. There could be integrity violation during transactions also. This is an example of a self-referential table which of persons which where every person's entry needs the name of the mother and the father which are also entries in this table. So, necessarily if you are entering a person record you need this feels to be populated and that can be populated only if those records already exists.

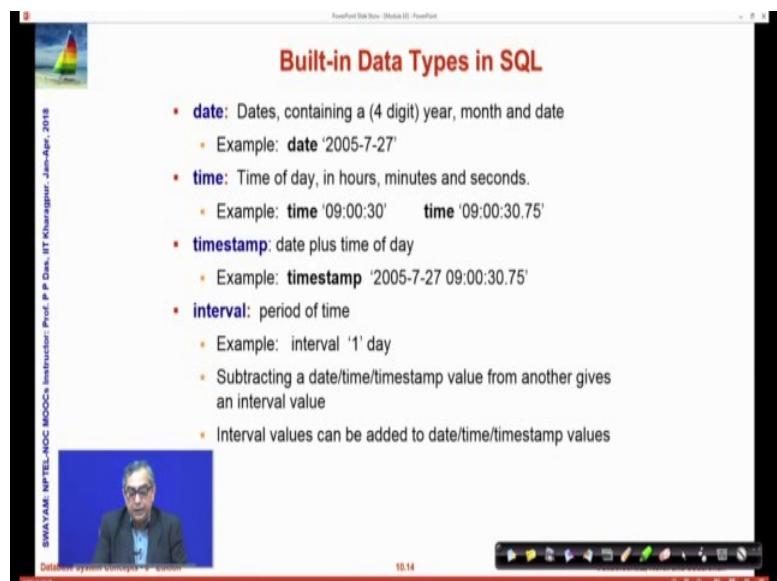
So, there is some order in which you have to enter the records or you have to set them as null and then update them in future and or some ways to say that well do not check this integrity now. We will talk about this integrity at a later point of time. So, these are the issues that necessarily you will have to be addressed.

(Refer Slide Time: 17:54)



Let us move on and look at the SQL data types and schemas.

(Refer Slide Time: 18:02)



So, in addition to the data types like char, varchar, int and all that you have an explicit date data type which gives you a year, month, date kind of format with a four-digit year. because date is very frequently required, you have a time type to give you hour, minute, second time format. You have a timestamp, which is date and time together and you have what is known as interval where you can do a date or time difference between two different dates, two different time, two different time stamps and so on. So, these are the

common added built in types which makes it very easy to handle the temporal aspects in SQL queries.

(Refer Slide Time: 18:54)

create table student
(*ID* varchar (5),
name varchar (20) not null,
dept_name varchar (20),
tot_cred numeric (3,0) default 0,
primary key (*ID*))

- create index studentID_index on student(*ID*)
- Indices are data structures used to speed up access to records with specified values for index attributes

```
select *  
from student  
where ID = '12345'
```

- can be executed by using the index to find the required record, without looking at all records of student
- More on indices in Chapter 11

In addition, the next that you can do is you can create an index. So, let us look at this. So, this create table definition you understand well by now. You can do this, I can say create index and give a name for the index and specify which field on which the index should be created. So, here we are saying that the index should happen here. This is the name of the relation; this is the name of the attribute name of the field. Now, this does not change any data neither does it change any schema, but it creates certain additional structure so that it becomes easier to search this particular table using IDs.

So, if I have a query like this that I am trying to find out all information about a particular student then as we have said that by default the different entries the rows of a relation are unordered. So, the only way to find out this particular row and in fact, whether it actually exist would be to go over all the relations one by one. But if we index it, then it creates some kind of an efficient data structure through which it can be searched out very efficiently very easily. with a later module we will talk about indexing. But just to give you the idea that this is similar to finding out a value in an unordered array if you are thinking of C.

In contrast, we all know that this can be done, but takes a whole lot of time it takes order in time. But I could keep those numbers in terms of some binary search tree, balanced

binary search tree like red black tree or two three four tree kind of, where the search can be conducted in a login time. Or I could keep it in terms of some efficient hashing mechanism where the search could happen in terms of an ordered one time also. So, indexing has a lot of importance and we will talk about that more, but this is how you create index in SQL.

(Refer Slide Time: 21:18)

The screenshot shows a PowerPoint slide with the title 'User-Defined Types' in red at the top. Below the title, there is a bulleted list of SQL statements. A blue oval has been drawn around the word 'Dollars' in the first statement, and a blue wavy underline has been applied to the word 'Dollars' in the second statement. The first statement is:

```
create type Dollars as numeric (12,2) final
```

The second statement is:

```
create table department  
(dept_name varchar (20),  
building varchar (15),  
budget Dollars);
```

At the bottom left of the slide, there is a small video window showing a man speaking. The video window has a caption that reads 'SWAYAM: NPTEL-ANOC-MOOCs Instruction: Prof. P. Das, IIT-Kharagpur'. The video player interface shows '10.18' in the bottom right corner.

You can have user defined types, you can say create type and use some specific. you know sub types of a type as and give it a name. So, you send numeric 12, 2. So, which is the 12 digit number with two decimal places of precision. you can call it dollar and then use that as a type name. So, type name doing this helps in to visit make sure that wherever you actually have to conceptually refer to dollars you are talking about dollars it is easier to understand and you are making sure that everywhere the same numeric precision is used.

(Refer Slide Time: 22:02)

The screenshot shows a PowerPoint slide with the title 'Domains'. The slide contains the following text:

- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person_name char(20) not null
```

Handwritten annotations include a blue wavy line under 'not null'.

- Types and domains are similar
- Domains can have constraints, such as **not null**, specified on them

```
create domain degree_level varchar(10)
constraint degree_level_test
check (value in ('Bachelors', 'Masters', 'Doctorate'));
```

Handwritten annotations include a blue wavy line under 'degree_level' and another under 'check'.

A video player window is visible at the bottom left, showing a man speaking. The video player interface includes controls like play, stop, and volume, and a timestamp of 10.17.

You can also actually go further and create domains which is very similar to create type. But domains are more powerful in the sense that, in a domain, you can also add constraints like NOT NULL and you say that this is person name, say that this once you have set that this person name is 20 characters long and it cannot be null then you do not specifically have to every time.

You define a field based on this domain type you do not have to specifically say that it is NOT NULL. You could also create a specific constraint in terms of the CHECK clause and make it easier. So, now if you say degree level you do not have to put CHECK clause explicitly in the SQL query, because it is already specified in the created domain.

(Refer Slide Time: 22:52)

The slide is titled "Large-Object Types" in red. It contains a bulleted list:

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*.
 - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
 - **clob**: character large object -- object is a large collection of character data
 - When a query returns a large object, a pointer is returned rather than the large object itself

Navigation icons at the bottom include back, forward, and search symbols.

SQL supports certain large-objects which are either called blobs if they are binary, or called clob if they are character objects. The only the major difference in terms of the large object types are they are not stored as a part of the table. They stored elsewhere and you actually maintain a kind of a reference a pointer to that large object. So, this is very useful in terms of handling photos, videos and no big binary files, character files also.

(Refer Slide Time: 23:21)

The slide has a red header "PPD". Below it, there is a bulleted list:

- Integrity Constraints
- SQL Data Types and Schemas
- Authorization

The main title "AUTHORIZATION" is displayed in large red capital letters. Navigation icons at the bottom include back, forward, and search symbols.

Let us move to authorization.

(Refer Slide Time: 23:29)

The screenshot shows a PowerPoint slide with the title 'Authorization' in red at the top center. The slide content is organized into two main bullet points:

- Forms of authorization on parts of the database:
 - **Read** - allows reading, but not modification of data
 - **Insert** - allows insertion of new data, but not modification of existing data
 - **Update** - allows modification, but not deletion of data
 - **Delete** - allows deletion of data
- Forms of authorization to modify the database schema
 - **Index** - allows creation and deletion of indices
 - **Resources** - allows creation of new relations
 - **Alteration** - allows addition or deletion of attributes in a relation
 - **Drop** - allows deletion of relations

Handwritten notes are present on the right side of the slide:

- A bracket groups the first four items under 'Read', 'Insert', 'Update', and 'Delete' with the handwritten note 'DBA Prof.'.
- A bracket groups the last three items under 'Index', 'Resources', 'Alteration', and 'Drop' with the handwritten note 'Analysts'.

Other slide details include:

- Left margin: 'SHIVAYAM: NPTEL MOOC Instructor: Prof. P. Das, IIT Kharagpur - Jatin-Das'
- Bottom left: 'Database System Concepts - 4. Control'
- Bottom center: '10.20'
- Bottom right: Standard Mac OS X window controls

Next, authorization is the process by which you restrict different users to be able to do different kind of operations. You recall in the early modules on database overview. We mentioned that there could be several types of users for a database. There could be absolutely application users who may necessarily do not feature as a part of the database development. But there could be application developers expectedly most of you would become application developers or there could be intermediate higher level of analyst who design databases, design constraints, decide on indices and so on and that could be database administrator.

And also in terms of different application programs and programmers there is a need to separate out who can access which part of the database. For example, if you look at a add a banking system, then while I am, by net banking application is accessing different information about my account, one part I need to ensure that I can only access my account.

And also what the database system needs to ensure is that a net banking application should in no way be able to access the information about the specific employees, because, in the same database information about the bank employees will also be there. It should not be possible for possible to access information about different physical information about the branches as to where, how many square feet of area that branch has and so on and so forth.

So, we need to put variety of restrictions and as we will see that authorisation or this process of restricting or allowing different access and different authority to operate is decided based on two different factors. One is what you want to do, and two is who wants to do that. So, what and who, so we identify different operations or different operations on certain tables or operations on certain attributes as what needs to be done. And on the other side, we will identify who in terms of specific individual user IDs or groups of user IDs or roles that exist.

So, here we will just try to show you how we can do that in SQL. So, the first part of the authorisation is being able to do different things with the database that means, the instances of the database. So, there are authorisations to read insert, update and delete. So, read is where you can access the data, but you cannot modify; insert is when you can add new data, but you do not insert rights authorisation, you cannot update an existing data, you can only insert data. You can have update rights, where we can change make modifications, but you may not be, you are not allowed to delete data, and you can that would be delete right where it allows you to delete data and my new this authorisations or not these are all independent authorisation.

So, you may have I mean certain authorisations may need certain other authorisations to be present. For example, if you are updating naturally evenly to read, but it is these are all independent authorisations, and you may have one or more of them to be able to do the appropriate actions. Similar set of another set of authorisations will exist, if you want to for those want to modify the database schema, naturally, this is primarily for the applications and application programmers, and this primarily would be for the analysts that you can index the different table you can do.

You can have authorisation for resources which mean you can create new relations, create new schemas, you can alter schemas, you can drop schemas and so on. So, these are the different kinds of authorisation that are possible.

(Refer Slide Time: 28:01)

The screenshot shows a PowerPoint slide with the title 'Authorization Specification in SQL' in red at the top. Below the title is a bulleted list of points about the GRANT statement:

- The **grant** statement is used to confer authorization
 - grant <privilege list>
 - on <relation name or view name> to <user list>
- <user list> is:
 - a user-id
 - **public**, which allows all valid users the privilege granted
 - A role (more on this later)
- Granting a privilege on a view does not imply granting any privileges on the underlying relations
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator)

At the bottom left of the slide, there is a small video player showing a man speaking, with the text 'Database System - Lecture 8' and '10.21'. The slide has a watermark on the left side that reads 'SWANAMI_NPTEL_ANOCA_Instructional_Protocol_P_01_Cap_1_SQL_Authorization_June-2016'.

So, let us see how it works the authorisation is specified in terms of a statement called grant. So, you grant an authorisation to a privilege list and on certain relation to a group of users. So, grant what kind of authorisation you are granting that is the privilege less list on what relation on view you are granting that is on condition and to whom are you granting those. So, user list could be a specific user ID or you could say public which in this case everybody will have that or this could be a role which will see, what a role is.

Granting a privilege on a view does not imply granting any privileges on the underline relation, please mind this one, because, you have seen that a view can be formed from multiple different relations. So, if somebody has been granted a particular privilege say read privilege on a view, then it does not mean that the corresponding underline relationship, you been granted a read privilege on faculty relation that we faculty view that we did that does not mean that the user will automatically get a read privilege on the underline in structure relation. So, that has to be kept in mind.

The grantor of the privilege must already hold the privilege that is you cannot grant. Naturally, grant will be done also by somebody in some of the users you may be, so that user who is granting must also have the privilege, same privilege on the specific item. So, you cannot grant privilege on something, some relation or view on which you yourself do not have that or it has to be the database administrative who naturally has privilege for everything.

(Refer Slide Time: 30:13)

The screenshot shows a PowerPoint slide with the title 'Privileges in SQL' in red at the top center. Below the title is a bulleted list of SQL privileges:

- **select**: allows read access to relation, or the ability to query using the view
 - Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *instructor* relation:
`grant select on instructor to U_1 , U_2 , U_3`
- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges

At the bottom left of the slide, there is a small video player showing a man speaking, with the text 'SWANAMI_NPTEL_ANOCC_MOCOCs_Instructor: Prof. P. Das, IIT Kharagpur' and 'Database system (University of Calicut)'. The video player has a timestamp of '10.22'.

The privileges on SQL there is a SELECT privilege, which is so you know this is the privilege list. SELECT privilege which basically means read access. So, it is saying SELECT privileges is read access on instructor and these are the different users. So, this is how typically, you can have insert to ability to insert tuples, the update privilege this should be easy to understand now delete privilege.

So, only thing is read is a called SELECT here, so these are the different privileges that you have in a SQL. And you have one all in compassing privilege which is called all privileges so as a short form of allowing all these allowable privileges.

(Refer Slide Time: 31:00)

The screenshot shows a PowerPoint slide with the title 'Revoking Authorization in SQL' in red at the top center. Below the title is a bulleted list of points about the REVOKE statement:

- The **revoke** statement is used to revoke authorization
revoke <privilege list>
on <relation name or view name> from <user list>
- Example:
revoke select on branch from U₁, U₂, U₃
- <privilege-list> may be all to revoke all privileges the revoker may hold
- If <revoker-list> includes **public**, all users lose the privilege except those granted it explicitly
- If the same privilege was granted twice to the same user by different grantors, the user may retain the privilege after the revocation
- All privileges that depend on the privilege being revoked are also revoked

At the bottom left of the slide, there is a vertical watermark-like text: 'DATA ANALYSIS NPTEL-NOCO-MOCOCs Instructional IP-Draft, IIT-Kharagpur - Jain-Agarwal - June-Aug-2016'. At the bottom right, there is a small navigation bar with icons.

Certainly, if you can grant an authorization, then needs to be a reverse process that is if we want to withdraw authorization of certain privileges on certain items for certain users, so that is known as a REVOKE statement. So, we can revoke it the structure looks exactly similar to the grant. So, you REVOKE a privilege list, SELECT, insert this kind of on certain relation and view from a set of users. So, you can say that REVOKE SELECT ON branch from this. So, once this is done, then U₁, U₂ and U₃ will not be able to read the branch relation or the branch view. The privilege list may be all to REVOKE all privileges. So, instead of revoking SELECT, insert separately, you can just say all and revoke all of that.

The list of revoking can include public which means all users lose that privilege and or though those were granted. If the same privileges granted twice, now it is possible that a user gets privilege granted to him or her by two different granting authorities, then if ones is one of them is revoked the other will still continue to remain; So, every privilege that is granted needs to be explicitly revoked that is a basic meaning.

So, all privileges that depend on the privilege being revoked are also revoked. So, some privilege which is dependent on some other privilege, if you REVOKE the update privilege then this SELECT privilege will remain. But, if you REVOKE the SELECT privilege then if you also had the update privilege that will certainly get revoked, because if you cannot read then naturally you cannot change.

(Refer Slide Time: 33:01)

The slide is titled "Roles". It contains the following bulleted list:

- `create role instructor;`
 `grant instructor to Amit;`
- Privileges can be granted to roles:
 `grant select on takes to instructor;`
- Roles can be granted to users, as well as to other roles
 `create role teaching_assistant`
 `grant teaching_assistant to instructor;`
 • Instructor inherits all privileges of *teaching_assistant*
- Chain of roles
 `create role dean;`
 `grant instructor to dean;`
 `grant dean to Satoshi;`

The SQL also allows you to create certain roles. Roles are kind of like virtual use that so, we all say that we all play certain role. So, I have an entity as an individual say I may be user called PPD, but I have a role as an instructor, I have a role as say the head of the department, I have a role as a chairman of committee and so on. So, often times it becomes easier to grant privileges to different roles.

You do not really care immediately about who that individual could be who that particular user could be who has that privilege, whoever plays that role gets that privilege, whoever becomes the director of IIT Kharagpur has the privilege to appoint faculty members it is of that kind. It does not specifically. So, role is of that kind of a concept. So, you CREATE ROLE we are saying that role is the role instructor is created.

And then you are saying that you grant instructor to Amit which says that Amit now plays the role of instructor. So, any privilege that the instructor role has Amit will enjoy that. So, let us see more of this the privileges can be grant to roles. So, earlier we said it could be public, it could be users, but now you are saying that it could be two roles. So, here this role was creates and the privilege is being granted to that. And since Amit plays that role it will mean that Amit with this GRANT SELECT ON takes to instructor, Amit will actually get a privilege of SELECT on takes relation that is the kind of derived structure that roles give you roles can be granted to users as well as to other roles.

So, roles are becoming like virtual users. So, you can CREATE ROLE teaching assistant and grant teaching assistant to instructor, which means that if you do that you are granting this. So, it means that any privilege the teaching assistant will have, the instructor will get those privileges, because, you have made in instructor to also play the role teaching assistant mind you instructor itself is virtual entity. So, if Amit is an instructor by this, then Amit plays this role and this role plays teaching assistant role and this teaching assistant role has certain privileges, so naturally through this chain process Amit will get those privileges. So, in an instructor inherits all privileges of teaching assistant.

So, this is what exactly, what I was talking of you can have a chain of roles create a role dean new one, you have created, then grant instructor to dean grant dean to Satoshi. So, which means; that once you grant dean to instructor; so anybody who plays the dean's role will get all privileges of instructor. Here you are saying that Satoshi is going to play the dean role. So, the Satoshi in terms of chaining gets all the privileges that instructor has.

(Refer Slide Time: 36:41)

Authorization on Views

```

create view geo_instructor as
(select *
from instructor
where dept_name = 'Geology');
grant select on geo_instructor to geo_staff

```

- Suppose that a geo_staff member issues

```

select *
from geo_instructor;

```

- What if
 - geo_staff does not have permissions on instructor?
 - creator of view did not have some permissions on instructor?

So, once this has been done, then you can have authorization on views as well. So, you have created a view here. So, this is the view created the geo instructor the Geology instructor. And on that view particularly you have given the privilege to the geo staff. So, a geo staff member would be able to access this view. And if this query is fired by a geo

staff member, which I am assuming is a role then this view will get executed and the results of all instructors in the department geology will be update. But, what if the geo staff does not have permission on instructor, does not matter that is the beauty of the whole thing. The geo staff may not have permission to do SELECT on instructor, but the geo staff has permission to SELECT on geo instructor. So, the geo staff will be able to execute this view, but the geo staff will not be able to do a SELECT on from the instructor database instructor table.

(Refer Slide Time: 38:16)

Other Authorization Features*

- references privilege to create foreign key
 - grant reference (*dept_name*) on *department* to Mariano;
 - why is this required?
- transfer of privileges
 - grant select on *department* to Amit with grant option;
 - revoke select on *department* from Amit, Satoshi cascade;
 - revoke select on *department* from Amit, Satoshi restrict;

There are several other authorization features, the references a privilege to create foreign key. So, we talked about basic read, write, data manipulation privileges, but there could be other privileges like whether you can create a foreign key, whether you can transfer of privilege. So, whether you can give one privilege to another, so whether you can cascade, whether you can restrict and so on.

So, transfer of privileges is also privilege. I am naturally will have to think of this is an authorization. So, actually what we can authorize is also a privilege that needs to be authorized. So, these are the derived privileges that I have.

(Refer Slide Time: 39:00)

The screenshot shows a Microsoft PowerPoint slide titled "Module Summary". The slide content includes a bulleted list of three items:

- Learnt SQL expressions for integrity constraints
- Familiarized with more data types in SQL
- Discussed authorization in SQL

At the bottom left of the slide, there is a vertical footer text: "DATA MAINTAINABILITY: INTEGRITY CONCEPTS", "Database System Concepts - 6th Edition", and "Prof. Dr. P. Gajjar, IIT Kharagpur". At the bottom right, it says "10.27". The slide has a standard Microsoft Office ribbon at the top and a toolbar at the bottom.

So, in summary, we have learnt about SQL expressions to deal with integrity constraints. We are familiarized with more data types particularly user defined types and domains creation of index and we have discussed about authorization in SQL. Each one of them particularly authorization has lot more details, but at this intermediate level we just wanted to get a basic idea about authorization to be able to deal with that.

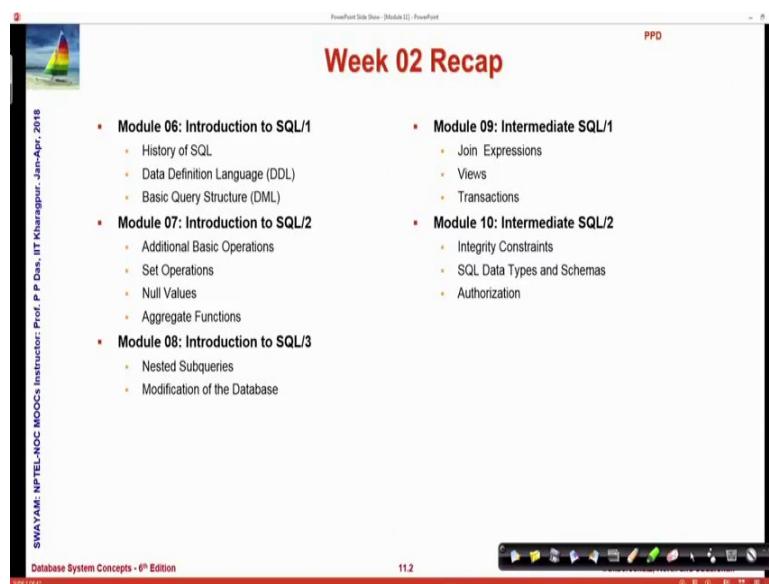
So, with this we close our discussion on the intermediate level SQL features. In the next module that we start next week, we will talk about some of the advanced SQL features.

Database Management System
Prof. Partha Pratim Das
Department of Computer Science & Engineering
Indian Institute of Technology, Kharagpur

Lecture - 11
Advanced SQL

Welcome to module eleven of database management system. This will be on advanced SQL.

(Refer Slide Time: 00:31)



The slide is titled "Week 02 Recap" in red at the top right. It features a small sailboat icon in the top left corner. The content is organized into two columns of bullet points:

<ul style="list-style-type: none">▪ Module 06: Introduction to SQL/1<ul style="list-style-type: none">▫ History of SQL▫ Data Definition Language (DDL)▫ Basic Query Structure (DML)▪ Module 07: Introduction to SQL/2<ul style="list-style-type: none">▫ Additional Basic Operations▫ Set Operations▫ Null Values▫ Aggregate Functions▪ Module 08: Introduction to SQL/3<ul style="list-style-type: none">▫ Nested Subqueries▫ Modification of the Database	<ul style="list-style-type: none">▪ Module 09: Intermediate SQL/1<ul style="list-style-type: none">▫ Join Expressions▫ Views▫ Transactions▪ Module 10: Intermediate SQL/2<ul style="list-style-type: none">▫ Integrity Constraints▫ SQL Data Types and Schemas▫ Authorization
--	--

At the bottom left, there is vertical text: "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P P Das, IIT Kharagpur, Jan-Apr., 2018". At the bottom center, it says "Database System Concepts - 6th Edition". The bottom right corner shows slide number "11.2".

Before we start let me quickly recap what we did last week in the five modules. Last week we totally spent on discussing first the introductory features of SQL how to create data and how to write basic queries, and we studied about all different kinds of SQL operations at theoretic operation, handling of null values aggregation, nested queries and so on. And then we did an intermediate level of SQL query formation in terms of joint expression, views and integrities different kinds of SQL data types and importantly authorization.

(Refer Slide Time: 01:24)

The slide is titled "Module Outline" in red. It contains a bulleted list of topics:

- Accessing SQL From a Programming Language
- Functions and Procedural Constructs
- Triggers

At the bottom left, it says "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. Das, IIT Kharagpur, Jan-Apr., 2018". At the bottom center, it says "Database System Concepts - 6th Edition". At the bottom right, it shows slide number 11.4.

In the context, in this context, we now take up some more of the SQL features which are somewhat advanced. And we will try to understand how SQL can be used from a programming language, and familiarize with functions and procedures in SQL. This will violate some of the basic premises that we started with in saying that SQL is a declarative language only because as you understand functions procedure as procedural language features we will see how to handle those, and we will take a look into another important feature of triggers.

(Refer Slide Time: 02:07)

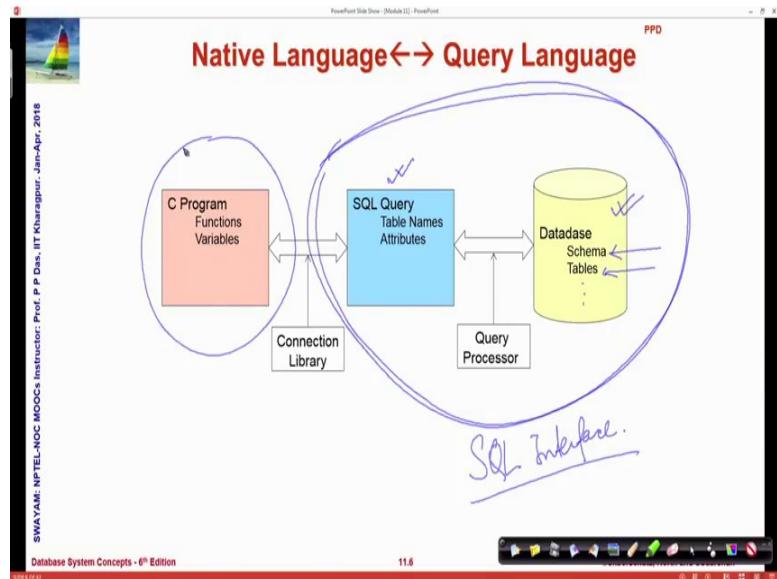
The slide has a large red title "ACCESSING SQL FROM A PROGRAMMING LANGUAGE" at the top. Below it is a bulleted list of topics:

- Accessing SQL From a Programming Language
- Functions and Procedural Constructs
- Triggers

At the bottom left, it says "SWAYAM: NPTEL-NOC MOOCs Instructor: Prof. P. Das, IIT Kharagpur, Jan-Apr., 2018". At the bottom center, it says "Database System Concepts - 6th Edition". At the bottom right, it shows slide number 11.5.

First with accessing SQL so, this is the module outline accessing SQL from a programming language.

(Refer Slide Time: 02:11)



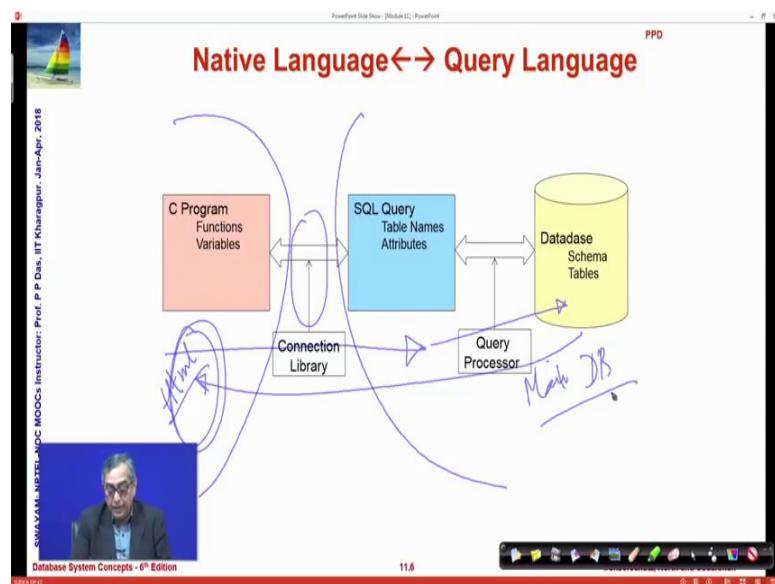
So, this is just kind of an abstract view that you can think of that what we have been doing so far naturally there is a database which has primarily two things schema and tables. Of course, there are many other things there are index, there is authorization that triggers and all that, but there is a database which stores everything.

And so far we have been dealing with only this part of the information that I can write certain SQL query to define table, using table names, attributes and certain logic and that goes through certain query processor works on the database to get me either create the desired effect. Either that is creating table instances or creating index or extracting certain relation values, defining some views and so on. So, all these have to happen through an SQL interface. So, this has to happen through an SQL interface.

So, whatever system we are using whether you are using MySQL, Postgres, Oracle, Sybase SQL server whatever you using that has some interface through which these SQL queries can be executed, so that is kind of a standalone complete system. But, in general, we will have lot of more requirements in terms of the application. For example, the application might require some graphics, SQL does not have support for graphics; application might require certain numerical algorithms to be executed might require some geometric computations to be done poly intersection of polygons to be computed

and so on and so forth. It might require some network programming and so on. So, there is a need to do all these and certainly these are best done in terms of certain native language that could be C++, java, C#, visual basic, python any of the native languages which has support for a variety of different tasks.

(Refer Slide Time: 04:20)



So, what is critical is can we make a bridge between these two that is can we take the advantages of the SQL domain and the advantages of the normal imperative procedural programming domain together. That is I can do some graphic representation and then certainly go to the database extract some information and then present it in the graphics and vice versa. For example, whenever we access say gmail, we get to see them in terms of an html presentation which is some kind of a graphics rendering, but the all the mail entry certainly come from some database.

So, somewhere there is a connection by which when I say get my inbox mails and somewhere the information goes over from this to the SQL query up to the database and the result is brought back to me. I see that in the html page here, but here there must be certain mail database, where all these information exists. So, what I am trying to come at is it is critical that the application programs or you know high level programming normal programming languages native languages should be able to interface with SQL. And what we discuss here is two different mechanisms for this interfacing. The first one is using a connection library and this is what I will specifically show you.

(Refer Slide Time: 05:50)

Accessing SQL From a Programming Language

- API (application-program interface) for a program to interact with a database server
- Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- Various tools:
 - JDBC (Java Database Connectivity) works with Java
 - ODBC (Open Database Connectivity) works with C, C++, C#, Visual Basic, and Python
 - Other API's such as ADO.NET sit on top of ODBC
 - Embedded SQL

Database System Concepts - 6th Edition

11.7

So, for if you are using a connection library then you have a set of APIs application programming interfaces these are basically functions in that library. So, that with that application can connect to the database because certainly the database is somewhere else is a different server. And send an SQL command to the database server so that you can say that this is what I want. And then a result is computed the result of that computation the table the can be brought back those tuples can be brought back to the application program. Mind you when we are here when you are sending the information in terms of the database server, we are talking SQL command, we are talking about attributes tables of the SQL space.

Whereas, when I want it in the program I want it in terms of program variables. So, there has to be certain correspondence made between them. There are a variety of tools available which allow you to do this JDBC is common very commonly known which is specific for java. We have an open database connectivity APIs which has different versions for different languages these are the common languages that it is used with. And as we will see later on that there is another mechanism for doing the same thing called the embedded SQL, we will come to that later.

(Refer Slide Time: 07:15)

The screenshot shows a PowerPoint slide with the title 'JDBC' in red at the top center. Below the title is a bulleted list of points about JDBC:

- JDBC is a Java API for communicating with database systems supporting SQL
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes
- Model for communicating with the database:
 - Open a connection
 - Create a "statement" object
 - Execute queries using the Statement object to send queries and fetch results
 - Exception mechanism to handle errors

At the bottom left of the slide, there is a small video window showing a man speaking. The video window has a blue background and the text 'Database System Concepts - 6th Edition' at the bottom. The overall interface is a standard Windows desktop with a taskbar at the bottom.

So, JDBC is a java API, I will not go into details here if you know java you should be able to quickly look up. So, it is a it communicates with the database by opening a connection creating what java calls a statement object and executes the query using the statement objects and that is used to send the query as well as to get back the result. So, java is object oriented as you know so statement object is used as a as an encapsulation which travels between the java program and the SQL query processor. And since the query gets back the result since the query gets back there is that kind of there is exception mechanism to handle errors which is common for java.

(Refer Slide Time: 08:00)

The screenshot shows a PowerPoint slide with the title 'ODBC' in red at the top center. Below the title is a bulleted list of points about ODBC:

- Open DataBase Connectivity (ODBC) standard
 - standard for application program to communicate with a database server
 - application program interface (API) to
 - open a connection with a database,
 - send queries and updates,
 - get back results
- Applications such as GUI, spreadsheets, etc. can use ODBC

At the bottom left of the slide, there is a small video window showing a man speaking. The video window has a blue background and the text 'Database System Concepts - 6th Edition' at the bottom. The overall interface is a standard Windows desktop with a taskbar at the bottom.

What you see in contrast does a similar thing, but since it is to cater for different languages it has got a softer model it has got less powerful model. It is a standard application program to communicate with database server. So, again it has to open a connection to the database, send queries and updates and get back results. So, these are these are the three basic things, these three other three basic things that certainly needs to be done if I want to easily work across the application programming domain and the database programming domain. So, applications such as GUI, spreadsheet, etcetera can use ODBC.

(Refer Slide Time: 08:39)

ODBC – Python Example

- The code uses a data source named "SQLS" from the odbc.ini file to connect and issue a query.
- It creates a table, inserts data using literal and parameterized statements and fetches the data

```

import pyodbc

conn = pyodbc.connect('DSN=SQLS;UID=test01;PWD=test01')
cursor=conn.cursor()
cursor.execute("create table rvtest (col1 int, col2 float,
col3 varchar(10))")
cursor.execute("insert into rvtest values(1, 10.0,
'ABC')")
cursor.execute("select * from rvtest")

while True:
    row=cursor.fetchone()
    if not row:
        break
    print(row)

cursor.execute("delete from rvtest")
cursor.execute("insert into rvtest values (?, ?, ?)", 2,
20.0, 'XYZ')
cursor.execute("select * from rvtest")

while True:
    row=cursor.fetchone()
    if not row:
        break
    print(row)

```

Source: <https://dzone.com/articles/tutorial-connecting-to-odbc-data-sources-with-python>

So, yeah I am just quickly showing you an example we will talk about application programming mode in a later module. So, this is a python example. So, python for this the ODBC for python is known as pyodbc library. So, you need to import that. And then using that you can connect. So, if you have to connect you have to say which database who is the user, what is the password, because authentication needs to happen. So, SQLS is a database here and with this user with this password is connecting you that is successful you get a conn object and on the conn object you have something what is known as a cursor. Cursor is nothing but if you think about it is kind of a pointer. So, it can be used to point to either a row or a whole query or a table.

So, you get back a cursor object. So, then this is if you look into this part, this part is nothing but a pure SQL query. So, you take that as a string and pass it on to the execute

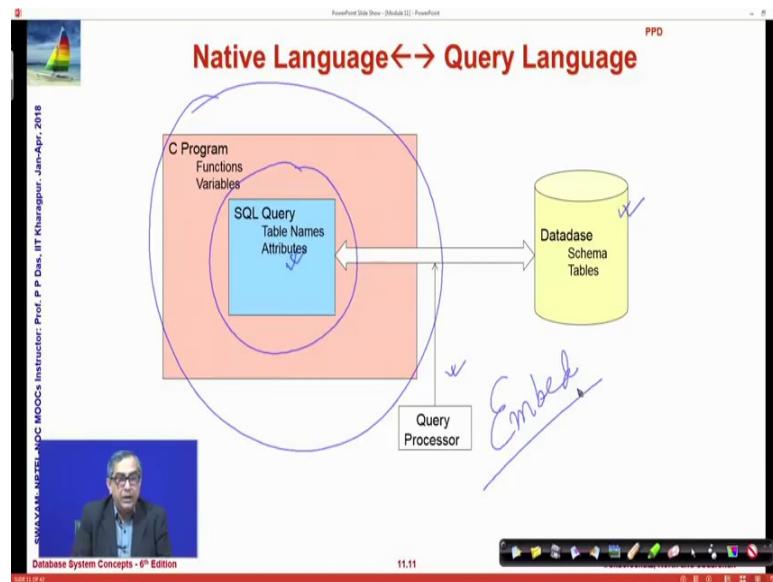
method of cursor. So, what it does is, so this is this has done the first task which is connecting to the database. Now, you are basically putting the query to the database in terms of doing saying that this execute. Similarly, so if that will get executed, so the table is created naturally there is no result to get if you have created a table. Now, you do an insert. So, again this particular record, you can see that within this double quote, this whole insert syntax you take that as a string and just give it to execute as a parameter.

In the third, that you do is do a select. So, you have done this we have inserted a record, created a table inserted a record in that and then now you are doing a select. So, certainly we expect one record to come back. So, these are the SQL executions and then your get back result. So, cursor has another method which is known as fetch one. So, what it will do that if your result table has multiple entries, then the cursor will if will be will start with the first row and fetch one will bring the whole of this first row as a row vector.

So, it will bring in as all the components as one as a python string. And then you check whether it is empty or not, if it is empty then the I mean there is nothing to bring back, so you are done. So, you break otherwise you simply print the row, so you are printing there. And then you go back again this is while true. So, what happens is the cursor will advance to the next row and get you the next row. Again you do the same thing go back it will come back to the same, but once you get an empty result, you know that there is nothing more to proceed and you break.

So, this is what is illustrated then there are some more this particular record is deleted, then again another record is inserted. And another select is done. So, this is how a native program here in this case python can interact with the database and do any of the SQL tasks that we were doing earlier with SQL interface, now we can be done from the python, so that is a basic ODBC mechanism. I particularly chose python to just give you a different flavour, you can we will have assignments on doing it for C and using jdbc on java as well.

(Refer Slide Time: 12:19)



Now, I am come I will come back to the same interaction issue, but this time you can see that I am using a different diagram. The database is the same; this part is the same; the query processing is same, but instead of having a connection library, now I have put the SQL query itself as a part of the native program, the C program. So, this is what is called embedding.

So, you say I embed I put the SQL as a part of C, but naturally SQL is not C. So, we need to put certain additional syntax in C, so that I can directly write SQL as a part of the C program. I am talking about C, again just as an example if this is true for other several other languages which can be used as used for embedding SQL within them.

(Refer Slide Time: 13:09)

The slide is titled "Embedded SQL". It contains the following text and bullet points:

EXEC SQL <embedded SQL statement>;

Note: this varies by language:

- In some languages, like COBOL, the semicolon is replaced with END-EXEC
- In Java embedding uses # SQL{ };

Below the slide, a video player window shows a man speaking, identified as Prof. P. Das, BT Kharagpur, Jan-Apr., 2018. The video player has a progress bar at 11.12 and a red bar at the bottom.

So, this is called the embedded SQL, it works for C, C++ , java fortran etcetera. And the language native language in which your embedding is called the is known as a host language. And basic form of these languages allow that the are come from the System R. So, what is important is this particular statement EXEC-SQL. This EXEC-SQL written inside the body of a C program will tell the C compiler that this part is not C; this part is actually embedded SQL.

And it will be treated differently; it will be compiled by the SQL compiler within the C compiler, so that is the basic structure, so EXEC-SQL. And then you put the pure SQL statement the embedded SQL statement. So, let us go forward and see some of this, there are different syntax for different native language embedding.

(Refer Slide Time: 14:11)

The slide is titled "Embedded SQL (Cont.)". It contains the following text:

- Before executing any SQL statements, the program must first connect to the database. This is done using:
EXEC-SQL connect to server user user-name using password;
Here, *server* identifies the server to which a connection is to be established
- Variables of the host language can be used within embedded SQL statements. They are preceded by a colon (:) to distinguish from SQL variables (e.g., :credit_amount)
- Variables used as above must be declared within DECLARE section, as illustrated below. The syntax for declaring the variables, however, follows the usual host language syntax

EXEC-SQL BEGIN DECLARE SECTION
int credit-amount ;
EXEC-SQL END DECLARE SECTION;

Below the slide, there is a video player window showing a man speaking. A red box highlights this video player.

So, to be able to connect you say EXEC-SQL and connect to server username using password. So, we saw similar things in terms of ODBC based connection these three information needs to be specified which database server, who is the user, what is the password. So, here you say that in this form. And this particular statement you can write as a part of the C program. Now, naturally the question here is in the in the earlier case when we are using the connection library, your results came back through the cursor which you then could deal because the cursors are necessarily objects in your native language. So, in python the cursor was a particular object in the pyodbc library.

But now you have embedded the SQL in terms of your c program. So, naturally the results or whatever you are doing in C which needs to have a communication with the SQL statement need to be differentiated from the SQL names themselves. So, any native language variable that needs to be treated in SQL will be marked with a colon preceding it. So, credit amount of this a colon credit amount, so it becomes a SQL variable provided credit amount itself is a C variable, so that is the basic you know connection mechanics. There are far more details in that, but I am just giving you the glimpse.

Now, any region where you write the SQL, you can write it as exact SQL begin declare section, and END declare section this is where you specify what are the different what are the different C variables C declarations that need to be used for this SQL definition. I will just show you an example soon so that.

(Refer Slide Time: 16:15)

The screenshot shows a PowerPoint slide titled "Embedded SQL (Cont.)". The slide content includes:

- To write an embedded SQL query, we use the `declare c cursor for <SQL query>` statement. The variable `c` is used to identify the query
- Example:
 - From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable `credit_amount` in the host language
 - Specify the query in SQL as follows:

Below the text, there is a block of SQL code:

```
EXEC SQL
declare c cursor for
select ID, name
from student
where tot_cred > :credit_amount
END_EXEC
```

Annotations include a blue curly brace under "declare c cursor for", a blue arrow pointing from "credit_amount" in the SQL code to the variable "credit_amount" in the explanatory text, and a blue circle highlighting "credit_amount" in the explanatory text.

Similar to the ODBC style, you have to declare a cursor. So, this is to write the embedded SQL, you use declared C cursor for such and such SQL queries, so then that C, variable C will become your handle in the C language to be able to answer handle the query results. So, here is an example. So, you can see that credit amount is a variable in the host language.

So, you are using it in SQL with colon credit amount which says that it is this host language variable. So, that you can set a particular credit amount and go with that. And this is the query, and you have set a cursor on that which you can make use of in the exact SQL.

(Refer Slide Time: 17:14)

The slide is titled "Embedded SQL (Cont.)". It contains the following text:

- Example:
 - From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable `credit_amount` in the host language
- Specify the query in SQL as follows:

```
EXEC SQL
declare c cursor for
select ID, name
from student
where tot_cred > :credit_amount
END_EXEC
```
- The variable `c` (used in the cursor declaration) is used to identify the query

A video player window is visible on the right side of the slide, showing a man speaking. The video player has a progress bar at 11.15 and a red status bar at the bottom.

So, let us this is the example continued.

(Refer Slide Time: 17:18)

The slide is titled "Embedded SQL (Cont.)". It contains the following text:

- The `open` statement for our example is as follows:

```
EXEC SQL open c;
```

This statement causes the database system to execute the query and to save the results within a temporary relation. The query uses the value of the host-language variable `credit-amount` at the time the `open` statement is executed.
- The `fetch` statement causes the values of one tuple in the query result to be placed on host language variables.

```
EXEC SQL fetch c into :si, :sn END_EXEC
```

Repeated calls to `fetch` get successive tuples in the query result

A video player window is visible on the right side of the slide, showing a man speaking. The video player has a progress bar at 11.16 and a red status bar at the bottom.

Let us look at other features. You can once you have set a query you can actually execute that by the `open` statement. So, you say `EXEC SQL open c`, the cursor. So, that will execute the query that you have associated with that cursor. And then once that has been done, then you can fetch the results into that cursor one by one; one tuple at a time.